



Best Practices



Best Practices

Publication Date: 2018-10-25

SUSE LLC

10 Canal Park Drive

Suite 200

Cambridge MA 02141

USA

<https://www.suse.com/documentation> 

Contents

1	Introduction	1
1.1	What's Covered in this Guide?	1
1.2	Prerequisites	1
1.3	Network Requirements	2
1.4	Hardware Recommendations	6
2	Managing Your Subscriptions	8
2.1	SUSE Customer Center (SCC)	8
2.2	Disconnected Setup with RMT or SMT (DMZ)	8
	Repository Management Tool (RMT) and Disconnected Setup (DMZ)	9
	• Repository Management Tool (SMT) and Disconnected Setup (DMZ)	10
	• Updating Repositories on Uyuni From Storage Media	11
	• Refreshing Data on the Storage Medium	12
3	Expanded Support	14
3.1	Managing Red Hat Enterprise Linux Clients	14
	Server Configuration for Red Hat Enterprise Linux Channels	14
	• Red Hat Enterprise Linux Channel Management Tips	15
	• Mirroring RHEL Media into a Channel	15
	• Registering RES Salt Minions with Uyuni	16
	• Register a Salt Minion via Bootstrap	18
	• Manual Salt Minion Registration	19
3.2	Preparing Channels and Repositories for CentOS Traditional Clients	19
3.3	Registering CentOS Salt Minions with Uyuni	22
4	Salt Formulas and Uyuni	24
4.1	What are Salt Formulas?	24
4.2	Installing Salt Formulas via RPM	24

- 4.3 File Structure Overview 26
- 4.4 Editing Pillar Data in Uyuni 27
 - Simple edit-group Example 33
- 4.5 Writing Salt Formulas 36
- 4.6 Separating Data 37
- 4.7 Uyuni Generated Pillar Data 39
- 4.8 Formula Requirements 40
- 4.9 Using Salt Formulas with Uyuni 40

5 Configuration Management with Salt 59

- 5.1 Configuration Management Overview 59
- 5.2 State Data: Levels of Hierarchy 59
- 5.3 Salt States Storage Locations 60
- 5.4 Uyuni States 60
- 5.5 Pillar Data 61
- 5.6 Group States 61

6 Salt Minion Scalability 63

- 6.1 Salt Minion Onboarding Rate 63
- 6.2 Minions Running with Unaccepted Salt Keys 63
- 6.3 Salt Timeouts 64
 - Background Information 64 • A Presence Ping Mechanism for Unreachable Salt Minions 65 • Overriding Salt Presence Timeout Values 65 • Salt SSH Minions (SSH Push) 65

7 Activation Key Management 66

- 7.1 What are Activation Keys? 66
- 7.2 Provisioning and Configuration 67

- 7.3 Activation Keys Best Practices 67
 - Key Label Naming 68 • Channels which will be Included 68
- 7.4 Combining Activation Keys 69
- 7.5 Using Activation Keys and Bootstrap with Traditional Clients (Non-Salt) 70
- 7.6 Using Activation Keys when Registering Salt Minions 71
 - Using an Activation Key and Custom Grains File 71 • Using an Activation Key in the Minion Configuration File 72

8 Contact Methods 73

- 8.1 Selecting a Contact Method 73
- 8.2 Default (the Uyuni Daemon rhnsd) 73
 - Configuring Uyuni rhnsd Daemon 74 • Viewing rhnsd Daemon Status 74
- 8.3 Push via SSH 75
 - Configuring the Server for Push via SSH 76 • Using sudo with Push via SSH 76 • Client Registration 77 • API Support for Push via SSH 78 • Proxy Support with Push via SSH 79
- 8.4 Push via Salt SSH 79
 - Overview 80 • Requirements 81 • Bootstrapping 81 • Configuration 82 • Action Execution 82 • Known Limitation 82 • For More Information 82
- 8.5 OSAD 83
 - Enabling and Configuring OSAD 83

9 Advanced Patch Lifecycle Management 86

10 Live Patching with SUSE Manager 87

- 10.1 Introduction 87
- 10.2 Initial Setup Requirements 87
- 10.3 Live Patching Setup 87
- 10.4 Cloning Channels 89

10.5	Removing Non-live Kernel Patches from the Cloned Channels	90
10.6	Promoting Channels	90
10.7	Applying Live Patches to a Kernel	91
11	Uyuni Server Migration	93
11.1	Service Pack Migration Introduction	93
11.2	Service Pack Migration	93
11.3	Upgrading PostgreSQL to Version 9.6	94
11.4	Updating Uyuni	94
11.5	Migrating Uyuni version 3.1 to 3.2	95
	Using YaST	96 • Using zypper 97
11.6	Uyuni Migration from Version 2.1 to Version 3	98
	Prerequisites	99 • Setup the Target Machine 100 • Performing the Migration 101 • Speeding up the Migration 102 • Packages on External Storage 103 • Troubleshooting a Broken Web UI after Migration 103 • Example Session 103
12	Client Migration	109
12.1	Upgrading SLE 12 SPx to version 15	109
	System Upgrade Preparation	110 • Sample Autoinstallation Script for System Upgrade (SLES 12 SP3 to SLES 15) 112
12.2	Migrating SLE 12 or later to version 12 SP3	113
13	PostgreSQL Database Migration	121
13.1	New Uyuni Installations	121
13.2	Migrating an Existing Installation	121
13.3	Performing a Fast Migration	122
13.4	Typical Migration Sample Session	123

14 Backup and Restore 127

- 14.1 Backing up Uyuni 127
- 14.2 Administering the Database with smdba 129
- 14.3 Database Backup with smdba 130
 - Performing a Manual Database Backup 130 • Scheduling Automatic Backups 131
- 14.4 Restoring from Backup 132
- 14.5 Archive Log Settings 133
- 14.6 Retrieving an Overview of Occupied Database Space 133
- 14.7 Moving the Database 134
- 14.8 Recovering from a Crashed Root Partition 135
- 14.9 Database Connection Information 136

15 Authentication Methods 137

- 15.1 Authentication Via PAM 137
- 15.2 Authentication Via eDirectory and PAM 138
- 15.3 Example Quest VAS Active Directory Authentication Template 139

16 Using a Custom SSL Certificate 140

- 16.1 Prerequisites 140
- 16.2 Setup 140
- 16.3 Using a Custom Certificate with SUSE Manager Proxy 141

17 Troubleshooting 142

- 17.1 Registering Cloned Salt Minions 142
- 17.2 Registering Cloned Traditional Systems 143

17.3	Typical OSAD/jabberd Challenges	145
	Open File Count Exceeded	145 • jabberd Database Corruption 146 • Capturing XMPP Network Data for Debugging Purposes 147 • Engineering Notes: Analyzing Captured Data 148
17.4	RPC Connection Timeout Settings	149
17.5	Client/Server Package Inconsistency	150
17.6	Corrupted Repository Data	150
17.7	Unable to Get Local Issuer Certificate	151
18	Additional Resources	152
18.1	Learning YAML Syntax for Salt States	152
18.2	Getting Started with Jinja Templates	152
18.3	Salt Best Practices	152
19	A Uyuni 2.1 and 3.2 Product Comparison	153
A	GNU Licenses	157
20	GNU Free Documentation License	158

1 Introduction

This document targets system administrators.

1.1 What's Covered in this Guide?

This document describes SUSE recommended best practices for Uyuni. This information has been collected from a large number of successful Uyuni real world implementations and includes feedback provided by product management, sales and engineering.



Note: Uyuni Version Information

In this manual if not other specified, Uyuni version 3.2 is assumed and this version is required if a feature is discussed. Uyuni 3.2 and Uyuni 3.2 Proxy were originally released as a SLES 12 SP3 extension. Whenever features of the Uyuni3.2 host operating system are documented and not other specified version 12 SP3 is assumed.

This chapter will discuss the following topics:

- Prerequisites
- Network Requirements
- Hardware Requirements

1.2 Prerequisites

Purchased Registration Keys. During initial setup Uyuni will request a product Registration Key. This key will be provided to you after purchasing the product. You can find your key located under your SUSE Customer Center account. Log-in with your SUSE Customer Center credentials or register for a new account. -<https://scc.suse.com>

Evaluation Keys. If you wish to run a test system (non-production) a 60 day evaluation key may be obtained. On the Uyuni product page click *TRY SUSE MANAGER*. The evaluation key limits the number of systems that may be registered with SUSE Manager to 10. For more information see:

- <https://www.suse.com/products/suse-manager/> 

SCC Organization Credentials. During setup you will also be asked to enter your SUSE Customer Center Organization Credentials.

Users and Passwords. During both the SUSE Linux Enterprise installation and setup of Uyuni several users and passwords will be created:

- SUSE Linux Enterprise root user account
- PostgreSQL database user and password
- Certificate of Authority password
- Uyuni administrator user and password



Tip: Safe Passwords

Maintain security by creating safe passwords. Store passwords within a secure location. Use the following guidelines when creating your passwords.

- At least 8 characters long
- Should contain uppercase characters A B C
- Should contain lowercase characters a b c
- Should contain numbers 1 2 3
- Should contain symbols ~ ! @ #

1.3 Network Requirements

Uyuni and SUSE Manager Proxy both contact several external addresses in order to maintain updates and subscriptions. The following lists provide the up-to-date hostnames for each service requiring permission when used in combination with corporate firewall and content filters.

SUSE CUSTOMER CENTER HOSTNAMES (REQUIRED)

- <https://scc.suse.com> 
- <https://updates.suse.com> 

NOVELL CUSTOMER CENTER HOSTNAMES (LEGACY)

- <https://secure-www.novell.com> ↗
- <https://nu.novell.com> ↗

For Uyuni to function properly it requires the following pre-configured components within your network.

Important: Websocket Support

If SUSE Manager is accessed via an HTTP proxy (Squid, etc) the proxy must support websocket connections.

Networking Hardware. The following table provides networking hardware info. As Uyuni will likely be managing a large number of systems (quite possibly numbering in hundreds or even thousands), networking hardware that increases bandwidth becomes increasingly more valuable.

Hardware	Recommended
100Mbits/s Link	Non-production test server
1Gb/s Link	Production Server

DHCP Server. The purpose of the Dynamic Host Configuration Protocol (DHCP) is to assign network settings centrally (from a server) rather than configuring them locally on each and every workstation. A host configured to use DHCP does not have control over its own static address. It is enabled to configure itself completely and automatically according to directions from the server. A DHCP server supplies not only the IP address and the netmask, but also the host name, domain name, gateway, and name server addresses for the client to use. For more information on configuring DHCP see also:

- https://www.suse.com/documentation/sles-12/book_sle_admin/data/cha_dhcp.html ↗

FQDN (Fully Qualified Domain Name). DNS assists in assigning an IP address to one or more names and assigning a name to an IP address. In Linux, this conversion is usually carried out by a special type of software known as bind. The machine that takes care of this conversion is called a *name server*. The names make up a hierarchical system in which each name component is

separated by a period. The name hierarchy is, however, independent of the IP address hierarchy described above. Consider a complete name, such as `jupiter.example.com`, written in the format `hostname.domain`. A full name, referred to as a *fully qualified domain name* (FQDN), consists of a host name and a domain name (`example.com`). For more information on configuring a name server see also:

- https://www.suse.com/documentation/sles-12/book_sle_admin/data/sec_basicnet_nameres.html ↗

DNS (Dynamic Name System) Server. A DNS Server is required for resolving domain names and host names into IP addresses. For example, the IP address `192.168.2.100` could be assigned to the host name `jupiter`. In the case of Uyuni the DNS server must be resolvable both via DNS and reverse lookup. For more information on configuring DNS see also:

- https://www.suse.com/documentation/sles-12/book_sle_admin/data/cha_dns.html ↗



Important: Microsoft NT Lan Manager Compatibility

Microsoft NT Lan Manager can be configured for use with basic authentication and will work with Uyuni but authentication using native (NTLM) Microsoft protocols is not supported.

Open Port List. During the setup process of Uyuni all required ports will be opened automatically. The following tables provide you with an overview of ports which are used by Uyuni.

TABLE 1.1: **REQUIRED SERVER PORTS**

Port	Protocol	Description
22	TCP	SSH
67	UDP	DHCP
69	UDP	TFTP, used to support PXE services
80	TCP	HTTP, used in some bootstrap cases
123	UDP	NTP time service

Port	Protocol	Description
443	TCP	HTTPS, used for Web UI, client, Proxy server, and API traffic
4505	TCP	Salt, used by the Salt-master to accept communication requests from minions
4506	TCP	Salt, used by the Salt-master to accept communication requests from minions
5222	TCP	XMPP client, used for communications with the <u>osad</u> daemon on traditional client systems
5269	TCP	XMPP server, used for pushing actions to SUSE Manager Proxy

For more information, see [Port Listing \(advanced.topics.app.ports.server.xml#at-ports\)](#).



Tip: Denying External Network Access

When your network requires denying external network access to and from Uyuni, an RMT or SMT Server may be registered against Uyuni. The RMT or SMT server can then be used to synchronize the necessary SUSE repositories. For more information on utilizing an RMT or SMT Server, see: [Section 2.2, “Disconnected Setup with RMT or SMT \(DMZ\)”](#).



Note: Blocking Port 80

Port 80 may be blocked as traffic is automatically redirected through port 443. It should be noted you will lose redirection. Keep in mind you will need additional ports open when using traditional clients in combination with osad (XMPP TCP 5222).

1.4 Hardware Recommendations

This section provides tested production recommendations for small to mid size networks that will be managed by Uyuni.

Hardware	Recommended
CPU	Multi-core 64bit CPU (x86_64, ppc64le).
RAM	Minimum 4 GB + for test server
	Minimum 16 GB + for base installation
	Minimum 32 GB + for a production server
Free Disk Space	Minimum 100 GB + for root partition
	Minimum 50 GB + for <u>/var/lib/pgsql</u>
	Minimum 50 GB per SUSE product, or 200 GB per Red Hat product <u>/var/pacewalk</u>

Advised Number of CPUs. Review the following list for CPU recommendations.

- Connecting 200 systems or less to Uyuni : 4 CPUs
- Connecting 500 systems or less to Uyuni : 4-8 CPUs
- When implementing RHEL channels: 8 CPUs

Disk Space. Uyuni stores information in several directories. For these directories it is strongly recommend that you create separate file-systems or use an NFS share. During installation one VG will be created that contains all disks selected during installation. Therefore the first disk should be large enough to contain the OS. Normally 20GB - 50GB is sufficient. A 50 GB partition would be the recommended size. The following directories should be created on a separate file-system.

- /var/pacewalk This directory will contain all rpm's. Each RPM will be stored only once. The needed size will depend on the number of channels and type of channels that will be downloaded. The general rule would be that per SUSE Service Pack (including SUSE Red-Hat Expanded Support) around 50 GB should be enough. An extra 150 GB for RES/CentOS

repositories should be added on top. If other non-enterprise distributions (eg OpenSUSE) are added, calculated 50 GB per distribution. This directory could also be stored on an NFS share.

- /var/lib/pgsql This directory contains the PostgreSQL database. Recommended is to create a file-system of 50 GB. This volume should be monitored, because a full file-system where the database is running on can cause unexpected errors (and this even months after it happened).
- /srv/tftpboot If PXE/cobbler is used, this directory will contain the images (initrd and linux) for all created auto-installation profiles. Each image is around 50 MB. Depending on the number of profiles a decision has to be made if it would be useful to move this directory to a separate file-system.
- /var/log As SUSE Manager writes a large number of logs, it is recommended to create a separate file-system for /var/log. The size should be around 20 GB.
- /var/pacewalk/db_backup For the backup of the PostgreSQL database, it is recommended to create a separate directory. As the database can be rather large, it is advised to mount it on a separate file-system. A safe estimate would be to provide twice space as for the directory created for /var/lib/pgsql.


Supported Databases. Uyuni 3 and later no longer provides support for an external Oracle database. The default database is an embedded PostgreSQL. During Uyuni setup the database will be created and configured.

2 Managing Your Subscriptions

There are two methods for managing your subscriptions. Both methods access SUSE Customer Center and provide specialized benefits.

- Directly connecting to SUSE Customer Center is the recommended default way of managing your Uyuni server.
- If you have special network security requirements which do not allow access from your internal network to the internet then you can use SUSE Linux Enterprise Server 12 running the Repository Management Tool (RMT) or Repository Management Tool (SMT). These tools will contact SUSE Customer Center from a system connected to the external network and obtain updates for your clients which you may then mount on your internal Uyuni server. This is the preferred method for managing client systems within a highly secure network infrastructure.

2.1 SUSE Customer Center (SCC)

SUSE Customer Center (SCC) is the central place to manage your purchased SUSE subscriptions, helping you access your update channels and get in contact with SUSE experts. The user-friendly interface gives you a centralized view of all your SUSE subscriptions, allowing you to easily find all subscription information you need. The improved registration provides faster access to your patches and updates. SUSE Customer Center is also designed to provide a common platform for your support requests and feedback. Discover a new way of managing your SUSE account and subscriptions via one interface—anytime, anywhere. For more information on using SUSE Customer Center , see <https://scc.suse.com/docs/userguide> .

2.2 Disconnected Setup with RMT or SMT (DMZ)

If it is not possible to connect Uyuni directly or via a proxy to the Internet, a disconnected setup in combination with RMT or SMT is the recommended solution. In this scenario, RMT or SMT stays in an “external” network with a connection to SUSE Customer Center and synchronizes the software channels and repositories on a removable storage medium. Then you separate the storage medium from RMT or SMT, and mount it locally on your Uyuni server to read the updated data.

RMT. The successor of SMT and currently runs on the following systems:

- SUSE Linux Enterprise 15 (when available)
- Temporarily (for testing only): 12 SP2, and 12 SP3
- Not officially supported: openSUSE Leap 42.2, Leap 42.3, and openSUSE Tumbleweed

RMT allows you to provision updates for all of your devices running a product based on SUSE Linux Enterprise 12 SPx and later as well as openSUSE Leap.

SMT. The predecessor of RMT and is no longer actively developed. It runs on SUSE Linux Enterprise Server 12 SPx and allows you to provision updates for products based on SUSE Linux Enterprise 12 SPx and earlier. You will still need it, if you want to update SUSE Linux Enterprise 11 clients.

2.2.1 Repository Management Tool (RMT) and Disconnected Setup (DMZ)

The following procedure will guide you through using RMT. It will work best with a dedicated RMT instance per Uyuni .

PROCEDURE: RMT: FETCHING REPOSITORY DATA FROM SUSE CUSTOMER CENTER

1. Configure RMT in the external network with SCC. For details about configuring RMT, see the official guide (when available).
 - a. Preparation work:
Run `rmt-cli sync` to download available products and repositories data for your organization from SCC.
Run `rmt-cli products list --all` to see the list of products that are available for your organization.
Run `rmt-cli repos list --all` to see the list of all repositories available.
 - b. With `rmt-cli repos enable` enable repositories you want to mirror.
 - c. With `rmt-cli products enable` enable products. For example, to enable SLES_15:

```
rmt-cli product enable sles/15/x86_64
```

2. Using RMT, mirror all required repositories.

3. Get the required JSON responses from SCC and save them as files at the specified path (for example, /mnt/usb).



Important: Write Permissions for RMT User

The directory being written to must be writeable for the same user as the rmt service. The rmt user setting is defined in the cli section of /etc/rmt.conf .

Enter:

```
{prompt.root}rmt-cli export data /mnt/usb
```

4. Export settings about repositories to mirror to the specified path (in this case, /mnt/usb); this command will create a repos.json file there:

```
{prompt.root}rmt-cli export settings /mnt/usb
```

5. Mirror the repositories according to the settings in the repos.json file to the specified path (in this case, /mnt/usb).

```
{prompt.root}rmt-cli export repos /mnt/usb
```


6. Unmount the storage medium and carry it securely to your Uyuni server.

On the Uyuni server, continue with *Section 2.2.3, “Updating Repositories on Uyuni From Storage Media”*.

2.2.2 Repository Management Tool (SMT) and Disconnected Setup (DMZ)

The following procedure will guide you through using SMT.

PROCEDURE: SMT: FETCHING REPOSITORY DATA FROM SUSE CUSTOMER CENTER

1. Configure SMT in the external network with SCC. For details about configuring SMT with SUSE Linux Enterprise 12, see https://www.suse.com/documentation/sles-12/book_smt/data/book_smt.html .
2. Using SMT, mirror all required repositories.
3. Create a “database replacement file” (for example, /tmp/dbrepl.xml).

```
{prompt.root}smt-sync --createdbreplacementfile /tmp/dbrepl.xml
```

1. Mount a removable storage medium such as an external hard disk or USB flash drive.
2. Export the data to the mounted medium:

```
smt-sync --todir /media/disk/  
smt-mirror --dbreplfile /tmp/dbrepl.xml --directory /media/disk \  
--fromlocalsmt -L /var/log/smt/smt-mirror-export.log
```



Important: Write Permissions for SMT User

The directory being written to must be writeable for the same user as the `smt` daemon (`user = smt`). The `smt` user setting is defined in `/etc/smt.conf`. You can check if the correct user is specified via the following command:

```
{prompt.root}egrep '^smtUser' /etc/smt.conf
```

+

+ .Keeping the Disconnected Server Up-to-date NOTE: `smt-sync` also exports your subscription data. To keep Uyuni up-to-date with your subscriptions, you must frequently import and export this data.

+

1. Unmount the storage medium and carry it securely to your Uyuni server.

On the Uyuni server, continue with [Section 2.2.3, “Updating Repositories on Uyuni From Storage Media”](#).

2.2.3 Updating Repositories on Uyuni From Storage Media

This procedure will show you how to update the repositories on the Uyuni server from the storage media.

PROCEDURE: UPDATING THE UYUNISERVER FROM THE STORAGE MEDIUM

1. Mount the storage medium on your Uyuni server (for example, at `/media/disk`).

2. Specify the local path on the Uyuni server in `/etc/rhn/rhn.conf`:

```
server.susemanager.fromdir = /media/disk
```

This setting is mandatory for SUSE Customer Center and `mgr-sync`.

3. Restart Tomcat:

```
systemctl restart tomcat
```

1. Before performing another operation on the server execute a full sync:

```
mgr-sync refresh # SCC (fromdir in rhn.conf required!)
```

2. `mgr-sync` can now be executed normally:

```
mgr-sync list channels  
mgr-sync add channel channel-label
```



Warning: Data Corruption

The disk must always be available at the same mount point. To avoid data corruption, do not trigger a sync, if the storage medium is not mounted. If you have already added a channel from a local repository path, you will not be able to change its URL to point to a different path afterwards.

Up-to-date data is now available on your Uyuni server and is ready for updating client systems. According to your maintenance windows or update schedule refresh the data on the storage medium with RMT or SMT.

2.2.4 Refreshing Data on the Storage Medium

PROCEDURE: REFRESHING DATA ON THE STORAGE MEDIUM FROM RMT OR SMT

1. On your Uyuni server, unmount the storage medium and carry it to your RMT or SMT.
2. On your RMT or SMT system, continue with the synchronization step.



Warning: Data Corruption

The storage medium must always be available at the same mount point. To avoid data corruption, do not trigger a sync if the storage medium is not mounted.

This concludes using RMT or SMT with Uyuni .

3 Expanded Support

In the following sections find information about Red Hat and Centos clients.

3.1 Managing Red Hat Enterprise Linux Clients

The following sections provide guidance on managing Red Hat Expanded Support clients, this includes Salt minions and traditional systems.

3.1.1 Server Configuration for Red Hat Enterprise Linux Channels

This section provides guidance on server configuration for Red Hat Enterprise Linux Channels provided by SUSE.

- Minimum of 8 GB RAM and at least two physical or virtual CPUs. Taskomatic will use one of these CPUs.
- Taskomatic requires of minimum of 3072 MB RAM. This should be set in /etc/rhn/rhn.conf:

```
taskomatic.java.maxmemory=3072
```

- Provision enough disk space. /var/spacwalk contains all mirrored RPMs. For example, Red Hat Enterprise Linux 6 x86_64 channels require 90 GB and more.
- LVM or an NFS mount is recommended.
- Access to RHEL 5/6/7 Subscription Media.



Warning: Access to RHEL Media or Repositories

Access to Red Hat base media repositories and RHEL installation media is the responsibility of the user. Ensure that all your RHEL systems obtain support from RHEL or all your RHEL systems obtain support from SUSE. If you do not follow these practices you may violate terms with Red Hat.

3.1.2 Red Hat Enterprise Linux Channel Management Tips

This section provides tips on Red Hat Enterprise Linux channel management.

- The base parent distribution Red Hat Enterprise Linux channel per architecture contains zero packages. No base media is provided by SUSE. The RHEL media or installation ISOs should be added as child channels of the Red Hat Enterprise Linux parent channel.
- The Red Hat Enterprise Linux and tools channels are provided by SUSE Customer Center (SCC) using `mgr-sync`.
- It can take up to 24 hours for an initial channel synchronization to complete.
- When you have completed the initial synchronization process of any Red Hat Enterprise Linux channel it is recommended to clone the channel before working with it. This provides you with a backup of the original synchronization.

3.1.3 Mirroring RHEL Media into a Channel

The following procedure guides you through setup of the RHEL media as a Uyuni channel. All packages on the RHEL media will be mirrored into a child channel located under RES 5/6/7 distribution per architecture.

PROCEDURE: MIRRORING RHEL MEDIA INTO A CHANNEL

1. Create a new Channel by log in to the Web UI and selecting *Channels > Manage Software Channels > Create Channel*.
2. Fill in basic channel details and add the channel as a child to the corresponding RES 5/6/7 distribution channel per architecture from SCC. The base parent channel should contain zero packages.
3. Modify the RES 5/6/7 activation key to include this new child channel.
4. As root on the Uyuni command line copy the ISO to the `/tmp` directory.
5. Create a directory to contain the media content:

```
{prompt.root}mkdir -p /srv/www/htdocs/pub/rhel
```

6. Mount the ISO:

```
{prompt.root}mount -o loop /tmp/name_of_iso /srv/www/htdocs/pub/rhel
```

7. Start `spacewalk-repo-sync` to synchronize Red Hat Enterprise Linux 7 packages:

```
{prompt.root}spacewalk-repo-sync -c channel_name -u https://127.0.0.1/pub/rhel/
Repo URL: https://127.0.0.1/pub/rhel/
Packages in repo:          [...]
Packages already synced:   [...]
Packages to sync:          [...]
[...]
```

To synchronize RES 5/6 packages:

```
{prompt.root}spacewalk-repo-sync -c channel_name -u https://127.0.0.1/pub/rhel/
Server/
Repo URL: https://127.0.0.1/pub/rhel/Server/
Packages in repo:          [...]
Packages already synced:   [...]
Packages to sync:          [...]
[...]
```

8. When the channel has completed the synchronization process you can use the channel as any normal Uyuni channel.

Attempting to synchronize the repository will sometimes fail with this error:

```
[Errno 256] No more mirrors to try.
```

To troubleshoot this error, look at the HTTP protocol to determine if `spacewalk-repo-sync` is running:

PROCEDURE: DEBUG SPACEWALK-REPO-SYNC

1. Start debugging mode with `export URLGRABBER_DEBUG=DEBUG`
2. Check the output of `/usr/bin/spacewalk-repo-sync --channel <channel-label> --type yum`
3. If you want to disable debug mode, use `unset URLGRABBER_DEBUG`

3.1.4 Registering RES Salt Minions with Uyuni

This section will guide you through registering RHEL minions with Uyuni.

This section assumes you have updated your server to the latest patch level.

3.1.4.1 Synchronizing Appropriate Red Hat Enterprise Linux Channels

Ensure you have the corresponding Red Hat Enterprise Linux product enabled and required channels have been fully synchronized:

RHEL 7.X

- Product: Red Hat Enterprise Linux 7
- Mandatory channels: rhel-x86_64-server-7 , res7-suse-manager-tools-x86_64 , res7-x86_64 `systemitem >`

RHEL 6.X

- Product: Red Hat Enterprise Linux 6
- Mandatory channels: rhel-x86_64-server-6 , res6-suse-manager-tools-x86_64 , res6-x86_64



Tip: Checking Synchronization Progress

To check if a channel has finished synchronizing you can do one of the following:

- From the UyuniWeb UI browse to *Admin > Setup Wizard* and select the *SUSE Products* tab. Here you will find a percent completion bar for each product.
- Alternatively, you may check the synchronization log file located under /var/log/rhn/reposync/channel-label.log using `cat` or the `tailf` command. Keep in mind that base channels can contain multiple child channels. Each of these child channels will generate its own log during the synchronization progress. Do not assume a channel has finished synchronizing until you have checked all relevant log files including base and child channels.

Create an activation key associated with the Red Hat Enterprise Linux channel.

3.1.4.2 Creating a Bootstrap Repository

The following procedure demonstrate creating a bootstrap repository for RHEL:

1. On the server command line as root, create a bootstrap repo for RHEL with the following command:

```
mgr-create-bootstrap-repo RHEL_activation_channel_key
```

2. Rename `bootstrap.sh` to `resversion-bootstrap.sh`:

```
{prompt.root}cp bootstrap.sh res7-bootstrap.sh
```

3.1.5 Register a Salt Minion via Bootstrap

The following procedure will guide you through registering a Salt minion using the bootstrap script.

PROCEDURE: REGISTRATION USING THE BOOTSTRAP SCRIPT

1. For your new minion download the bootstrap script from the Uyuni server:

```
wget --no-check-certificate https://`server`/pub/bootstrap/res7-bootstrap.sh
```

2. Add the appropriate `res-gpg-pubkey--key` to the `ORG_GPG_KEY` key parameter, comma delimited in your `res7-bootstrap.sh` script. These are located on your Uyuni server at:

```
http://`server`/pub/
```

3. Make the `res7-bootstrap.sh` script executable and run it. This will install necessary Salt packages from the bootstrap repository and start the Salt minion service:

```
{prompt.root}chmod +x res7-bootstrap.sh{prompt.root}./res7-bootstrap.sh
```

4. From the Uyuni Web UI select *Salt* > *Keys* and accept the new minion's key.

Important: Troubleshooting Bootstrap

If bootstrapping a minion fails it is usually caused by missing packages. These missing packages are contained on the RHEL installation media. The RHEL installation media should be loop mounted and added as a child channel to the Red Hat Enterprise Linux channel. See the warning in [Section 3.1, “Managing Red Hat Enterprise Linux Clients”](#) on access to RHEL Media.

3.1.6 Manual Salt Minion Registration

The following procedure will guide you through the registration of a Salt minion manually.

1. Add the bootstrap repository:

```
yum-config-manager --add-repo https://`server`/pub/repositories/res/7/bootstrap
```

2. Install the salt-minion package:

```
{prompt.root}yum install salt-minion
```

3. Edit the Salt minion configuration file to point to the Uyuni server:

```
{prompt.root}mkdir /etc/salt/minion.d{prompt.root}echo "master:`server_fqdn`" > /  
etc/salt/minion.d/susemanager.conf
```

4. Start the minion service:

```
{prompt.root}systemctl start salt-minion
```

5. From the Uyuni Web UI select the *Salt* > *Keys* and accept the new minion's key.

3.2 Preparing Channels and Repositories for CentOS Traditional Clients

This following section provides an example procedure for configuring CentOS channels and repositories and finally registering a CentOS client with Uyuni.

These steps will be identical for Scientific Linux and Fedora.

PROCEDURE: PREPARING CHANNELS AND REPOSITORIES

1. As root install spacewalk-utils on your Uyuni server:

```
zypper in spacewalk-utils
```

Important: Supported Tools

The spacewalk-utils package contains a collection of upstream command line tools which provide assistance with spacewalk administrative operations. You will be using the `spacewalk-common-channels` tool. Keep in mind SUSE only provides support for `spacewalk-clone-by-date` and `spacewalk-manage-channel-lifecycle` tools.

2. Run the `spacewalk-common-channels` script to add the CentOS7 base, updates, and Spacewalk client channels.

```
{prompt.root}spacewalk-common-channels -u admin -p`secret`-a x86_64  
'centos7'{prompt.root}spacewalk-common-channels -u admin -p`secret`-a x86_64  
'centos7-updates'{prompt.root}spacewalk-common-channels -u admin -p`secret`-a  
x86_64 'spacewalk26-client-centos7'
```

Note: Required Channel References

The `/etc/rhn/spacewalk-common-channels.ini` must contain the channel references to be added. If a channel is not listed, check the latest version here for updates:

<https://github.com/spacewalkproject/spacewalk/tree/master/utils> ↗

3. From the Web UI select *Main Menu > Software > Manage Software Channels > Overview*. Select the base channel you want to synchronize, in this case `CentOS7 (x86_64)`. Select *Repositories > Sync*. Check the channels you want to synchronize and then click the *Sync Now* button or, optionally, schedule a regular synchronization time.
4. Copy all relevant GPG keys to `/srv/www/htdocs/pub`. Depending on what distribution you are interested in managing these could include an EPEL key, SUSE keys, Red Hat keys, and CentOS keys. After copying these you can reference them in a comma-delimited list within your bootstrap script (see *Procedure: Preparing the Bootstrap Script*).

- CentOS7 key files: <http://mirror.centos.org/centos/RPM-GPG-KEY-CentOS-7> ↗
- EPEL key file: <http://mirrors.kernel.org/fedora-epel/RPM-GPG-KEY-EPEL-7> ↗
- Spacewalk key: <http://spacewalk.redhat.com/yum/RPM-GPG-KEY-spacewalk-2015> ↗
- Red Hat keys: <http://www.redhat.com/contact/security-response-team/gpg-keys.html> ↗

5. Install and setup a CentOS 7 client with the default installation packages.
6. Ensure the client machine can resolve itself and your Uyuni server via DNS. Validate that there is an entry in `/etc/hosts` for the real IP address of the client.
7. Create an activation key (`centos7`) on the Uyuni server that points to the correct parent/child channels, including the CentOS base repo, updates, and Spacewalk client.

Now prepare the bootstrap script.

PROCEDURE: PREPARING THE BOOTSTRAP SCRIPT

1. Create/edit your bootstrap script to correctly reflect the following:

```
# can be edited, but probably correct (unless created during initial install):

# NOTE: ACTIVATION_KEYS *must* be used to bootstrap a client machine.

ACTIVATION_KEYS=1-centos7

ORG_GPG_KEY=res.key,RPM-GPG-KEY-CentOS-7,suse-307E3D54.key,suse-9C800ACA.key,RPM-
GPG-KEY-spacewalk-2015

FULLY_UPDATE_THIS_BOX=0

yum clean all
# Install the prerequisites
yum -y install yum-rhn-plugin rhn-setup
```

2. Add the following lines to the bottom of your script, (just before `echo "-bootstrap complete -"`):

```
# This section is for commands to be executed after registration
mv /etc/yum.repos.d/Cent* /root/
yum clean all
chkconfig rhnsd on
```

```
chkconfig osad on
service rhnsd restart
service osad restart
```

3. Continue by following normal bootstrap procedures to bootstrap the new client.

3.3 Registering CentOS Salt Minions with Uyuni

The following procedure will guide you through registering a CentOS Minion.



Warning: Support for CentOS Patches

CentOS uses patches originating from CentOS is not officially supported by SUSE . See the matrix of Uyuni clients on the main page of the Uyuni wiki, linked from the *Quick Links* section: https://wiki.microfocus.com/index.php?title=SUSE_Manager ↗

PROCEDURE: REGISTER A CENTOS 7 MINION

1. Add the Open Build Service repo for Salt:

```
{prompt.root}yum-config-manager --add-repo http://download.opensuse.org/
repositories/systemsmanagement:/saltstack:/products/RHEL_7/
```

2. Import the repo key:

```
{prompt.root}rpm --import http://download.opensuse.org/repositories/
systemsmanagement:/saltstack:/products/RHEL_7/repodata/repomd.xml.key
```

3. Check if there is a different repository that contains Salt. If there is more than one repository listed disable the repository that contains Salt apart from the OBS one.

```
{prompt.root}yum list --showduplicates salt
```

4. Install the Salt minion:

```
{prompt.root}yum install salt salt-minion
```

5. Change the Salt configuration to point to the Uyuni server:

```
{prompt.root}mkdir -p /etc/salt/minion.d{prompt.root}echo "master:`server_fqdn`" > /
etc/salt/minion.d/susemanager.conf
```

6. Restart the minion

```
{prompt.root}systemctl restart salt-minion
```

7. Proceed to *Main Menu* > *Salt* > *Keys* from the Web UI and accept the minion's key.

4 Salt Formulas and Uyuni

This chapter provides an introduction for using Salt Formulas with Uyuni. Creation of custom formulas will also be introduced.

4.1 What are Salt Formulas?

Formulas are collections of Salt States that have been pre-written by other Salt users and contain generic parameter fields. Formulas allow for reliable reproduction of a specific configuration again and again. Formulas can be installed from RPM packages or an external git repository. The following list will help you in making the decision to use a state or a formula.

FORMULA TIPS

- When writing states for trivial tasks, formulas are probably not worth the time investment.
- For large, non-trivial configurations use formulas.
- Formulas and States both act as a kind of configuration documentation. Once written and stored you will have a snapshot of what your infrastructure should look like.
- You can greatly decrease the amount of time required for writing a formula by grabbing one of the many pre-written formulas on github. There are pre-written formulas for hundreds of configurations. You can find these located at <https://github.com/saltstack-formulas> ↗.



Note: Formula with Forms Improvements

- Formula data can now be managed via the XMLRPC API
- New types were added to help SUSE Manager handle almost any kind of upstream formula.

4.2 Installing Salt Formulas via RPM

SUSE releases formulas as RPM packages.



Note: Formula Channel Location

Available formulas can be located within the SUSE-Manager-Server-3.2-Pool channel.

PROCEDURE: INSTALLING SALT FORMULAS FROM AN RPM

1. To search for available formulas, execute the following command on your Uyuni server:

```
zypper se --type package formula
```

You will see a list of available Salt formulas:

S	Name	Summary
Type		

	locale-formula	Locale Salt Formula for SUSE Manager
package		

2. For more information about a formula, run the following command:

```
zypper info locale-formula
```

```
Information for package locale-formula:
-----
Repository: SUSE-Manager-Server-{productnumber}-Pool
Name:      locale-formula
Version:   0.2-1.1
Arch:      noarch
Vendor:    SUSE LLC <https://www.suse.com/>
Support Level: Level 3
Status:    not installed
Installed Size: 47.9 KiB
Installed: No
Source package : locale-formula-0.2-1.1.src
Summary       : Locale Salt Formula for SUSE Manager
Description    : Salt Formula for SUSE Manager. Sets up the locale.
```

3. To install a formula run as root:

```
zypper in locale-formula
```

4.3 File Structure Overview

RPM-based formulas must be placed in a specific directory structure to ensure proper functionality. A formula always consists of two separate directories: The states directory and the metadata directory. Folders in these directories need to have an exactly matching name, for example locale.

The Formula State Directory

The formula states directory contains anything necessary for a Salt state to work independently. This includes .sls files, a map.jinja file and any other required files. This directory should only be modified by RPMs and should not be edited manually. For example, the locale-formula states directory is located in:

```
/usr/share/susemanager/formulas/states/locale/
```

The Formula Metadata Directory

The metadata directory contains a form.yml file which defines the forms for Uyuni and an optional metadata.yml file that can contain additional information about a formula. For example, the locale-formula metadata directory is located in:

```
/usr/share/susemanager/formulas/metadata/locale/
```

Custom Formulas

Custom formula data or (non-RPM) formulas need to be placed into any state directory configured as a Salt file root:

State directory

Custom state formula data need to be placed in:

```
/srv/salt/custom/
```

Metadata Directory

Custom metadata (information) need to be placed in:

```
/srv/formula_metadata/custom/
```

All custom folders located in the following directories needs to contain a form.yml file. These files are detected as form recipes and may be applied to groups and systems from the Web UI:

```
/srv/salt/formula_metadata/`custom-formula-name`/form.yml
```

4.4 Editing Pillar Data in Uyuni

Uyuni requires a file called `form.yml`, to describe how formula data should look within the Web UI. `form.yml` is used by Uyuni to generate the desired form, with values editable by a user.

For example, the `form.yml` that is included with the locale-formula is placed in:

```
/usr/share/susemanager/formulas/metadata/locale/form.yml
```

See part of the following locale-formula example:

```
# This file is part of locale-formula.
#
# Foobar is free software: you can redistribute it and/or modify
# it under the terms of the GNU General Public License as published by
# the Free Software Foundation, either version 3 of the License, or
# (at your option) any later version.
#
# Foobar is distributed in the hope that it will be useful,
# but WITHOUT ANY WARRANTY; without even the implied warranty of
# MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
# GNU General Public License for more details.
#
# You should have received a copy of the GNU General Public License
# along with Foobar. If not, see <http://www.gnu.org/licenses/>.

timezone:
  $type: group

  name:
    $type: select
    $values: ["CET",
              "CST6CDT",
              "EET",
              "EST",
              "EST5EDT",
              "GMT",
              "GMT+0",
              "GMT-0",
              "GMT0",
              "Greenwich",
              "HST",
              "MET",
              "MST",
              "MST7MDT",
              "NZ",
              "NZ-CHAT",
```

```

    "Navajo",
    "PST8PDT",
    "UCT",
    "UTC",
    "Universal",
    "W-SU",
    "WET",
    "Zulu",
    "Etc/GMT+1",
    "Etc/GMT+2",
    "Etc/GMT+3",
    "Etc/GMT+4",
    "Etc/GMT+5",
    "Etc/GMT+6",
    "Etc/GMT+7",
    "Etc/GMT+8",
    "Etc/GMT+9",
    "Etc/GMT+10",
    "Etc/GMT+11",
    "Etc/GMT+12",
    "Etc/GMT-1",
    "Etc/GMT-2",
    "Etc/GMT-3",
    "Etc/GMT-4",
    "Etc/GMT-5",
    "Etc/GMT-6",
    "Etc/GMT-7",
    "Etc/GMT-8",
    "Etc/GMT-9",
    "Etc/GMT-10",
    "Etc/GMT-11",
    "Etc/GMT-12",
    "Etc/GMT-13",
    "Etc/GMT-14",
    "Etc/GMT",
    "Etc/GMT+0",
    "Etc/GMT-0",
    "Etc/GMT0",
    "Etc/Greenwich",
    "Etc/UCT",
    "Etc/UTC",
    "Etc/Universal",
    "Etc/Zulu"
]
$default: CET

hardware_clock_set_to_utc:

```

```
$type: boolean
$default: True
...
```

`form.yml` contains additional information that describes how the form for a pillar should look for Uyuni. This information is contained in attributes that always start with a `$` sign.

Important: Ignored Values

All values that start with a `$` sign are annotations used to display the UI that users interact with. These annotations are not part of pillar data itself and are handled as metadata.

The following are valid attributes.

`$type`

The most important attribute is the `$type` attribute. It defines the type of the pillar value and the form-field that is generated. The following represent the supported types:

- `text`
- `password`
- `number`
- `url`
- `email`
- `date`
- `time`
- `datetime`
- `boolean`
- `color`
- `select`
- `group`
- `edit-group`
- `namespace`
- `hidden-group` (obsolete, renamed to `namespace`)



Note: Text Attribute

The text attribute is the default and does not need to be specified explicitly.

Many of these values are self-explanatory:

- The text type generates a simple text field
- The password type generates a password field
- The color type generates a color picker.

The group, edit-group, and namespace (formerly hidden-group) types do not generate an editable field and are used to structure form and pillar data. The difference between group and namespace is group generates a visible border with a heading, and namespace shows nothing visually (and is only used to structure pillar data). The difference between group and edit-group is edit-group allows to structure and restrict editable fields in a more flexible way. For example, edit-group supports the \$minItems and \$maxItems attributes, and thus it simplifies complex and repeatable input structures. For an edit-group example, see [Section 4.4.1, “Simple edit-group Example”](#).

\$default

\$default allows you to specify a default value that is displayed and used, if no other value is entered. In an edit-group it allows to create initial members of the group and populate them with specified data.

\$name

\$name allows you to specify the name of a value that is shown in the form. If this value is not set, the pillar name is used and capitalized without underscores and dashes. You reference it in the same section with \${name}.

\$help and \$placeholder

The \$help and \$placeholder attributes are used to give a user a better understanding of what the value should be.

- \$help defines the message a user sees when hovering over a field
- \$placeholder displays a gray placeholder text in the field.

\$placeholder may only be used with text fields like text, password, email or date. It does not make sense to add a placeholder if you also use \$default as this will hide the placeholder.

\$minItems and \$maxItems

In an edit-group, \$minItems and \$maxItems allow you to specify the lowest and highest number the group can occur.

\$itemName

In an edit-group, \$itemName allows to define a template for the name to be used for the members of the group.

\$prototype

In an edit-group, \$prototype allows to define default (or pre-filled) values for newly added members in the group.

\$scope

\$scope allows you to specify a hierarchy level at which a value may be edited. Possible values are system, group and readonly.

The default \$scope: system allows values to be edited at group and system levels. A value can be entered for each system but if no value is entered the system will fall back to the group default.

If using \$scope: group, a value may only be edited for a group. On the system level you will be able to see the value, but not edit it.

The \$scope: readonly option makes a field read-only. It can be used to show a user data which should be known, but should not be editable. This option only makes sense in combination with the \$default attribute.

\$visibleIf

\$visibleIf allows you to show a field or group if a simple condition is met. A condition always looks similar to the following example:

```
some_group#another_group#my_checkbox == true
```

The left part of the above statement is the path to another value, and groups are separated by \$ signs. The middle section of the command should be either == for a value to be equal or != for values that should be not equal. The last field in the statement can be any value which a field should have or not have.

The field with this attribute associated with it will now be shown only when the condition is met. In this example the field will be shown only if my_checkbox is checked. The ability to use conditional statements is not limited to check boxes. It may also be used to check values of select-fields, text-fields, etc.

A check box should be structured like the following example:

```
some_group:
  $type: group

another_group:
  $type: group

  my_checkbox:
    $type: boolean
```

Relative paths can be specified using prefix dots. One dot means sibling, 2 dots mean parent, etc. This is mostly useful for edit-group.

```
some_group:
  $type: group

another_group:
  $type: group

  my_checkbox:
    $type: boolean

  my_text:
    $visibleIf: .my_checkbox

yet_another_group:
  $type: group

  my_text2:
    $visibleIf: ..another_group#my_checkbox
```

By using multiple groups with the attribute, you can allow a user to select an option and show a completely different form, dependent upon the selected value.

Values from hidden fields may be merged into the pillar data and sent to the minion. A formula must check the condition again and use the appropriate data. For example:

```
show_option:
  $type: checkbox
some_text:
  $visibleIf: show_option == true
```

```
{% if pillar.show_option %}
do_something:
  with: {{ pillar.some_text }}
```



```
{% endif %}
```

\$values

\$values can only be used together with \$type: select to specify the different options in the select-field. \$values must be a list of possible values to select. For example:

```
select_something:
  $type: select
  $values: ["option1", "option2"]
```

Or alternatively:

```
select_something:
  $type: select
  $values:
    - option1
    - option2
```

4.4.1 Simple edit-group Example

See the following **edit-group** example:

```
partitions:
  $name: "Hard Disk Partitions"
  $type: "edit-group"
  $minItems: 1
  $maxItems: 4
  $itemName: "Partition ${name}"
  $prototype:
    name:
      $default: "New partition"
    mountpoint:
      $default: "/var"
    size:
      $type: "number"
      $name: "Size in GB"
  $default:
    - name: "Boot"
      mountpoint: "/boot"
    - name: "Root"
      mountpoint: "/"
      size: 5000
```

After clicking *Add* for one time you will see *Figure 4.1, "edit-group Example in the Web UI"* filled with the default values. The formula itself is called hd-partitions and will appear as Hd Partitions in the Web UI.

This is a feature preview: On this page you can configure [Salt formulas](#) to automatically install and configure software. We would be glad to receive your feedback.

← Prev Next →

Hd Partitions

Hard Disk Partitions

Partition Boot

Name:

Mountpoint:

Size in GB:

Partition Root

Name:

Mountpoint:

Size in GB:

Partition New partition

Name:

Mountpoint:

Size in GB:

FIGURE 4.1: edit-group EXAMPLE IN THE WEB UI

+ Add Item

To remove definition of a partition click the minus symbol in the title line of an inner group. When form fields are properly filled confirm with clicking *Save Formula* in the upper right corner of the formula.

4.5 Writing Salt Formulas

Salt formulas are pre-written Salt states, which may be configured with pillar data. You can parametrize state files using Jinja. Jinja allows you to access pillar data by using the following syntax. (This syntax works best when your uncertain a pillar value exists as it will throw an error):

```
pillar.some.value
```

When you are sure a pillar exists may also use the following syntax:

```
salt['pillar.get']('some:value', 'default value')
```

You may also replace the pillar value with grains (for example, grains.some.value) allowing access to grains.

Using data this way allows you to make a formula configurable. The following code snippet will install a package specified in the pillar package_name. For example:

```
install_a_package:
  pkg.installed:
    - name: {{ pillar.package_name }}
```

You may also use more complex constructs such as if/else and for-loops to provide greater functionality. For Example:

```
{% if pillar.installSomething %}
something:
  pkg.installed
{% else %}
anotherPackage:
  pkg.installed
{% endif %}
```

Another example:

```
{% for service in pillar.services %}
```

```
start_{{ service }}:
  service.running:
    - name: {{ service }}
{% endfor %}
```

Jinja also provides other helpful functions. For example, you can iterate over a dictionary:

```
{% for key, value in some_dictionary.items() %}
do_something_with_{{ key }}: {{ value }}
{% endfor %}
```

You may want to have Salt manage your files (for example, configuration files for a program), and you can change these with pillar data. For example, the following snippet shows how you can manage a file using Salt:

```
/etc/my_program/my_program.conf:
  file.managed:
    - source: salt://my_state/files/my_program.conf
    - template: jinja
```

Salt will copy the file `salt-file_roots/my_state/files/my_program.conf` on the salt master to `/etc/my_program/my_program.conf` on the minion and template it with Jinja. This allows you to use Jinja in the file, exactly like shown above for states:

```
some_config_option = {{ pillar.config_option_a }}
```

4.6 Separating Data

It is often a good idea to separate data from a state to increase its flexibility and add re-usability value. This is often done by writing values into a separate file named `map.jinja`. This file should be placed within the same directory as your state files.

The following example will set `data` to a dictionary with different values, depending on which system the state runs on. It will also merge data with the pillar using the `some.pillar.data` value so you can access `some.pillar.data.value` by just using `data.value`.

You can also choose to override defined values from pillars (for example, by overriding `some.pillar.data.package` in the example).

```
{% set data = salt['grains.filter_by']({
  'Suse': {
    'package': 'packageA',
```

```

        'service': 'serviceA'
    },
    'RedHat': {
        'package': 'package_a',
        'service': 'service_a'
    }
}, merge=salt['pillar.get']('some:pillar:data')) %}

```

After creating a map file like the above example, you can easily maintain compatibility with multiple system types while accessing "deep" pillar data in a simpler way. Now you can import and use data in any file. For example:

```

{% from "some_folder/map.jinja" import data with context %}

install_package_a:
  pkg.installed:
    - name: {{ data.package }}

```

You can also define multiple variables by copying the {% set ...%} statement with different values and then merge it with other pillars. For example:

```

{% set server = salt['grains.filter_by']({
    'Suse': {
        'package': 'my-server-pkg'
    }
}, merge=salt['pillar.get']('myFormula:server')) %}
{% set client = salt['grains.filter_by']({
    'Suse': {
        'package': 'my-client-pkg'
    }
}, merge=salt['pillar.get']('myFormula:client')) %}

```

To import multiple variables, separate them with a comma. For Example:

```

{% from "map.jinja" import server, client with context %}

```

Formulas utilized with Uyuni should follow formula conventions listed in the official documentation:

- <https://docs.saltstack.com/en/latest/topics/development/conventions/formulas.html> 

4.7 Uyuni Generated Pillar Data

When pillar data is generated (for example, after applying the highstate) the following external pillar script generates pillar data for packages, group ids, etc. and includes all pillar data for a system:

```
/usr/share/susemanager/modules/pillar/suma_minion.py
```

The process is executed as follows:

1. The `suma_minion.py` script starts and finds all formulas for a system (by checking the `group_formulas.json` and `server_formulas.json` files).
2. `suma_minion.py` loads the values for each formula (groups and from the system) and merges them with the highstate (default: if no values are found, a group overrides a system if `$scope: group` etc.).
3. `suma_minion.py` also includes a list of formulas applied to the system in a pillar named `formulas`. This structure makes it possible to include states. The top file (in this case specifically generated by the `mgr_master_tops.py` script) includes a state called `formulas` for each system. This includes the `formulas.sls` file located in:

```
/usr/share/susemanager/formulas/states/
```

The content looks similar to the following:

```
include: {{ pillar["formulas"] }}
```

This pillar includes all formulas, that are specified in pillar data generated from the external pillar script.

4.8 Formula Requirements

Formulas should be designed/created directly after a Uyuni installation, but if you encounter any issues check the following:

- The external pillar script (`suma_minion.py`) must include formula data.
- Data is saved to `/srv/susemanager/formula_data` and the `pillar` and `group_pillar` sub-directories. These should be automatically generated by the server.
- Formulas must be included for every minion listed in the top file. Currently this process is initiated by the `mgr_master_tops.py` script which includes the `formulas.sls` file located in:

```
/usr/share/susemanager/formulas/states/
```

This directory must be a salt file root. File roots are configured on the salt-master (Uyuni) located in:

```
/etc/salt/master.d/susemanager.conf
```


4.9 Using Salt Formulas with Uyuni

The following procedure provides an overview on using Salt Formulas with Uyuni.

1. Official formulas may be installed as RPMs. Place the custom states within `/srv/salt/your-formula-name/` and the metadata (`form.yml` and `metadata.yml`) in `/srv/formula_metadata/your-formula-name/`. After installing your formulas they will appear in *Salt > Formula Catalog*.
2. To begin using a formula, apply it to a group or system. Apply a formula to a group or system by selecting the *System Details > Formulas* tab of a *System Details* page or *System Group*. From the *System Details > Formulas* page you can select any formulas you wish to apply to a group or system. Click the *Save* button to save your changes to the database.
3. After applying one or more formulas to a group or system, additional tabs will become available from the top menu, one for each formula selected. From these tabs you may configure your formulas.

4. When you have finished customizing your formula values you will need to apply the highstate for them to take effect. Applying the highstate will execute the state associated with the formula and configure targeted systems. You can use the *Apply Highstate* button from any formulas page of a group.
5. When a change to any of your values is required or you need to re-apply the formula state because of a failure or bug, change values located on your formula pages and re-apply the highstate. Salt will ensure that only modified values are adjusted and restart or reinstall services only when necessary.

This conclude your introduction to Salt Formulas. For additional information, see:

- <https://docs.saltstack.com/en/latest/topics/development/conventions/formulas.html> 
- < < < < < < HEAD

= SUSE Manager for Retail Salt Formulas

This section provides an introduction to Salt Formulas shipped with SUSE Manager for Retail. These formulas such as the PXE boot, branch server network, or saltboot formulas are used to fine-tune the SUSE Manager for Retail infrastructure.

= = Pxe Formula The Pxe formula ([pxe-formula](#)) for installing, setting up, and uninstalling syslinux PXE boot on the POS branchserver.

This is a feature preview: On this page you can configure [Salt formulas](#) to automatically install and configure your feedback via the [forum](#).

[← Prev](#)
[Next →](#)

Pxe

Kernel filename:

Initrd filename:

Kernel commandline parameters:

PXE root directory:

Branch id:

FIGURE 4.2: PXE FORMULA

= = Branch Network Formula

The Branch Network formula ([branch-network-formula](#)) for configuring the branch server network.

This is a feature preview: On this page you can configure [Salt formulas](#) to automatically install and configure network settings. Please provide feedback via the [forum](#).

[← Prev](#)
[Next →](#)

Branch Network

Dedicated Nic:



Nic:

eth1

Ip:

192.168.128.1

Netmask:

255.255.255.0

Enable Route:



Enable Nat:



Forwarder:

bind

Forwarder Fallback:



server directory:

/srv/saltboot

server directory user:

saltboot

server directory group:

saltboot

Using Salt Formulas with Uyuni

FIGURE 4.3: BRANCH NETWORK FORMULA

= = Saltboot Formula

The Saltboot formula (saltboot-formula) is a formula for configuring a boot image of a POS terminal.

Details

Software

Configuration

Provisioning

Groups

Audit

States

Overview

Properties

Remote Command

Hardware

Migrate

Notes

Formulas

Branch Network

Dhcpd

Pxe

Saltboot

Tftpd

Vsftpd

This is a feature preview: On this page you can configure [Salt formulas](#) to automatically install and configure Saltboot. Please provide your feedback via the [forum](#).

← Prev

Next →

Saltboot

Disk partitioning

Disk 1

Disk symbolic ID:

disk

Device type:

DISK

Disk device (valid only for DISK devices):

RAID level:

Symbolic IDs of devices to used for RAID device type

+ Add Item

Disk Label:

gpt

FIGURE 4.4: SALTBOOT FORMULA

Then you configure one or more partitions:

The screenshot shows a web-based configuration interface for Saltboot Formula Partitions. The main container is titled "Partitions" and contains a list of partition configurations. The first partition, "Partition 1", is expanded, showing the following fields:

- Partition symbolic ID:** A text input field containing the letter "p".
- Partition size (MiB):** An empty text input field.
- Device mount point:** An empty text input field.
- Filesystem format:** An empty text input field.
- OS Image to deploy:** An empty text input field.
- Image Version:** An empty text input field.
- Partition encryption password:** An empty text input field.
- Partition flags:** An empty text input field.

Below the list of partitions, there is a button labeled "+ Add Item". At the bottom of the interface, there is another button labeled "+ Add Item".

FIGURE 4.5: SALTBOOT FORMULA PARTITIONS

= = Image Sync Formula

The Image Sync formula ([image-sync-formula](#)) is a formula for syncing images to a Branch Server.

For now, there is nothing configurable and it is not part of the highstate. This means it is not visible in the Web UI. Apply it from the command-line or via cron as follows (replace *<branch-server>* with the name of your branch server):

```
salt <branchserver> state.apply image-sync
```



= = Salt Formulas Coming with SUSE Manager

For general information, see the [Salt Formulas installation and usage instructions](https://docs.saltstack.com/en/latest/topics/development/conventions/formulas.html) at <https://docs.saltstack.com/en/latest/topics/development/conventions/formulas.html> .

= = = Domain Name System (Bind)

With the bind formula you set up and configure a Domain Name System (DNS) server. For technical information about the bind formula and low-level pillar data, see the [README.rst](#) file on the Uyuni server: [/usr/share/susemanager/formulas/metadata/bind/README.rst](#).

DNS is needed to resolve the domain names and host names into IP addresses. For more information about DNS, see the SLES Administration Guide, Services, The Domain Name System.

 d186.suse.de 

Details

Software

Configuration

Provisioning

Groups

Audit

States

Formulas

Bind

This is a feature preview: On this page you can configure [Salt formulas](#) to automatically install and configure [bind](#).
feedback via the [forum](#).

← Prev

Next →

Bind

Config

Options

+ Add Item

Include Forwarders:

☐

Configured Zones

Zone 1

Name:

FIGURE 4.6: BIND FORMULA

In the Config group you can set arbitrary options such as directory where are the zone data files (usually /var/lib/named/) or forwarders. Click *Add Item* to provide more Key/Value fields for configuration.

Check Include Forwarders if you want to rely on an external DNS server if your DNS is down (or is otherwise not able to resolve an address).

At least, you will configure one zone. In Configured Zones define your zone; for example, example.com. Then in Available Zones configure this zone: as Name enter your zone (in this case example.com) and the File to which this configuration should be written (example.com.txt). Enter the mandatory SOA record (start of authority), and the A, NS, and CNAME Records you need.

On the other hand, if no records entry exists, the zone file is not generated by this state rather than taken from salt://zones. For how to overwrite this URL, see pillar.example.

← Prev

Next →

Configured Zones

Zone 1

Name:

Type:

master

Notify:

☐

+ Add Item

Available Zones

Zone 1

Name:

File:

SOA

NS:

ns@zone

Contact:

admin@domain

FIGURE 4.7: BIND-02-ZONES

← Prev

Next →

Available Zones

Zone 1

Name:

File:

SOA

NS:

ns@zone

Contact:

admin@domain

Serial:

auto

Class:

IN

Refresh:

8600

Retry:

900

Expiry:

86000

NXDOMAIN:

500

TTL:

8600

FIGURE 4.8: BIND-03-RECORDS

← Prev

Next →

Records

A

+ Add Item

NS

@

+ Add Item

CNAME

+ Add Item

Generate Reverse

Network:

For Zones

FIGURE 4.9: BIND-03-RECORDS2

In Generate Reverse, and define reverse mapping and for which zones:

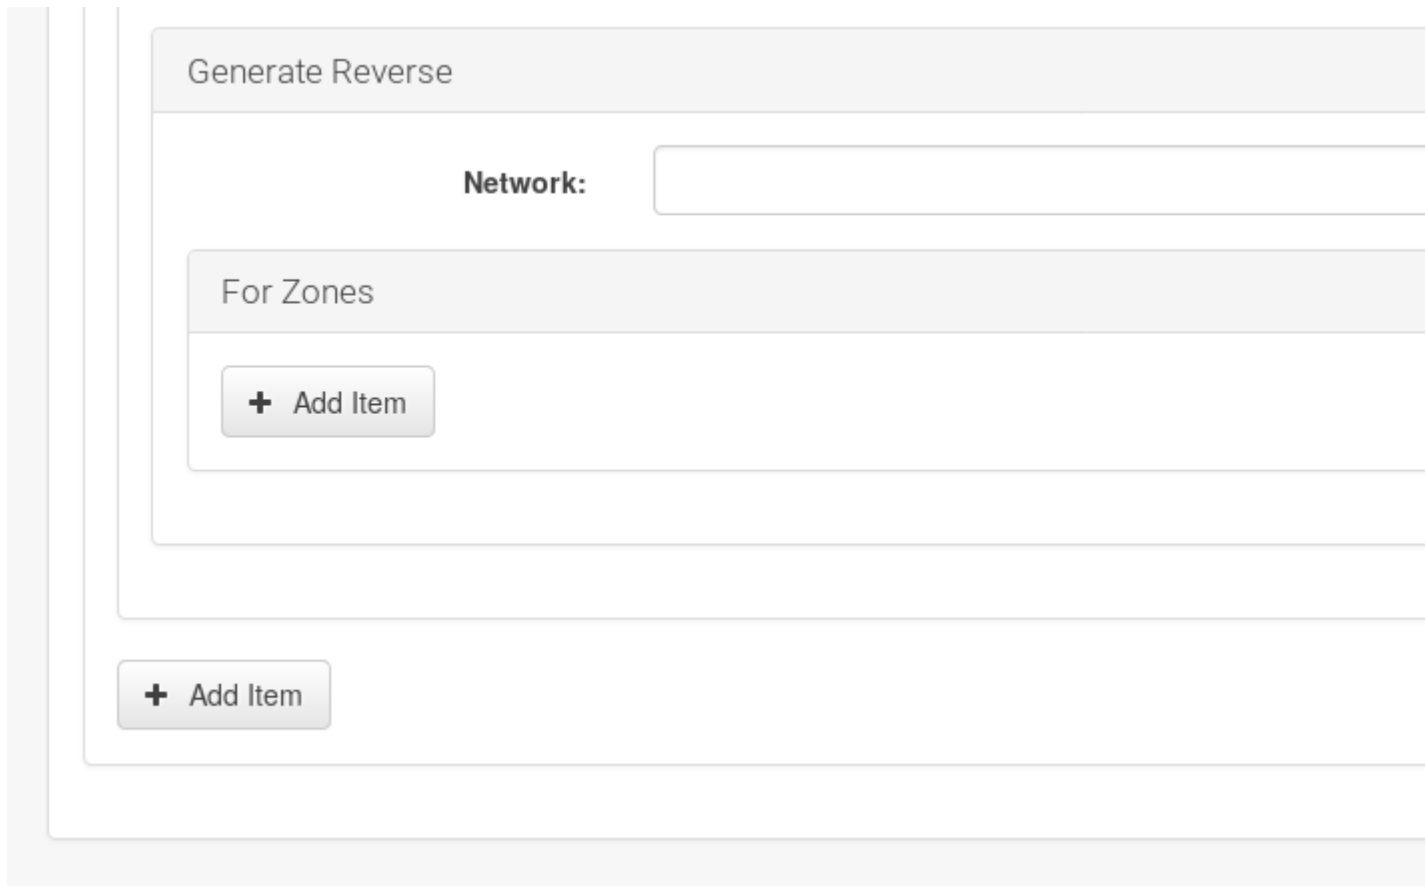


FIGURE 4.10: **BIND-04-REVERSE**

When saved, data is written to `/srv/susemanager/formula_data/pillar/<salt-minion.example.com>_bind.json`

If you apply the highstate (*System Details > States > Highstate*), it first ensures that `bind` and all required packages will get installed. Then it will start the DNS service (`named`).

=== Dhcpcd

With the `dhcpcd` formula you set up and configure a DHCP server (Dynamic Host Configuration Protocol). For technical information about the `dhcpcd` formula and low-level pillar data, see the Pillar example file `/usr/share/susemanager/formulas/metadata/dhcpcd/pillar.example`.

DHCP is needed to define network settings centrally (on a server) and let clients retrieve and use this information for local host configuration. For more information about DHCP, see the SLES Administration Guide, Services, DHCP.

This is a feature preview: On this page you can configure [Salt formulas](#) to automatically install and configure your feedback via the [forum](#).

← Prev Next →

Dhcpd

Domain Name:

example.org

Domain Name Servers

example.org

– Remove

+ Add Item

Listen Interfaces

eth1

– Remove

+ Add Item

Authoritative:



FIGURE 4.11: DHCPD FORMULA

Domain Name.

Domain Name Servers. One or more Domain Name Service (DNS) servers.

On which interface(s) the DHCP server should listen (Listen interfaces). Set option for this interface: Authoritative: Max Lease Time: Default Lease Time:

Next is at least one network in the Network configuration (subnet) group (with IP address, netmask, etc.). You define every network with Dynamic IP range, Routers, and Hosts with static IP addresses (with defaults from subnet) (optionally).


And finally Hosts with static IP addresses (with global defaults).

If you apply the highstate (*System Details > States > Highstate*), it first ensures that dhcp-server and all required packages will get installed. Then it will start the DHCP service (dhcpd).

= = = Tftpd

With the tftpd formula you set up and configure a TFTP server (Trivial File Transfer Protocol). A TFTP server is a component that provides infrastructure for booting with PXE.

For more information about setting up TFTP, see the SLES Deployment Guide, Preparing Network Boot Environment, Setting Up a TFTP Server.

 d186.suse.de 

Details

Software

Configuration

Provisioning

Groups

Audit

States

Formulas

Dhcpd

Tftpd

Vsftpd

Bind

This is a feature preview: On this page you can configure [Salt formulas](#) to automatically install and configure your feedback via the [forum](#).

← Prev

Next →

Tftpd

Internal Network Address:

TFTP base directory:

/srv/tftpboot

run TFTP under user:

tftp

FIGURE 4.12: TFTP FORMULA

For setting up a TFTP server, specify the Internal Network Address, TFTP base directory (default: /srv/tftpboot), and run TFTP under user (default: sftp).

If you apply the highstate (*System Details > States > Highstate*), it first ensures that atftp and all required packages will get installed. Then it will start TFTP (atftpd).

== Vsftpd

With the vsftpd formula you set up and configure Vsftpd. Vsftpd is an FTP server or daemon, written with security in mind. "vs" in its name stands for "Very Secure".

This is a feature preview: On this page you can configure [Salt formulas](#) to automatically install and configure your feedback via the [forum](#).

← Prev Next →

Vsftpd

FTP server directory:

Internal Network Address:

Enable ssl:

☐

Chroot dir:

Allow anonymous FTP:

☒

Allow SSL for anonymous:

☒

Run standalone:

☒

Allow local users:

☒

Activate directory messages:

☒

Use localtime:

☒

FIGURE 4.13: VSFTPD FORMULA

For configuring a VSFTP server, specify the settings and options in the Vsftpd formula. There are settings such as FTP server directory, Internal Network Address Enable ssl, etc.

If you apply the highstate (*System Details* > *States* > *Highstate*), it first ensures that vsftpd and all required packages will get installed. Then it will start the VSFTP service (vsftpd).

For more information about setting up and tuning Vsftpd, see the documentation coming with the vsftpd package (/usr/share/doc/packages/vsftpd/ when the package is installed).

<<<<<<< HEAD >>>>>>> eadc4168... Removed entities from top of files, added a single include, added missing entities for retail

>>>>>>> eadc4168... Removed entities from top of files, added a single include, added missing entities for retail

5 Configuration Management with Salt

5.1 Configuration Management Overview

Salt is capable of applying states by matching minions with relevant state data. This data comes from Uyuni in the form of package and custom states.

5.2 State Data: Levels of Hierarchy

State data comes from Uyuni in the form of package and custom states and targets minions at three specific levels of hierarchy. The state hierarchy is defined by the following order or priority: individual minions have priority on packages and custom states over groups; next a group has priority over the organization.

- Minion Level
Systems > Specific Minion > States
- Group Level
Systems > System Groups
- Organization Level
Systems > Manage System Types: > My Organization

For example:

- Org1 requires that vim version 1 is installed
- Group1 requires that vim version 2 is installed
- Group2 requires any version installed

This would lead to the following order of hierarchy:

- Minion1 part of [Org1, Group1] wants vim removed, vim is removed (Minion Level)
- Minion2 part of [Org1, Group1] wants vim version 2 gets version 2 (Group Level)
- Minion3 part of [Org1, Group1] wants any version, gets version 2 (Org Level)
- Minion4 part of [Org1, Group2] wants any version, gets vim version 1 (Org Level)

5.3 Salt States Storage Locations

The Uyuni salt-master reads its state data from three file root locations.

The directory `/usr/share/susemanager/salt` is used by Uyuni and comes from the `susemanager-sls`. It is shipped and updated together with Uyuni and includes certificate setup and common state logic to be applied to packages and channels.

The directory `/srv/susemanager/salt` is generated by Uyuni and based on assigned channels and packages for minions, groups and organizations. This file will be overwritten and regenerated. This could be thought of as the Uyuni database translated into salt directives.

The third directory `/srv/salt` is for custom state data, modules etc. Uyuni does not operate within or utilize this directory. However the state data placed here affects the Highstate of minions and is merged with the total state result generated by Uyuni.

5.4 Uyuni States

All sls files created by users will be saved to disk on the salt-master server. These files will be placed in `/srv/susemanager/salt/` and each organization will be placed within its own directory. Although these states are custom, these states are created using Uyuni. The following provides an overview of directory structure:

```
|— manager_org_DEVEL
|   |— files
|   |   ... files needed by states (uploaded by users)...
|   |— state.sls
|       ... other sls files (created by users)...
E.g.:
|— manager_org_TESTING
|   |— files
|   |   |— motd      # user created
|   |   ... other files needed by states ...
|   |— motd.sls     # user created
|       ... other sls files ...
```

5.5 Pillar Data

SUSE Manager exposes a small amount of internal data as Pillars which can be used with custom SUSE Linux Enterprise Server states. Data that is exposed includes group membership, organization membership, and file roots. These are managed either automatically by Uyuni, or manually by the user.

To avoid hard-coding organization IDs within SUSE Linux Enterprise Server files, a pillar entry is added for each organization:

```
org-files-dir: relative_path_to_files
```

The specified file is available for all minions which belong to the organization.

This is an example of a Pillar located at /etc/motd:

```
file.managed:
  - source: salt://{{ pillar['org-files-dir'] }}/motd
  - user: root
  - group: root
  - mode: 644
```

5.6 Group States

Pillar data can be used to perform bulk actions, like applying all assigned states to minions within the group. This section contains some example of bulk actions that you can take using group states.

In order to perform these actions, you will need to determine the ID of the group that you want to manipulate. You can determine the Group ID by using the spacecmd command:

```
spacecmd group_details
```

In these examples we will use an example Group ID of GID.

To apply all states assigned to the group:

```
salt -I 'group_ids:GID' state.apply custom.group_GID
```

To apply any state (whether or not it is assigned to the group):

```
salt -I 'group_ids:GID' state.apply ``state``
```

To apply a custom state:

```
salt -I 'group_ids:2130' state.apply manager_org_1.`customstate`
```

Apply the highstate to all minions in the group:

```
salt -I 'group_ids:GID' state.apply
```

6 Salt Minion Scalability

6.1 Salt Minion Onboarding Rate

The rate at which SUSE Manager can on-board minions (accept Salt keys) is limited and depends on hardware resources. On-boarding minions at a faster rate than SUSE Manager is configured for will build up a backlog of unprocessed keys slowing the process and potentially exhausting resources. It is recommended to limit the acceptance key rate programmatically. A safe starting point would be to on-board a minion every 15 seconds, which can be implemented via the following command:

```
for k in $(salt-key -l un|grep -v Unaccepted); do salt-key -y -a $k; sleep 15; done
```

6.2 Minions Running with Unaccepted Salt Keys

Minions which have not been on-boarded, (minions running with unaccepted Salt keys) consume resources, in particular inbound network bandwidth for ~ 2.5 Kb/s per minion. 1000 idle minions will consume around ~ 2.5 Mb/s, and this number will drop to almost 0 once on-boarding has been completed. Limit non-onboarded systems for optimal performance.

Salt's official documentation suggests the maximum number of opened files should be set to at least $2 \times$ the minion count. Current default is 16384, which is sufficient for ~ 8000 minions. For larger installations, this number may be increased by editing the following line in /usr/lib/systemd/system/salt-master.service :

```
LimitNOFILE=16384
```

6.3 Salt Timeouts

6.3.1 Background Information

Salt features two timeout parameters called `timeout` and `gather_job_timeout` that are relevant during the execution of Salt commands and jobs—it does not matter whether they are triggered using the command line interface or API. These two parameters are explained in the following article.

This is a normal workflow when all minions are well reachable:

- A salt command or job is executed:

```
salt '*' test.ping
```

- Salt master publishes the job with the targeted minions into the Salt PUB channel.
- Minions take that job and start working on it.
- Salt master is looking at the Salt RET channel to gather responses from the minions.
- If Salt master gets all responses from targeted minions, then everything is completed and Salt master will return a response containing all the minion responses.

If some of the minions are down during this process, the workflow continues as follows:

1. If `timeout` is reached before getting all expected responses from the minions, then Salt master would trigger an additional job (a Salt `find_job` job) targeting only pending minions to check whether the job is already running on the minion.
2. Now `gather_job_timeout` is evaluated. A new counter is now triggered.
3. If this new `find_job` job responses that the original job is actually running on the minion, then Salt master will wait for that minion's response.
4. In case of reaching `gather_job_timeout` without having any response from the minion (neither for the initial `test.ping` nor for the `find_job` job), Salt master will return with only the gathered responses from the responding minions.

By default, Uyuni globally sets `timeout` and `gather_job_timeout` to 120 seconds. So, in the worst case, a Salt call targeting unreachable minions will end up *with 240 seconds of waiting* until getting a response.

6.3.2 A Presence Ping Mechanism for Unreachable Salt Minions

In order to prevent waiting until timeouts are reached when some minions are down, SUSE introduced a so-called "presence mechanism" for Salt minions.

This presence mechanism checks for unreachable Salt minions when Uyuni is performing synchronous calls to these minions, and it excludes unreachable minions from that call. Synchronous calls are going to be displaced in favor of asynchronous calls but currently still being used during some workflows.

The presence mechanism triggers a Salt `test.ping` with a custom and fixed short Salt timeout values. Default Salt values for the presence ping are: `timeout = 4` and `gather_job_timeout = 1`. This way, we can quickly detect which targeted minions are unreachable, and then exclude them from the synchronous call.

6.3.3 Overriding Salt Presence Timeout Values

Uyuni administrators can increase or decrease default presence ping timeout values by removing the comment markers (`\#`) and setting the desired values for `salt_presence_ping_timeout` and `salt_presence_ping_gather_job_timeout` options in `/etc/rhn/rhn.conf`:

```
# SUSE Manager presence timeouts for Salt minions
# salt_presence_ping_timeout = 4
# salt_presence_ping_gather_job_timeout = 1
```

6.3.4 Salt SSH Minions (SSH Push)

Salt SSH minions are slightly different than regular minions (zeromq). Salt SSH minions do not use Salt PUB/RET channels but a wrapper Salt command inside of an SSH call. Salt `timeout` and `gather_job_timeout` are not playing a role here.

Uyuni defines a timeout for SSH connections in `/etc/rhn/rhn.conf`:

```
# salt_ssh_connect_timeout = 180
```

The presence ping mechanism is also working with SSH minions. In this case, Uyuni will use `salt_presence_ping_timeout` to override the default timeout value for SSH connections.

7 Activation Key Management

7.1 What are Activation Keys?

An activation key in Uyuni is a group of configuration settings with a label. You can apply all configuration settings associated with an activation key by adding its label as a parameter to a bootstrap script. Under normal operating conditions best practices suggest using an activation key label in combination with a bootstrap script.

An activation key can specify:

- Channel Assignment
- System Types (Traditionally called Add-on Entitlements)
- Contact Method
- Configuration Files
- Packages to be Installed
- System Group Assignment

Activation keys are just a collection of configuration settings which have been given a label name and then added to a bootstrap script. When the bootstrap script is executed all configuration settings associated with the label are applied to the system the script is run on.

7.2 Provisioning and Configuration




FIGURE 7.1: PROVISIONING AND CONFIGURATION OVERVIEW

7.3 Activation Keys Best Practices

There are a few important concepts which should be kept in mind when creating activation keys. The following sections provide insight when creating and naming your activation keys.

7.3.1 Key Label Naming

One of the most important things to consider during activation key creation is label naming. Creating names which are associated with your organization's infrastructure will make it easier for you when performing more complex operations. When naming key labels keep the following in mind:

- OS naming (mandatory): Keys should always refer to the OS they provide settings for
- Architecture naming (recommended): Unless your company is running on one architecture only, for example x86_64, then providing labels with an architecture type is a good idea.
- Server type naming: What is, or what will this server be used for?
- Location naming: Where is the server located? Room, building, or department?
- Date naming: Maintenance windows, quarter, etc.
- Custom naming: What naming scheme suits your organizations needs?

Example activation key label names:

```
sles12-sp2-web_server-room_129-x86_64
```

```
sles12-sp2-test_packages-blg_502-room_21-ppc64le
```

7.3.2 Channels which will be Included

When creating activation keys you also need to keep in mind which channels (software sources) will be associated with it.



Important: Default Base Channel

Keys should have a specific base channel assigned to it, for example SLES12-SP2-Pool-x86_64. If this is not the case Uyuni cannot use specific stages. Using the default base channel is not recommended and may cause problems.

- Channels to be included:
 - suse-manager-tools
- Typical packages to be included:
 - osad (pushing tasks)
 - Installs python-jabberpy and pyxml as dependencies
 - rhncfg-actions (Remote Command, Configuration Management)
 - Installs rhncfg and rhncfg-client as dependencies

7.4 Combining Activation Keys

You can combine activation keys when executing the bootstrap script on your clients. Combining keys allows for more control on what is installed on your systems and reduces duplication of keys for large complex environments.




FIGURE 7.2: COMBINING ACTIVATION KEYS

Combining Activation Keys

Server OS / Stage Key

Base Channels: sles12sp0_3prod-sles12-pool-x86_64 (01.06.2015)

Any system registered using this activation key will be subscribed to the selected child channels.

The following child channels of **sles12sp0_3prod-sles12-pool-x86_64 (01.06.2015)** can be associated with this activation key.

- sles12sp0_3prod-obs-home-packages-x86_64
- sles12sp0_3prod-obs-server-packages-x86_64
- sles12sp0_3prod-sle-12-ga-desktop-amd-driver-x86_64-we
- sles12sp0_3prod-sle-12-ga-desktop-midia-driver-x86_64-we
- sles12sp0_3prod-sle-ha12-pool-x86_64
- sles12sp0_3prod-sle-ha12-updates-x86_64
- sles12sp0_3prod-sle-manager-tools12-pool-x86_64**
- sles12sp0_3prod-sle-manager-tools12-updates-x86_64**
- sles12sp0_3prod-sle-module-adv-systems-management12-pool-x86_64
- sles12sp0_3prod-sle-module-adv-systems-management12-updates-x86_64
- sles12sp0_3prod-sle-module-legacy12-pool-x86_64
- sles12sp0_3prod-sle-module-legacy12-updates-x86_64
- sles12sp0_3prod-sle-module-public-cloud12-pool-x86_64
- sles12sp0_3prod-sle-module-public-cloud12-updates-x86_64
- sles12sp0_3prod-sle-module-web-scripting12-pool-x86_64
- sles12sp0_3prod-sle-module-web-scripting12-updates-x86_64
- sles12sp0_3prod-sles12-updates-x86_64**
- sles12sp0_3prod-sle-sdk12-pool-x86_64
- sles12sp0_3prod-sle-sdk12-updates-x86_64
- sles12sp0_3prod-sle-we12-pool-x86_64

Update Key

Any other type of key

Base Channels: SUSE Manager Default

Any system registered using this activation key will be subscribed to the selected child channels.

- sles12sp0_3prod-sle-module-legacy12-pool-x86_64
- sles12sp0_3prod-sle-module-legacy12-updates-x86_64
- sles12sp0_3prod-sle-module-public-cloud12-pool-x86_64
- sles12sp0_3prod-sle-module-public-cloud12-updates-x86_64
- sles12sp0_3prod-sle-module-web-scripting12-pool-x86_64**
- sles12sp0_3prod-sle-module-web-scripting12-updates-x86_64**
- sles12sp0_3prod-sles12-updates-x86_64
- sles12sp0_3prod-sle-sdk12-pool-x86_64**
- sles12sp0_3prod-sle-sdk12-updates-x86_64**
- sles12sp0_3prod-sle-we12-updates-x86_64
- SUSE-Manager-Proxy-1.7-Pool for x86_64
- SLES11-SP1-Pool for x86_64 Proxy 1.7
- SLES11-SP1-Updates for x86_64 Proxy 1.7
- SLES11-SP2-Core for x86_64 Proxy 1.7
- SLES11-SP2-Updates for x86_64 Proxy 1.7
- SUSE-Manager-Proxy-1.7-Updates for x86_64
- SUSE-Manager-Proxy-2.1-Pool for x86_64
- SLES11-SP3-Pool for x86_64 Proxy 2.1
- SLES11-SP3-Updates for x86_64 Proxy 2.1
- SUSE-Manager-Proxy-2.1-Updates for x86_64

Update Key

FIGURE 7.3: COMBINING ACTIVATION KEYS 2

7.5 Using Activation Keys and Bootstrap with Traditional Clients (Non-Salt)

Create the initial bootstrap script template from the command line on the Uyuni server with:

```
# mgr-bootstrap
```

This command will generate the bootstrap script and place them in /srv/www/htdocs/pub/bootstrap .

Alternatively you may use the Web UI to create your bootstrap script template. For more information, see .

Use the Web UI to create your keys. From the Web UI proceed to *Overview > Tasks > Manage Activation Keys* .

7.6 Using Activation Keys when Registering Salt Minions

With the addition of Salt to Uyuni 3 states should now be considered best practice over the more traditional way of combining activation keys. Although states allow for more configuration options you need to place the new system within the correct group so the desired states will be applied to the system. Using an activation key on your minions will place the system within the correct group automatically.

You should be aware of a few facts when working with Salt over traditional activation keys:

- Currently we do not support specifying an activation key on the minion on-boarding page.
- Activation keys used with Salt minions are the same as those used with traditional systems and may be shared.
- The equivalent of specifying a key using the traditional bootstrap method is to place the desired key in the grain of a minion. For more information on grains, see <https://docs.saltstack.com/en/latest/topics/targeting/grains.html> ↗
- Once a minion has been accepted either from the *Salt > Keys* page located in the Web UI or from the command line, all configurations specified by the activation key placed within a salt grain will be applied.
- Currently you may only use one activation key when working with salt. You cannot combine them, despite this, salt states allow for even more control.

7.6.1 Using an Activation Key and Custom Grains File

Create a custom grains file and place it on the minion here:

```
# /etc/salt/grains
```

Then add the following lines to the grains file replacing 1-sles12-sp2 with your activation key label:

```
susemanager:  
  activation_key: 1-sles12-sp2
```

Now restart the minion with:

```
# systemctl restart salt-minion
```

7.6.2 Using an Activation Key in the Minion Configuration File

You may also place the activation key grain within the minion configuration file located in:

```
# /etc/salt/minion
```

Now add the following lines to the minion configuration file replacing 1-sles12-sp2 with your activation key label:

```
grains:
  susemanager:
    activation_key: 1-sles12-sp2
```

Reboot the minion with:

```
# systemctl restart salt-minion
```


8 Contact Methods

8.1 Selecting a Contact Method

Uyuni provides several methods for communication between client and server. All commands your Uyuni server sends its clients to do will be routed through one of them. Which one you select will depend on your network infrastructure. The following sections provide a starting point for selecting a method which best suits your network environment.



Note: Contact Methods and Salt

This chapter is only relevant for traditional clients as Salt clients (minions) utilize a Salt specific contact method. For general information about Salt clients, see .

8.2 Default (the Uyuni Daemon `rhnsd`)

The Uyuni daemon (`rhnsd`) runs on client systems and periodically connects with Uyuni to check for new updates and notifications. The daemon, which runs in the background, is started by `rhnsd.service` . By default, it will check every 4 hours for new actions, therefore it may take some time for your clients to begin updating after actions have been scheduled for them.

To check for updates, `rhnsd` runs the external `mgr_check` program located in `/usr/sbin/` . This is a small application that establishes the network connection to Uyuni. The SUSE Manager daemon does not listen on any network ports or talk to the network directly. All network activity is done via the `mgr_check` utility.



Warning: Auto accepting (EULAs)

When new packages or updates are installed on the client using Uyuni, any end user licence agreements (EULAs) are automatically accepted. To review a package EULA, open the package detail page in the Web UI.

This figure provides an overview of the default `rhnsd` process path. All items left of the `Python XMLRPC server` block represent processes running on a Uyuni client.




FIGURE 8.1: RHNSD CONTACT METHOD

8.2.1 Configuring Uyuni rhnsd Daemon

The Uyuni daemon can be configured by editing the file on the client:

```
/etc/sysconfig/rhn/rhnsd
```

This is the configuration file the rhnsd initialization script uses. An important parameter for the daemon is its check-in frequency. The default interval time is four hours (240 minutes). If you modify the configuration file, you must as root restart the daemon with `systemctl rhnsd restart`.



Important: Minimum Allowed Check-in Parameter

The minimum allowed time interval is one hour (60 minutes). If you set the interval below one hour, it will change back to the default of 4 hours (240 minutes).

8.2.2 Viewing rhnsd Daemon Status

You can view the status of rhnsd by typing the command `systemctl status rhnsd` as root.

8.3 Push via SSH

Push via SSH is intended to be used in environments where your clients cannot reach the Uyuni server directly to regularly check in and, for example, fetch package updates.

In detail, this feature enables a Uyuni located within an internal network to manage clients located on a “Demilitarized Zone” (DMZ) outside of the firewall protected network. Due to security reasons, no system on a DMZ is authorized to open a connection to the internal network and therefore your Uyuni server. The solution is to configure Push via SSH which utilizes an encrypted tunnel from your Uyuni server on the internal network to the clients located on the DMZ. After all actions/events are executed, the tunnel is closed. The server will contact the clients in regular intervals (using SSH) to check in and perform all actions and events.

! Important: Push via SSH Unsupported Actions

Certain actions are currently not supported on scheduled clients which are managed via Push via SSH. This includes re-installation of systems using the provisioning module.

The following figure provides an overview of the Push via SSH process path. All items left of the Taskomatic block represent processes running on a Uyuni client.

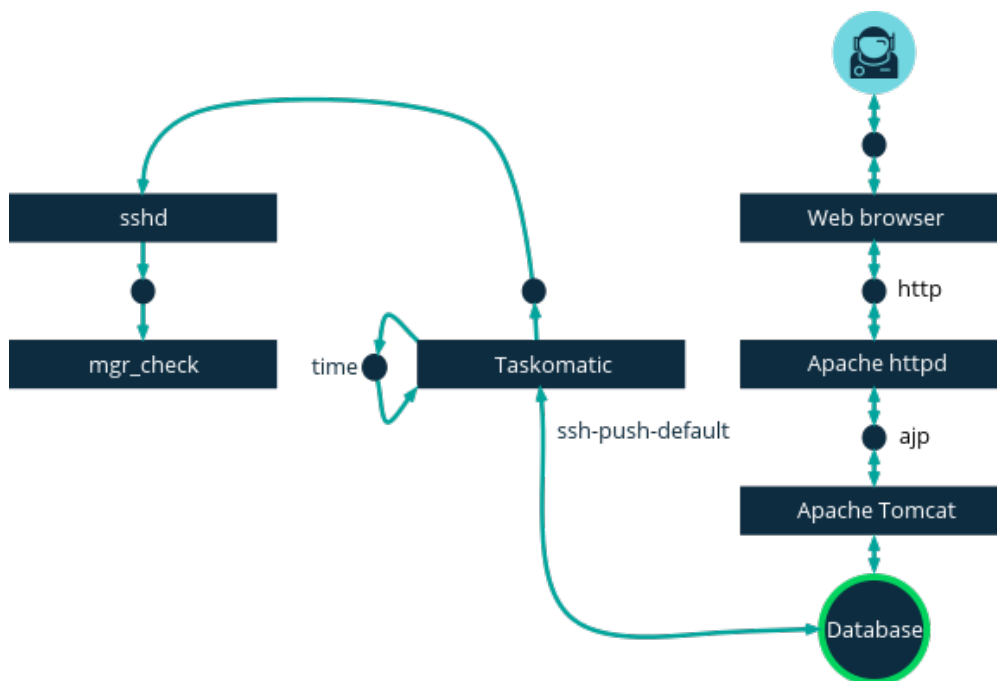


FIGURE 8.2: PUSH VIA SSH CONTACT METHOD

8.3.1 Configuring the Server for Push via SSH

For tunneling connections via SSH, two available port numbers are required, one for tunneling HTTP and the second for tunneling via HTTPS (HTTP is only necessary during the registration process). The port numbers used by default are 1232 and 1233. To overwrite these, add two custom port numbers greater than 1024 to /etc/rhn/rhn.conf like this:

```
ssh_push_port_http = high port 1
ssh_push_port_https = high port 2
```

If you would like your clients to be contacted via their hostnames instead of an IP address, set the following option:

```
ssh_push_use_hostname = true
```

It is also possible to adjust the number of threads to use for opening client connections in parallel. By default two parallel threads are used. Set taskomatic.ssh_push_workers in /etc/rhn/rhn.conf like this:

```
taskomatic.ssh_push_workers = number
```

8.3.2 Using sudo with Push via SSH

For security reasons you may desire to use sudo and SSH into a system as a user other than root. The following procedure will guide you through configuring sudo for use with Push via SSH.



Note: sudo Requirements

The packages spacewalk-taskomatic >= 2.1.165.19 and spacewalk-certs-tools ⇒ 2.1.6.7 are required for using sudo with Push via SSH.

PROCEDURE: CONFIGURING SUDO

1. Set the following parameter on the server located in /etc/rhn/rhn.conf.

```
ssh_push_sudo_user = `user`
```

The server will use sudo to ssh as the configured user.

2. You must create the user specified in *Procedure: Configuring sudo* on each of your clients and the following parameters should be commented out within each client's `/etc/sudoers` file:

```
#Defaults targetpw # ask for the password of the target user i.e. root
#ALL ALL=(ALL) ALL # WARNING! Only use this together with 'Defaults targetpw'!
```

3. Add the following lines beneath the `\## User privilege specification` section of each client's `/etc/sudoers` file:

```
<user> ALL=(ALL) NOPASSWD:/usr/sbin/mgr_check
<user> ALL=(ALL) NOPASSWD:/home/<user>/enable.sh
<user> ALL=(ALL) NOPASSWD:/home/<user>/bootstrap.sh
```

4. On each client add the following two lines to the `/home/user/.bashrc` file:

```
PATH=$PATH:/usr/sbin
export PATH
```

8.3.3 Client Registration

As your clients cannot reach the server, you will need to register your clients from the server. A tool for performing registration of clients from the server is included with Uyuni and is called `mgr-ssh-push-init`. This tool expects a client's hostname or IP address and the path to a valid bootstrap script located in the server's filesystem for registration as parameters.



Important: Specifying Ports for Tunneling before Registering Clients

The ports for tunneling need to be specified before the first client is registered. Clients already registered before changing the port numbers must be registered again, otherwise the server will not be able to contact them anymore.



Note: `mgr-ssh-push-init` Disables `rhnsd`

The `mgr-ssh-push-init` command disables the `rhnsd` daemon which normally checks for updates every 4 hours. Because your clients cannot reach the server without using the Push via SSH contact method, the `rhnsd` daemon is disabled.

For registration of systems which should be managed via the Push via SSH tunnel contact method, it is required to use an activation key that is configured to use this method. Normal Push via SSH is unable to reach the server. For managing activation keys, see [Chapter 7, Activation Key Management](#).

Run the following command as root on the server to register a client:

```
# mgr-ssh-push-init --client client --register \  
/srv/www/htdocs/pub/bootstrap/bootstrap_script --tunnel
```

To enable a client to be managed using Push via SSH (without tunneling), the same script may be used. Registration is optional since it can also be done from within the client in this case. mgr-ssh-push-init will also automatically generate the necessary SSH key pair if it does not yet exist on the server:

```
# mgr-ssh-push-init --client`client`--register bootstrap_script
```

When using the Push via SSH tunnel contact method, the client is configured to connect Uyuni via the high ports mentioned above (see /etc/sysconfig/rhn/up2date). Tools like rhnc and zypper will need an active SSH session with the proper port forwarding options in order to access the Uyuni API. To verify the Push via SSH tunnel connection manually, run the following command on the Uyuni server:

```
# ssh -i /root/.ssh/id_susemanager -R high port: susemanager :443`client`zypper ref
```

8.3.4 API Support for Push via SSH

The contact method to be used for managing a server can also be modified via the API. The following example code (python) shows how to set a system's contact method to ssh-push. Valid values are:

- default (pull)
- ssh-push
- ssh-push-tunnel

```
client = xmlrpclib.Server(SUMA_HOST + "/rpc/api", verbose=0)  
key = client.auth.login(SUMA_LOGIN, SUMA_PASSWORD)  
client.system.setDetails(key, 1000012345, {'contact_method' : 'ssh-push'})
```



Note: Migration and Management via Push via SSH

When a system should be migrated and managed using Push via SSH, it requires setup using the `mgr-ssh-push-init` script before the server can connect via SSH. This separate command requires human interaction to install the server's SSH key onto the managed client (root password). The following procedure illustrates how to migrate an already registered system:

PROCEDURE: MIGRATING REGISTERED SYSTEMS

1. Setup the client using the `mgr-ssh-push-init` script (without `--register`).
2. Change the client's contact method to `ssh-push` or `ssh-push-tunnel` respectively (via API or Web UI).

Existing activation keys can also be edited via API to use the Push via SSH contact method for clients registered with these keys:

```
client.activationkey.setDetails(key, '1-mykey', {'contact_method' : 'ssh-push'})
```

8.3.5 Proxy Support with Push via SSH

It is possible to use Push via SSH to manage systems that are connected to the Uyuni server via a proxy. To register a system, run `mgr-ssh-push-init` on the proxy system for each client you wish to register. Update your proxy with the latest packages to ensure the registration tool is available. It is necessary to copy the ssh key to your proxy. This can be achieved by executing the following command from the server:

```
{prompt.root}mgr-ssh-push-init --client`proxy`
```

8.4 Push via Salt SSH

Push via Salt SSH is intended to be used in environments where your Salt clients cannot reach the Uyuni server directly to regularly checking in and, for example, fetch package updates.



Note: Push via SSH

This feature is not related to Push via SSH for the traditional clients. For Push via SSH, see [Section 8.3, “Push via SSH”](#).

8.4.1 Overview




FIGURE 8.3: PUSH VIA SALT SSH CONTACT METHOD

Salt provides “Salt SSH” (`salt-ssh`), a feature to manage clients from a server. It works without installing Salt related software on clients. Using Salt SSH there is no need to have minions connected to the Salt master. Using this as a Uyuni connect method, this feature provides similar functionality for Salt clients as the traditional Push via SSH feature for traditional clients.

This feature allows:

- Managing Salt entitled systems with the Push via SSH contact method using Salt SSH.
- Bootstrapping such systems.

8.4.2 Requirements

- SSH daemon must be running on the remote system and reachable by the `salt-api` daemon (typically running on the Uyuni server).
- Python must be available on the remote system (Python must be supported by the installed Salt). Currently: python 2.6.



Note: Unsupported Systems

Red Hat Enterprise Linux and CentOS versions ≤ 5 are not supported because they do not have Python 2.6 by default.

8.4.3 Bootstrapping

To bootstrap a Salt SSH system, proceed as follows:

1. Open the *Bootstrap Minions >]* dialog in the Web UI (menu: `Systems[Bootstrapping]`).
2. Fill out the required fields. Select an *Activation Key >]* with the menu: `Push via SSH[` contact method configured. For more information about activation keys, see .
3. Check the *Manage system completely via SSH* option.
4. Confirm with clicking the *Bootstrap* button.

Now the system will be bootstrapped and registered in Uyuni. If done successfully, it will appear in the *Systems* list.

8.4.4 Configuration

There are two kinds of parameters for Push via Salt SSH:

- Bootstrap-time parameters - configured in the *Bootstrapping* page:
 - Host
 - Activation key
 - Password - used only for bootstrapping, not saved anywhere; all future SSH sessions are authorized via a key/certificate pair
- Persistent parameters - configured Uyuni-wide:
 - sudo user - same as in *Section 8.3.2, "Using sudo with Push via SSH"*.

8.4.5 Action Execution

The Push via Salt SSH feature uses a taskomatic job to execute scheduled actions using `salt-ssh`. The taskomatic job periodically checks for scheduled actions and executes them. While on traditional clients with SSH push configured only `rhn_check` is executed via SSH, the Salt SSH push job executes a complete `salt-ssh` call based on the scheduled action.

8.4.6 Known Limitation

- OpenSCAP auditing is not available on Salt SSH minions.
- Beacons do not work with Salt SSH.
 - Installing a package on a system using `zypper` will not invoke the package refresh.
 - Virtual Host functions (for example, a host to guests) will not work if the virtual host system is Salt SSH-based.

8.4.7 For More Information

For more information, see

- https://wiki.microfocus.com/index.php/SUSE_Manager/SaltSSHServerPush ↗
- <https://docs.saltstack.com/en/latest/topics/ssh/> ↗

8.5 OSAD

OSAD is an alternative contact method between Uyuni and its clients. By default, Uyuni uses rhnsd, which contacts the server every four hours to execute scheduled actions. OSAD allows registered client systems to execute scheduled actions immediately.

OSAD has several distinct components:

- The osa-dispatcher service runs on the server, and uses database checks to determine if clients need to be pinged, or if actions need to be executed.
- The osad service runs on the client. It responds to pings from osa-dispatcher and runs mgr_check to execute actions when directed to do so.
- The jabberd service is a daemon that uses the XMPP protocol for communication between the client and the server. The jabberd service also handles authentication.
- The mgr_check tool runs on the client to execute actions. It is triggered by communication from the osa-dispatcher service.

The osa-dispatcher periodically runs a query to check when clients last showed network activity. If it finds a client that has not shown activity recently, it will use jabberd to ping all osad instances running on all clients registered with your Uyuni server. The osad instances respond to the ping using jabberd, which is running in the background on the server. When the osa-dispatcher receives the response, it marks the client as online. If the osa-dispatcher fails to receive a response within a certain period of time, it marks the client as offline.

When you schedule actions on an OSAD-enabled system, the task will be carried out immediately. The osa-dispatcher periodically checks clients for actions that need to be executed. If an outstanding action is found, it uses jabberd to execute mgr_check on the client, which will then execute the action.

8.5.1 Enabling and Configuring OSAD

This section covers enabling the osa-dispatcher and osad services, and performing initial setup.

OSAD clients use the fully qualified domain name (FQDN) of the server to communicate with the `osa-dispatcher` service.

SSL is required for `osad` communication. If SSL certificates are not available, the daemon on your client systems will fail to connect. Make sure your firewall rules are set to allow the required ports. For more information, see [Table 1.1, “Required Server Ports”](#).

PROCEDURE: ENABLING OSAD

1. On your Uyuni server, as the root user, start the `osa-dispatcher` service:

```
systemctl start osa-dispatcher
```

2. On each client machine, install the `osad` package from the `Tools` child channel. The `osad` package should be installed on clients only. If you install the `osad` package on your Uyuni Server, it will conflict with the `osa-dispatcher` package.
3. On the client systems, as the root user, start the `osad` service:

```
systemctl start osad
```

Because `osad` and `osa-dispatcher` are run as services, you can use standard commands to manage them, including `stop`, `restart`, and `status`.

Configuration and Log Files. Each OSAD component is configured by local configuration files. We recommend you keep the default configuration parameters for all OSAD components.

Component	Location	Path to Configuration File
<code>osa-dispatcher</code>	Server	<code>/etc/rhn/rhn.conf</code> Section: <code>OSA configuration</code>
<code>osad</code>	Client	<code>/etc/sysconfig/rhn/osad.conf</code> <code>/etc/sysconfig/rhn/up2date</code>
<code>osad</code> log file	Client	<code>/var/log/osad</code>
<code>jabberd</code> log file	Both	<code>/var/log/messages</code>

Troubleshooting OSAD. If your OSAD clients cannot connect to the server, or if the `jabberd` service takes a lot of time responding to port 5552, it could be because you have exceeded the open file count.

Every client needs one always-open TCP connection to the server, which consumes a single file handler. If the number of file handlers currently open exceeds the maximum number of files that jabberd is allowed to use, jabberd will queue the requests, and refuse connections.

To resolve this issue, you can increase the file limits for jabberd by editing the /etc/security/limits.conf configuration file and adding these lines:

```
jabbersoftnofile5100
jabberhardnofile6000
```


Calculate the limits required for your environment by adding 100 to the number of clients for the soft limit, and 1000 to the current number of clients for the soft limit. In the example above, we have assumed 500 current clients, so the soft limit is 5100, and the hard limit is 6000.

You will also need to update the max_fds parameter in the /etc/jabberd/c2s.xml file with your chosen hard limit:

```
<max_fds>6000</max_fds>
```

9 Advanced Patch Lifecycle Management

Keeping systems patched and secure remains one of the greatest ongoing challenges that you will face as an administrator. Both proprietary and open-source companies are constantly working to provide updates which fix flaws discovered within their software products.

For the official *Best Practice Guide* on “Advanced Patch Lifecycle Management” , see <https://www.suse.com/documentation/suse-best-practices/susemanager/data/susemanager.html> .

10 Live Patching with SUSE Manager

10.1 Introduction

Under normal circumstances a system needs to be rebooted after a kernel update. SLE Live Patching allows you skipping the reboot by applying a subset of Linux kernel releases injected via kGraft live patching technology.

In the following sections you will learn how to use SLE Live Patching to avoid the typical reboot requirement after updating a system kernel.

For in depth information covering kGraft use, see https://www.suse.com/documentation/sles-12/singlehtml/book_sle_admin/book_sle_admin.html#cha.kgraft.

10.2 Initial Setup Requirements

To work with SLE Live Patching the following expectations are assumed:

- Uyuni fully updated.
- At least 1 Salt Minion running SLES 12 SP1 or later and registered with Uyuni .
- The matching SLES 12 SPx channels including the SLE Live Patching child channel fully synced.

10.3 Live Patching Setup

1. Subscribe all systems to be managed via live patching to your fully synced live patching child channels within your systems base channel by browsing to *Software > Software Channels* . Select both live patching channels and change subscription.



Note

When subscribing to a channel that contains a product, the product package will automatically be installed on traditionally registered systems and added to the package state on Salt managed systems. For Salt managed systems please apply the highstate to push these changes to your systems.

doctesting-minsles12sp2_DEV.tf.local

Delete System Remove from SSM

Details Software Groups Audit States Events

Patches Packages Software Channels SP Migration

When subscribing to a channel that contains a product, the product package will automatically be installed on traditional registered systems and added to the package state on salt managed system. For salt managed systems please apply the highstate to push these changes to the system.

Software Channel Subscriptions

This system is subscribed to the base channel, listed at top, and to the checked channels beneath, if any. Disabled checkboxes indicate channels that can't be manually subscribed or unsubscribed from.

- SLES12-SP2-Pool for x86_64
 - ☒ SLE-Live-Patching12-Pool for x86_64 SP2
 - ☒ SLE-Live-Patching12-Updates for x86_64 SP2
 - ☒ SLE-Manager-Tools12-Pool x86_64 SP2
 - ☒ SLE-Manager-Tools12-Updates x86_64 SP2
 - ☒ SLES12-SP2-Updates for x86_64
 - ☐ SUSE-Manager-Proxy-3.1-Pool for x86_64
 - ☐ SUSE-Manager-Proxy-3.1-Updates for x86_64

Change Subscriptions

1. Use the search field listed under *Software > Packages > Install* to install the latest kgrapt package to all systems to be managed via live patching.

doctesting-minsles12sp2_DEV.tf.local

Delete System Remove from SSM

Details Software Groups Audit States Events

Patches Packages Software Channels SP Migration

List / Remove Upgrade Install Non Compliant

Installable Packages

The following packages may be installed on this system.

Select All 1 - 25 of 78 (1 selected) Initial Selected Packages

The list of 78 item(s) below is filtered. Clear filter to see all 3,402 items.

kgrapt Alphabetical index 25 items per page

Package Name	Architecture
<input checked="" type="checkbox"/> kgrapt-1.0-22.1	x86_64
<input type="checkbox"/> kgrapt-devel-1.0-22.1	x86_64
<input type="checkbox"/> kgrapt-manual_en-12.8.2	noarch
<input type="checkbox"/> kgrapt-patch-3.12_32-25-default-1-2.7	x86_64
<input type="checkbox"/> kgrapt-patch-3.12_32-25-xen-1-2.7	x86_64

1. Apply the highstate to enable live patching:

doctesting-minsles12sp2_DEV.tf.local

Delete System Remove from SSM

Details Software Groups Audit States Events

Highstate Packages Custom

Apply Highstate

Earliest: 6/13/17 3:09 pm CEST

1. Once the highstate has been applied on Salt systems or the package has been installed on traditional systems browse to the systems details page for confirmation that live patching has been enabled. You can check the live patching state listed under the *System Info* > *Kernel* table field:

The screenshot shows the Spacewalk web interface for a system named 'doctesting-minsles12sp2_DEV.tf.local'. The interface includes a top navigation bar with tabs like Details, Software, Groups, Audit, States, and Events. Below this, there's a 'System Status' section indicating the system is up to date. The main content area is divided into several panels: 'System Info' (showing Hostname, IP Address, IPv6 Address, Virtualization, UUID, Kernel, and SUSE Manager System ID), 'System Events' (showing Checked In, Registered, and Last Booted times), 'System Properties' (showing System Types, Contact Method, Auto Patch Update, System Name, Description, and Location), and 'Subscribed Channels' (listing various SLES and SLE channels for x86_64 SP2).

10.4 Cloning Channels

It is considered best practice to clone a vendor channel that will be modified into a new channel with one of the following prefix names: `dev`, `testing`, and `prod`. In the following procedure you will clone the default vendor channel into a new channel named `dev-sles12-sp3-pool-x86_64` using the command line.

1. Open a terminal and as root enter:

```
# spacewalk-manage-channel-lifecycle --list-channels
Spacewalk Username: admin
Spacewalk Password:
Channel tree:

1. sles{sles-version}-sp{sp-ver}-pool-x86_64
   \__ sle-live-patching{sles-version}-pool-x86_64-sp{sp-ver}\__ sle-live-
   patching{sles-version}-updates-x86_64-sp{sp-ver}\__ sle-manager-tools{sles-version}-
```

```
pool-x86_64-sp{sp-ver}\__ sle-manager-tools{sles-version}-updates-x86_64-sp{sp-ver}\__ sles{sles-version}-sp{sp-ver}-updates-x86_64
```

2. Now use the `--init` argument to automatically create a new development clone of the original vendor channel:

```
spacewalk-manage-channel-lifecycle --init -c sles{sles-version}-sp{sp-ver}-pool-x86_64
```

10.5 Removing Non-live Kernel Patches from the Cloned Channels

In the following procedure you will remove all kernel patch updates located in the `dev-sles12-spSP3 -updates-x86_64` channel that require a reboot. You created `dev-sles12 -spSP3 -updates-x86_64` during [???TITLE???](#).

1. Check the current kernel version in use on your client:

```
# uname -r
3.12.62-60.64.8-default
```

2. From the Uyuni Web UI select *Software > Manage Software Channels > Overview > dev-sles12-sp3-updates-x86_64 > Patches > List/Remove* . Type kernel in the search field. Find the kernel version that matches the kernel in use on your minion.
3. Remove all kernel update versions that are later than the current kernel.
4. Your channel is now ready to promote for testing SLE Live Patching.

10.6 Promoting Channels

The following procedure will guide you through promoting and cloning a development channel to a testing channel. You will change the subscription from the dev repositories on your client to the new testing channel repositories. You will also add the SLE Live Patching child channels to your client.

1. Promote and clone the `dev-sles12 -spSP3 -pool-x86_64` to a new testing channel:

```
{prompt.root}spacewalk-manage-channel-lifecycle -promote -c dev-sles{sles-version}-sp{sp-ver}-pool-x86_64
```

2. From the Uyuni Web UI under the *Systems >]* tab select your client system to view the *menu:System Details[* page. Select *Software > Software Channels* . From the Software Channels page you can edit which channels a system is subscribed to. Select the new base software channel, in this case it should be `test-sles12-sp3-pool-x86_64` . Click the *Confirm >]* button to switch the Base Software Channel and finalize it by clicking the *menu:Modify Base Software Channel[* button.
3. From the *Software Channels >]* page select and add both *SLE Live Patching child channels* by clicking the *menu:Change Subscriptions[* button.

10.7 Applying Live Patches to a Kernel

The following procedure will guide you through selecting and viewing available CVE Patches (Common Vulnerabilities and Exposures) then applying these kernel updates using the new SLE Live Patching feature.

1. Select your SLES 12 SPSP3 minion from the *Systems >]* page to view its *menu:System Details[* . Once you have added the SLES 12 SPSP3 Updates child channel to your client, you should see several Critical software updates available. Click on Critical to see a list of available patches. Select any of these patches listed with the following synopsis: *Important: Security update for the Linux kernel*. All fixed security bugs will be listed along with their number. For example:(CVE-2016-8666)



Important: Reboot Icon

Normal or non-live kernel patches always require a reboot. In Uyuni these are represented by a Reboot Required icon located next to the Security shield icon.

2. You can search for individual CVE's by selecting the *Audit* tab from the navigation menu. Try searching for CVE-2016-8666 . You will see that the patch is available in the vendor update channel and the systems it applies to will be listed.

Important: CVE Availability

Not all security issues can be fixed by applying a live patch. Some security issues can only be fixed by applying a full kernel update and will require a reboot. The assigned CVE numbers for these issues are not included in live patches. A CVE audit will display this requirement.

11 Uyuni Server Migration

11.1 Service Pack Migration Introduction

You can upgrade the underlying operating system and also migrate Uyuni server from one patch level to the other (SP migration) or from one version to the next. This works for migrating Uyuni server 3.0 to version 3.1, or version 3.1 to version 3.2. For migrating version 3.0 to version 3.1, see the product documentation for Uyuni 3.1:

<https://www.suse.com/documentation/suse-manager-3/>

11.2 Service Pack Migration

SUSE Manager uses SUSE Linux Enterprise Server 12 as its underlying operating system. Therefore Service Pack migration (for example, from version 12 SP1 to 12 SP3) may be performed in the same way as a typical SLES migration.



Warning: Upgrading PostgreSQL to Version 9.6 Before Migrating to SLES12 SP3

Before migrating the underlying system to SUSE Linux Enterprise 12 SP3 you must upgrade PostgreSQL to version 9.6.

The migration needs PostgreSQL 9.4 and 9.6 installed in parallel and PostgreSQL 9.4 is only available in SUSE Linux Enterprise 12 SP2. For more information, see [Section 11.3, “Upgrading PostgreSQL to Version 9.6”](#).

SUSE offers a graphical and command line tool for upgrading to a new service pack. Comprehensive documentation for executing service pack migration scenarios is located in the SUSE Linux Enterprise Server documentation chapter https://www.suse.com/documentation/sles-12/book_sle_deployment/data/cha_update_sle.html.

11.3 Upgrading PostgreSQL to Version 9.6



Warning: Migrating to SLES 12 SP3

Uyuni Server 3.1 must *not* be migrated to SLES 12 SP3 before upgrading PostgreSQL to version 9.6.

The upgrade needs PostgreSQL 9.4 and 9.6 installed in parallel. PostgreSQL 9.4 is only available in SLES 12 SP2.

Before starting the update, prepare an up-to-date backup of your database.

On existing installations of Uyuni Server 3.1 you must migrate from PostgreSQL 9.4 to version 9.6.

```
/usr/lib/susemanager/bin/pg-migrate.sh
```

During the upgrade your Uyuni Server will not be accessible.

The upgrade will create a copy of the database under `/var/lib/pgsql` and thus needs sufficient disk space to hold two copies (versions 9.4 and 9.6) of the database. Because it does a full copy of the database, it also needs considerable time depending on the size of the database and the I/O speed of the storage system.

If your system is short on disk space you can do a fast, in-place upgrade by running

```
/usr/lib/susemanager/bin/pg-migrate.sh fast
```

The fast upgrade usually only takes a few minutes and uses no additional disk space. However, if the upgrade fails, you will need to restore the database from a backup.

For more information, see https://wiki.microfocus.com/index.php?title=SUSE_Manager/postgresql96.

11.4 Updating Uyuni

This section provides information on performing regular updates and running a `space-walk-schema-upgrade` on your PostgreSQL database.

PROCEDURE: UPDATING UYUNI

1. As the root user, stop Spacewalk services:

```
spacewalk-service stop
```

2. Apply the latest patches:

```
zypper patch
```

3. You will be informed if a new database schema was included in the latest patch. Ensure the database service is running:

```
rcpostgresql start
```

4. Perform the upgrade:

```
spacewalk-schema-upgrade
```

5. Restart Spacewalk services:

```
spacewalk-service start
```



Important: Restart of Services and Applications

Services affected by a package update are not automatically restarted after an update. You need to restart these services manually to avoid potential failures.

You may use `zypper ps` to check for any applications which may be using old code. Restart these applications.

11.5 Migrating Uyuni version 3.1 to 3.2

The migration can either be done with the Online Migration tool (YaST) or with the Zypper command line tool.

Requirements. Uyuni 3.2 requires SLES 12 SP3, with PostgreSQL version 9.6. Check the release notes for more information about these requirements. If you want to upgrade from an earlier version of Uyuni, check the relevant product documentation.



Note: Reduce Installation Size

When performing the migration, YaST will install all recommended packages. Especially in the case of custom minimal installations, this may increase the installation size of the system significantly.

To change this default behavior and allow only required packages, adjust `/etc/zypp/zypp.conf` and set the following variable:

```
solver.onlyRequires = true
installRecommends=false # or commented
```

This changes the behavior of all package operations, such as the installation of patches or new packages.

11.5.1 Using YaST



Warning: Checking PostgreSQL Version

Before migrating to SLES 12 SP3, check whether PostgreSQL is already updated to version 9.6. For more information, see [Section 11.3, “Upgrading PostgreSQL to Version 9.6”](#).

To perform the migration with YaST, use the *Online Migration* tool:

PROCEDURE: MIGRATING USING YAST

1. If you are logged into a GNOME session running on the machine you are going to update, switch to a text console. Running the update from within a GNOME session is not recommended. This does not apply when being logged in from a remote machine (unless you are running a VNC session with GNOME).
2. Start in YaSTSystem > *Online Migration* (`yast2 migration`). YaST will show possible migration targets with detailed summaries.

In case of trouble, resolve the following issues first:

- If the *Online Migration* is not available, install the `yast2-migration` package and its dependencies. Restart YaST, otherwise the newly installed module will not be shown in the control center.
- If there are “old” online updates available for installation, the migration tool will warn and ask to install them now before starting the actual migration. It is recommended to install all updates before proceeding.

11.5.2 Using zypper



Warning: Checking PostgreSQL Version

Before migrating to SLES 12 SP3, check whether PostgreSQL is already updated to version 9.6. For more information, see [Section 11.3, “Upgrading PostgreSQL to Version 9.6”](#).

To perform the migration with Zypper on the command-line, use the `zypper migration` subcommand tool:

PROCEDURE: MIGRATING USING `zypper migration`

1. If you are logged into a GNOME session running on the machine you are going to update, switch to a text console. Running the update from within a GNOME session is not recommended. This does not apply when being logged in from a remote machine (unless you are running a VNC session with GNOME).
2. The `zypper migration` subcommand show possible migration targets and a summary.
In case of trouble, resolve the following issues first:
 - If the `migration` subcommand is not available install the `zypper-migration-plugin` package and its dependencies.
 - If there are “old” online updates available for installation, the migration tool will warn and ask to install them now before starting the actual migration. It is recommended to install all updates before proceeding.
3. If more than one migration target is available for your system, select one from the list (specify the number).

4. Read the notification and update the Uyuni database schema as described ([spacewalk-schema-upgrade](#)).
5. Make sure Uyuni is up and running ([spacewalk-service start](#)).

After finishing the migration procedure Uyuni 3.2 on SLES 12 SP3 is available to be used.

11.6 Uyuni Migration from Version 2.1 to Version 3

The migration from Uyuni 2.1 to Uyuni 3 works in the same way as a migration from Red Hat Satellite to Uyuni. The migration happens from the original machine to a new one. There is no in-place migration. While this has the drawback that you temporarily need two machines, it also has the advantage that the original machine will remain fully functional in case something goes wrong.



Important: Migration Process

The whole process may be tricky, so it is strongly advised that the migration is done by an experienced consultant.

Given the complexity of the product, the migration is an “all-or-nothing” procedure- if something goes wrong you will need to start all over. Error handling is very limited. Nevertheless it should work more or less out of the box if all the steps are carefully executed as documented.



Note: Time-Consuming Operation

The migration involves dumping the whole database on the source machine and restoring it on the target machine. Also all of the channels and packages need to be copied to the new machine, so expect the whole migration to take several hours,

11.6.1 Prerequisites



Warning: Latest Updates

The source machine needs to run Uyuni 2.1 with all the latest updates applied. Before starting the migration process, make sure that the machine is up to date and all updates have been installed successfully.

Only machines running with the embedded PostgreSQL database may be migrated in one go. For the migration of an Oracle based installation, a two-step migration is required: First the installation needs to get migrated from Oracle to PostgreSQL (by means of a separate tool) and afterwards the migration to Uyuni 3 can be performed as documented here.

Uyuni 3 no longer supports Novell Customer Center but only SCC (SUSE Customer Center). Therefore, you can migrate a machine only after it has been switched to SCC. The migration script will check if the installation has already been switched to SCC and will terminate if this is not the case. Switch to SCC on the source machine and repeat the migration. During migration the database from the source machine needs to get dumped and this dump needs to be temporarily stored on the target system. The dump gets compressed with `gzip` using the default compression options (maximum compression only yields about 10% of space savings but costs a lot of runtime); so check the disk usage of the database with:

```
{prompt.root}du -sch /var/lib/pgsql/data
```

This will ensure that you have at least 30 % of this value available in `/var/spacewalk/tmp` . These values from a test migration should aid in illustrating space requirements:

```
suma21:/var/lib/pgsql# du -sch data
1,8G    data
1,8G    total
suma21:/var/spacewalk/tmp# ls -lh susemanager.dmp.gz
-rw-r--r-- 1 root root 506M Jan 12 14:58 susemanager.dmp.gz
```

This is a small test installation; for bigger installations the ratio might be better (space required for the database dump might be less than 30%). The dump will be written to the directory `/var/spacewalk/tmp` , the directory will be created if it does not exist yet. If you want the dump to be stored somewhere else, change the definition of the variable `$TMPDIR` on the beginning of the script to suit your needs.

11.6.2 Setup the Target Machine

To prepare the target machine (with the example host name `suma30`) proceed as follows:

PROCEDURE: SETUP TARGET MACHINE

1. On the target machine install SUSE Linux Enterprise Server 12 SP2 including the extension product “Uyuni” .



Important: Background Information on Required Target Machine

It is actually required to install version 12 SP2 on the target machine. On that version you will upgrade the PostgreSQL database from version_9.4 to 9.6. For more information about the PostgreSQL upgrade, see [Section 11.3, “Upgrading PostgreSQL to Version 9.6”](#).

2. Initiate `yast2 susemanagersetup` as you would normally do for an installation of Uyuni. For more information about installing Uyuni, see Book “Getting Started”, Chapter 2 “JeOS Installation”.
3. On the first Uyuni setup screen, ensure that *Migrate a SUSE Manager compatible server >]* is marked instead of *menu:Set up SUSE Manager from scratch[* .
4. On the second screen, enter the name of the source system as *Hostname of source SUSE Manager Server* as well as the domain name. Also enter the database credentials of the source system.
5. On the next screen, you will need to specify the IP address of the Uyuni 3 target system. Normally this value should be pre-set to the correct value and you only should need to press `Enter` . Only in the case of multiple IP addresses you might need to specify the one that should be used during migration.

Important: Faking the Host Name

During the migration process, the target system will fake its host name to be the same as the source system, this is necessary as the host name of a Uyuni installation is vital and should not be changed once set. Therefore do not be confused when logging in to your systems during migration; they both will present you with the same host name.

6. Continue by following the normal installation steps.

Important: Database Parameters

Specify the database parameters using the same database parameters as the source system is recommended. At least, using the the same database credentials as when creating the source or original Uyuni database is mandatory.

Enter your SCC credentials. After all the data has been gathered, YaST will terminate.

The actual migration will not start automatically but needs to be triggered manually as outlined in [Section 11.6.3, “Performing the Migration”](#).

11.6.3 Performing the Migration

A migration is performed by executing the following command:

```
/usr/lib/susemanager/bin/mgr-setup -m
```

This command reads the data gathered during [Procedure: Setup Target Machine](#), sets up Uyuni onto a new target machine and transfers all of the data from the source machine. As several operations need to be performed on the source machine via SSH, you will be prompted once for the root password of the source machine. A temporary SSH key named `migration-key` is created and installed on the source machine, so you need to give the root password only once. The temporary SSH key will be deleted after successful migration. Ideally, this is all you will need to do.

Depending on the size of the installation, the actual migration will take up to several hours. Once finished, you will be prompted to shutdown the source machine, re-configure the network of the target machine to use the same IP address and host name as the original machine and restart it. It should now be a fully functional replacement for your previous Uyuni 2.1 installation. The following numbers illustrate the runtime for dumping and importing a small database:

```
14:53:37   Dumping remote database to /var/spacewalk/tmp/susemanager.dmp.gz on target
           system. Please wait...
14:58:14   Database successfully dumped. Size is: 506M
14:58:29   Importing database dump. Please wait...
15:05:11   Database dump successfully imported.
```

For this example dumping the database takes around five minutes to complete. Importing the dump onto the target system will take an additional seven minutes. For big installations this can take up to several hours. You should also account for the time it takes to copy all the package data to the new machine. Depending on your network infrastructure and hardware, this can also take a significant amount of time.

11.6.4 Speeding up the Migration

A complete migration can consume a lot of time. This is caused by the amount of data that must be copied. Total migration time can be greatly decreased by eliminating the need to copy data prior to performing the migration (for example, channels, packages, auto-install images, and any additional data). You can gather all data via YaST by running the command `mgr-setup -r`. Executing `mgr-setup -r` will copy the data from the old server to the new one. This command may be run at any time and your current server will remain fully functional. Once the migration has been initiated only data changed since running `mgr-setup -r` will need to be transferred which will significantly reduces downtime.

On large installations transferring the database (which involves dumping the database onto the source machine and then importing the dump onto the target system) will still take some time. During the database transfer no write operations should occur therefore the migration script will shut down any Uyuni database services running on the source machine.

11.6.5 Packages on External Storage

Some installations may store the package data on external storage (for example, NFS mount on /var/spacewalk/packages). You do not need to copy this data to the new machine. Edit the script located in /usr/lib/susemanager/bin/mgr-setup and remove the respective rsync command (located around line 345).



Important: Mounting External Storage

Make sure your external storage is mounted on the new machine before starting the system for the first time. Analogue handling for /srv/www/htdocs/pub if appropriate.

In general, all needed files and directories, not copied by the migration tool, should be copied to the new server manually.

11.6.6 Troubleshooting a Broken Web UI after Migration

It is possible that the Web UI may break during migration. This behavior is not a bug, but a browser caching issue. The new machine has the same host name and IP address as the old machine. This duplication can confuse some Web browsers. If you experience this issue reload the page. For example, in Firefox pressing the key combination **Ctrl**–**F5** should resume normal functionality.

11.6.7 Example Session

This is the output of a typical migration:

```
suma30# /usr/lib/susemanager/bin/mgr-setup -m
  Filesystem type for /var/spacewalk is ext4 - ok.
  Open needed firewall ports...
  Migration needs to execute several commands on the remote machine.
  Please enter the root password of the remote machine.
Password:
  Remote machine is SUSE Manager
  Remote system is already migrated to SCC. Good.
  Shutting down remote spacewalk services...
  Shutting down spacewalk services...
  Stopping Taskomatic...
  Stopped Taskomatic.
  Stopping cobbler daemon: ..done
```

```

Stopping rhn-search...
Stopped rhn-search.
Stopping MonitoringScout ...
[ OK ]
Stopping Monitoring ...
[ OK ]
Shutting down osa-dispatcher: ..done
Shutting down httpd2 (waiting for all children to terminate) ..done
Shutting down Tomcat (/usr/share/tomcat6)
..done
Terminating jabberd processes...
    Stopping router ..done
    Stopping sm ..done
    Stopping c2s ..done
    Stopping s2s ..done
Done.
CREATE ROLE
* Loading answer file: /root/spacewalk-answers.
** Database: Setting up database connection for PostgreSQL backend.
** Database: Populating database.
** Database: Skipping database population.
* Configuring tomcat.
* Setting up users and groups.
** GPG: Initializing GPG and importing key.
* Performing initial configuration.
* Configuring apache SSL virtual host.
** /etc/apache2/vhosts.d/vhost-ssl.conf has been backed up to vhost-ssl.conf-swsave
* Configuring jabberd.
* Creating SSL certificates.
** Skipping SSL certificate generation.
* Deploying configuration files.
* Setting up Cobbler..
* Setting up Salt Master.
11:26:47   Dumping remote database. Please wait...
11:26:50   Database successfully dumped.
Copy remote database dump to local machine...
Delete remote database dump...
11:26:50   Importing database dump. Please wait...
11:28:55   Database dump successfully imported.
Schema upgrade: [susemanager-schema-2.1.50.14-3.2.devel21] -> [susemanager-
schema-3.0.5-5.1.develHead]
Searching for upgrade path to: [susemanager-schema-3.0.5-5.1]
Searching for upgrade path to: [susemanager-schema-3.0.5]
Searching for upgrade path to: [susemanager-schema-3.0]
Searching for start path: [susemanager-schema-2.1.50.14-3.2]
Searching for start path: [susemanager-schema-2.1.50.14]

```



```

The path: [susemanager-schema-2.1.50.14] -> [susemanager-schema-2.1.50.15] ->
[susemanager-schema-2.1.51] -> [susemanager-schema-3.0]
Planning to run schema upgrade with dir '/var/log/spacewalk/schema-upgrade/schema-
from-20160112-112856'
Executing spacewalk-sql, the log is in [/var/log/spacewalk/schema-upgrade/schema-
from-20160112-112856-to-susemanager-schema-3.0.log].
(248/248) apply upgrade [schema-from-20160112-112856/99_9999-upgrade-end.sql]      e-
suse-channels-to-public-channel-family.sql.postgresql]
The database schema was upgraded to version [susemanager-schema-3.0.5-5.1.develHead].
Copy files from old SUSE Manager...
receiving incremental file list
./
packages/

sent 18 bytes  received 66 bytes  168.00 bytes/sec
total size is 0  speedup is 0.00
receiving incremental file list
./
RHN-ORG-TRUSTED-SSL-CERT
res.key
rhn-org-trusted-ssl-cert-1.0-1.noarch.rpm
suse-307E3D54.key
suse-39DB7C82.key
suse-9C800ACA.key
bootstrap/
bootstrap/bootstrap.sh
bootstrap/client-config-overrides.txt
bootstrap/sm-client-tools.rpm

sent 189 bytes  received 66,701 bytes  44,593.33 bytes/sec
total size is 72,427  speedup is 1.08
receiving incremental file list
./
.mtime
lock
web.ss
config/
config/distros.d/
config/images.d/
config/profiles.d/
config/repos.d/
config/systems.d/
kickstarts/
kickstarts/autoyast_sample.xml
loaders/
snippets/
triggers/

```

```
triggers/add/
triggers/add/distro/
triggers/add/distro/post/
triggers/add/distro/pre/
triggers/add/profile/
triggers/add/profile/post/
triggers/add/profile/pre/
triggers/add/repo/
triggers/add/repo/post/
triggers/add/repo/pre/
triggers/add/system/
triggers/add/system/post/
triggers/add/system/pre/
triggers/change/
triggers/delete/
triggers/delete/distro/
triggers/delete/distro/post/
triggers/delete/distro/pre/
triggers/delete/profile/
triggers/delete/profile/post/
triggers/delete/profile/pre/
triggers/delete/repo/
triggers/delete/repo/post/
triggers/delete/repo/pre/
triggers/delete/system/
triggers/delete/system/post/
triggers/delete/system/pre/
triggers/install/
triggers/install/post/
triggers/install/pre/
triggers/sync/
triggers/sync/post/
triggers/sync/pre/

sent 262 bytes  received 3,446 bytes  7,416.00 bytes/sec
total size is 70,742  speedup is 19.08
receiving incremental file list
kickstarts/
kickstarts/snippets/
kickstarts/snippets/default_motd
kickstarts/snippets/keep_system_id
kickstarts/snippets/post_delete_system
kickstarts/snippets/post_reactivation_key
kickstarts/snippets/redhat_register
kickstarts/snippets/sles_no_signature_checks
kickstarts/snippets/sles_register
kickstarts/snippets/sles_register_script
```

```
kickstarts/snippets/wait_for_networkmanager_script
kickstarts/upload/
kickstarts/wizard/

sent 324 bytes  received 1,063 bytes  2,774.00 bytes/sec
total size is 12,133  speedup is 8.75
receiving incremental file list
ssl-build/
ssl-build/RHN-ORG-PRIVATE-SSL-KEY
ssl-build/RHN-ORG-TRUSTED-SSL-CERT
ssl-build/index.txt
ssl-build/index.txt.attr
ssl-build/latest.txt
ssl-build/rhn-ca-openssl.cnf
ssl-build/rhn-ca-openssl.cnf.1
ssl-build/rhn-org-trusted-ssl-cert-1.0-1.noarch.rpm
ssl-build/rhn-org-trusted-ssl-cert-1.0-1.src.rpm
ssl-build/serial
ssl-build/d248/
ssl-build/d248/latest.txt
ssl-build/d248/rhn-org-httpd-ssl-archive-d248-1.0-1.tar
ssl-build/d248/rhn-org-httpd-ssl-key-pair-d248-1.0-1.noarch.rpm
ssl-build/d248/rhn-org-httpd-ssl-key-pair-d248-1.0-1.src.rpm
ssl-build/d248/rhn-server-openssl.cnf
ssl-build/d248/server.crt
ssl-build/d248/server.csr
ssl-build/d248/server.key
ssl-build/d248/server.pem
```

```
sent 380 bytes  received 50,377 bytes  101,514.00 bytes/sec
total size is 90,001  speedup is 1.77
SUSE Manager Database Control. Version 1.5.2
Copyright (c) 2012 by SUSE Linux Products GmbH
```

```
INFO: Database configuration has been changed.
INFO: Wrote new general configuration. Backup as /var/lib/pgsql/data/
postgresql.2016-01-12-11-29-42.conf
INFO: Wrote new client auth configuration. Backup as /var/lib/pgsql/data/
pg_hba.2016-01-12-11-29-42.conf
INFO: New configuration has been applied.
Database is online
System check finished
```

```
=====
Migration complete.
Please shut down the old SUSE Manager server now.
Reboot the new server and make sure it uses the same IP address and hostname
```

as the old SUSE Manager server!

IMPORTANT: Make sure, if applicable, that your external storage is mounted in the new server as well as the ISO images needed for distributions before rebooting the new server!

=====

12 Client Migration

Upgrading from SLE 12 with the latest service pack (SP) to SLE 15 can be automated, but requires some preparation steps.

To upgrade the SP version on SLE 12 (for example, upgrading from SLE 12, or SLE 12 SP1 or SP2, to SLE 12 SP3 or later) can be fully automated and requires no additional preparation.

12.1 Upgrading SLE 12 SPx to version 15

SLE 12 SPx clients can be auto-upgraded to SLE 15 with YaST auto-installation. This also applies for other supported products based on SLE 12.



Note: Supported Upgrade Paths

For generally supported SUSE Linux Enterprise upgrade paths, see https://www.suse.com/documentation/sles-15/book_sle_upgrade/data/sec_upgrade-paths_supported.html (SUSE Linux Enterprise Upgrade Guide, Chapter “Supported Upgrade Paths to SLE 15”). It is important that you migrate the client to the latest available SP first. In the case of SLE 12, this is currently SP3 (August 2018).

This procedure will work for traditionally managed systems (system type `management`). It is not currently available for systems using Salt (system type `salt`).

During the procedure, the machine reboots and performs the system upgrade. The process is controlled by YaST and AutoYaST, not by `zypper` commands.



Important

Only perform this migration procedure on client systems managed by Uyuni servers. For upgrading the Uyuni server itself, see *Chapter 11, Uyuni Server Migration*. This is a viable method for major version upgrades such as an upgrade from SUSE Linux Enterprise 12 to 15.

12.1.1 System Upgrade Preparation

Make sure your Uyuni and all the clients you want to upgrade have installed all available updates, including the Uyuni tools. This is absolutely necessary, otherwise the system upgrade will fail.

The preparation process contains several steps:

1. Download and save installation media
2. Create an auto-installation distribution
3. Create an activation key
4. Upload an AutoYaST profile

PROCEDURE: DOWNLOAD AND SAVE INSTALLATION MEDIA

1. On the Uyuni server, create a local directory for the SLE 15 installation media.
2. Download an ISO image with the installation sources, and mount the ISO image on your server:

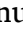
```
mkdir /mnt/sle15
mount -o loop DVD1.iso /mnt/sle15
```

Procedure: Create an Auto-Installation Distribution. For all distributions you want to upgrade, create a SLE 15 distribution in Uyuni.

1. In the Uyuni Web UI, click *Main Menu > Systems > Autoinstallation > Distributions*.
2. Enter a Distribution Label for your distribution (for example, autumn2018)
3. Specify the Tree Path, which is the root directory of the SLE 15 installation sources (for example, /mnt/sle15).
4. For Base Channel, use the update target distribution SLE-Product-SLES15-Pool for x86_64.
5. Confirm with Create Autoinstallable Distribution.

For more information about Autoinstallation, see *Book "Reference Manual", Chapter 8 "Autoinstallation"*.

Procedure: Create an Activation Key. In order to switch from the old SLE 12 SP3 base channel to the new SLE 15 channel, you need an activation key.

1. Go to *Main Menu* > *Systems* > *Activation Keys* and click Create Key.
2. Enter a description for your key.
3. Enter a key or leave it blank to generate an automatic key.
4. If you want to limit the usage, enter your value in the Usage text field.
5. Select the SLE-Product-SLES15-Pool for x86_64 base channel.
6. Decide about Add-On System Types. If in doubt, see https://www.suse.com/documentation/sles-15/book_quickstarts/data/art_modules.html  (SUSE Linux Enterprise Modules & Extensions Quick Start).
7. Click *Create Activation Key*.
8. Click the Child Channels tab and select the required channels. Finish with *Update Key*.

Procedure: Upload an AutoYaST Profile. Create an AutoYaST XML file according to [Section 12.1.2, "Sample Autoinstallation Script for System Upgrade \(SLES 12 SP3 to SLES 15\)"](#). For more information about AutoYaST, see .

1. Go to *Main Menu* > *Systems* > *Autoinstallation* and click Upload Kickstart/Autoyast File.
2. Paste the XML content in the text area or select the file to upload and click *Create*.
3. Add autoupgrade=1 in the Kernel Options of the Details tab and click *Update*.
4. Switch to the Variable tab.
5. In the text field registration_key= enter the key from the preparation above.
6. Click *Update Variables*.

After you have successfully finished this process, you are ready to perform the upgrade. For upgrade instruction, see [Warning: Synchronizing Target Channels](#).



Warning: Synchronizing Target Channels

Before successfully initializing the product migration, make sure that the migration target channels are completely mirrored. For the upgrade to SUSE Linux Enterprise 15, at least the SLE-Product-SLES15-Pool base channel with the SLE-Manager-Tools15-

Pool child channel for your architecture is required. The matching update channels such as SLE-Manager-Tools15-Updates and SLE-Product-SLES15-Updates are recommended. Watch the mirroring progress in /var/log/rhn/reposync/sles15-pool-x86_64.log.

1. Go to the system via *Main Menu > Systems* and click the name of the system. Then click *System Details > Provisioning > Autoinstallation > Schedule*, and choose the AutoYaST XML profile you have uploaded above.
2. Click *Schedule Autoinstallation and Finish*.

Next time the machine asks the Uyuni server for jobs, it will receive a reinstallation job which fetches kernel and initrd and writes a new /boot/grub/menu.lst (containing pointers to the new kernel and initrd).

When the machine boots, it will use the Grub configuration and boots the new kernel with its initrd. No PXE boot is required for this process. A shutdown of the machine is initiated as well, effectively 3 minutes after the job was fetched.

12.1.2 Sample Autoinstallation Script for System Upgrade (SLES 12 SP3 to SLES 15)

```
<?xml version="1.0"?>
<!DOCTYPE profile>
<profile xmlns="http://www.suse.com/1.0/yast2ns"
  xmlns:config="http://www.suse.com/1.0/configns">
  <general>
    $SNIPPET('spacewalk/sles_no_signature_checks')
    <mode>
      <confirm config:type="boolean">false</confirm>
    </mode>
  </general>
  <add-on>
    <add_on_products config:type="list">
      <listentry>
        <ask_on_error config:type="boolean">true</ask_on_error>
        <media_url>http://$redhat_management_server/ks/dist/child/sle15-updates-x86_64/
autumn2018</media_url>
        <name>SLE15 Updates</name>
        <product>SLE15</product>
        <product_dir>/</product_dir>
      </listentry>
    </add_on_products>
  </add-on>
</profile>
```



```

    <listentry>
      <ask_on_error config:type="boolean">true</ask_on_error>
      <media_url>http://$redhat_management_server/ks/dist/child/sle-manager-tools15-
pool-x86_64/summer2017</media_url>
      <name>SLE15 Manager Tools Pool</name>
      <product>SLE15</product>
      <product_dir></product_dir>
    </listentry>
    <listentry>
      <ask_on_error config:type="boolean">true</ask_on_error>
      <media_url>http://$redhat_management_server/ks/dist/child/sle-manager-tools15-
updates-x86_64/summer2017</media_url>
      <name>SLE15 Manager Tools Updates</name>
      <product>SLE15</product>
      <product_dir></product_dir>
    </listentry>
  </add_on_products>
</add-on>
<upgrade>
  <only_installed_packages config:type="boolean">>false</only_installed_packages>
  <stop_on_solver_conflict config:type="boolean">true</stop_on_solver_conflict>
</upgrade>
<backup>
  <sysconfig config:type="boolean">true</sysconfig>
  <modified config:type="boolean">true</modified>
  <remove_old config:type="boolean">>false</remove_old>
</backup>
<networking>
  <keep_install_network config:type="boolean">true</keep_install_network>
  <start_immediately config:type="boolean">true</start_immediately>
</networking>
<scripts>
  <init-scripts config:type="list">
    $SNIPPET('spacewalk/sles_register_script')
  </init-scripts>
</scripts>
</profile>

```

12.2 Migrating SLE 12 or later to version 12 SP3

Existing SLE 12 clients (SLE) may be upgraded to SP3 with the *SP Migration* procedure provided by the Web UI. The same applies for other supported products based on SUSE Linux Enterprise 12.



Warning: Synchronizing Target Channels

Before successfully initializing the product migration, you first must make sure that the migration target channels are completely mirrored. For the upgrade to SLE 12 SP3, at least the SLES12-SP3-Pool base channel with the SLE-Manager-Tools12-Pool child channel for your architecture is required. The matching update channels such as SLE-Manager-Tools12-Updates and SLES12-SP3-Updates are recommended.

PROCEDURE: MIGRATING SLE 12 CLIENT TO SP3

1. Direct your browser to the Uyuni Web UI where your client is registered, and login.
2. On the *Systems > All* page select your client system from the table.

f211.suse.de ?

Details Software Configuration Provisioning Groups Virtualization

Overview Properties Remote Command Connection Reactivation Hardware

System Status

Software Updates Available **Critical: 10** **Non-Critical: 9** **Packages: 29**

System Info

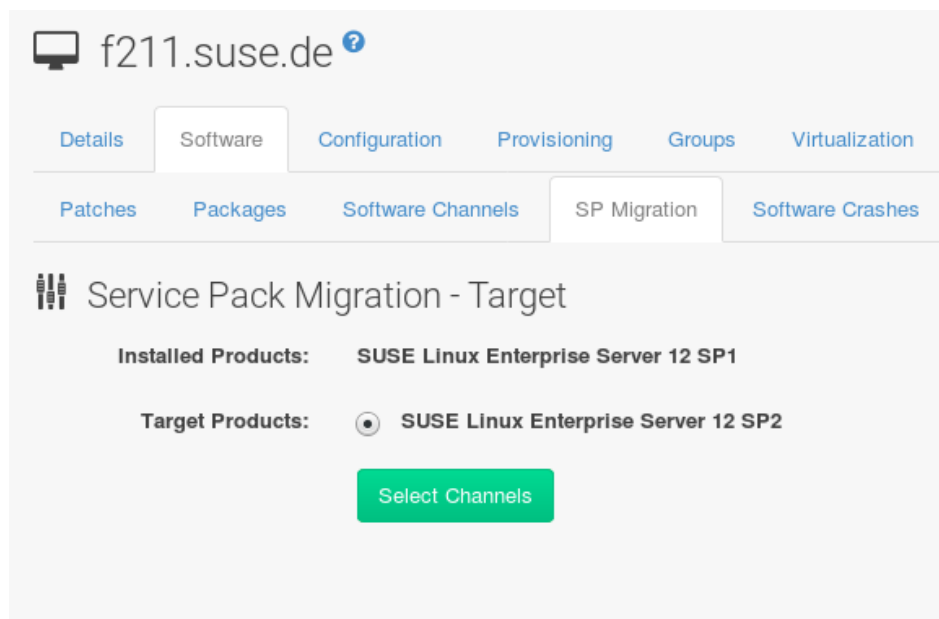
Hostname:	f211.suse.de
IP Address:	10.160.66.211
IPv6 Address:	::1
Kernel:	3.12.69-60.64.35-default
SUSE Manager System ID:	1000010005
Activation Key:	1-SLES-12-SP1
Installed Products:	SUSE Linux Enterprise Server 12 SP1
Lock Status:	System is unlocked (Lock system)

Subscribed Channels ([Alter Channel Subscriptions](#))

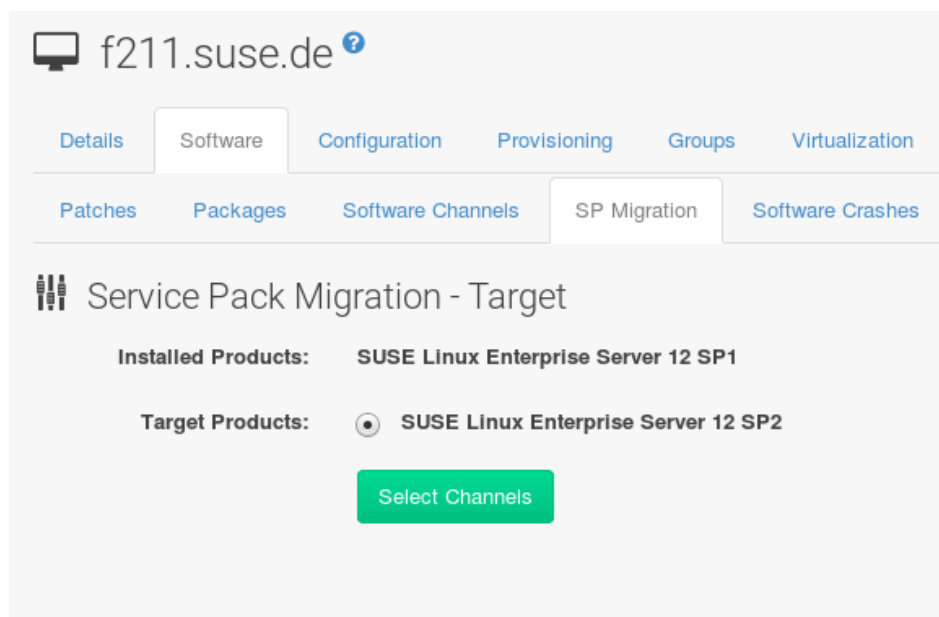
- [SLES12-SP1-Pool for x86_64](#)
- [SLE-Manager-Tools12-Pool x86_64 SP1](#)
- [SLE-Manager-Tools12-Updates x86_64 SP1](#)
- [SLES12-SP1-Updates for x86_64](#)

If there are *Software Updates Available* >] in the menu: *System Status*[notification install these updates first to avoid trouble during the migration process.



3. On the system's detail page select the *Software* >] tab > then the menu: *SP Migration*[tab.



4. From this tab you will see the installed products listed on your client. Select the wanted *Target Products* (if there is more than one), which is SUSE Linux Enterprise Server 12 SP3 .



Then confirm with *Select Channels* .

f211.suse.de   Delete System

Details Software Configuration Provisioning Groups Virtualization

Patches Packages Software Channels **SP Migration** Software Crashes

Service Pack Migration - Channels

Installed Products: SUSE Linux Enterprise Server 12 SP1

Target Products: SUSE Linux Enterprise Server 12 SP2

Target Base Channel:

- **Mandatory Child Channels:**
 - ☒ SLE-Manager-Tools12-Pool x86_64 SP2
 - ☒ SLE-Manager-Tools12-Updates x86_64 SP2
 - ☒ SLES12-SP2-Updates for x86_64
- **Optional Child Channels:**
 - ☐ SLE-Module-Adv-Systems-Management12-Pool for x86_64
 - ☐ SLE-Module-Adv-Systems-Management12-Updates for x86_64
 - ☐ SUSE-Manager-Proxy-3.0-Pool for x86_64 SP2
 - ☐ SUSE-Manager-Proxy-3.0-Updates for x86_64 SP2
 - ☐ SUSE-Manager-Proxy-3.1-Pool for x86_64
 - ☐ SUSE-Manager-Proxy-3.1-Updates for x86_64

[Schedule Migration](#)

5. Select *Schedule Migration* >] > and then menu: *Confirm* [.

f211.suse.de

Delete System | Add to SSM

Details

Software

Configuration

Provisioning

Groups

Virtualization

Audit

Events

Patches

Packages

Software Channels

SP Migration

Software Crashes

Service Pack Migration - Confirm

Please confirm below to schedule the migration of this system to the following products:

SUSE Linux Enterprise Server 12 SP2

Channel subscriptions after the migration:

SLES12-SP2-Pool for x86_64

SLE-Manager-Tools12-Pool x86_64 SP2

SLE-Manager-Tools12-Updates x86_64 SP2

SLES12-SP2-Updates for x86_64

Schedule action for no sooner than:

5/3/17

10:42 am

CEST

In order to detect any possible problems it is recommended to always do a **Dry Run** before scheduling the actual Service Pack Migration.

Go Back

Dry Run

Confirm

Check the *System Status* on the system's details when the migration is done.

Details

Software

Configuration

Provisioning

Groups

Virt

Overview

Properties

Remote Command

Connection

Reactive

System Status

 Software Updates Available **Critical:** 6 **Non-Critical:** 5 **Packages:** 26

System Info

Hostname:	f211.suse.de
IP Address:	10.160.66.211
IPv6 Address:	::1
Kernel:	4.4.49-92.14-default
SUSE Manager System ID:	1000010005
Activation Key:	1-SLES-12-SP1

Installed Products:

SUSE Linux Enterprise Server 12 SP2

Lock Status:  System is **unlocked**
([Lock system](#))

Subscribed Channels ([Alter Channel Subscriptions](#))

- [SLES12-SP2-Pool for x86_64](#)
- [SLE-Manager-Tools12-Pool x86_64 SP2](#)
- [SLE-Manager-Tools12-Updates x86_64 SP2](#)
- [SLES12-SP2-Updates for x86_64](#)

If the *System Status > J* notification does not report a successful migration but lists menu: *Software Updates Available* , install the update now and then check again.

Finally, consider to schedule a reboot.

13 PostgreSQL Database Migration

Uyuni 3 uses postgresql database version 9.4. Postgresql version 9.6 has been officially released for SUSE Linux Enterprise Server 12 SP3. In the near future postgresql 9.6 will become the base version provided by Uyuni. Currently version 9.4 is hardcoded into SUSE Manager, therefore when installing Uyuni it will explicitly use this version. This chapter provides guidance on migrating an existing 9.4 database to 9.6 on your Uyuni Server.

13.1 New Uyuni Installations

Once support for postgresql version 9.6 has been officially released for SUSE Manager, no action will be required for new installations. The Uyuni extension will pick up the latest version during installation on SLES 12 SP3. This will be fully transparent to the user. Check the active postgresql version with the following command:

```
suse-manager-example-srv:~ # psql --version
psql (PostgreSQL) 9.6.3
```

13.2 Migrating an Existing Installation

Before migrating to the new database version, ensure SUSE Manager is fully patched to the latest version. You can check if the system is ready to use postgresql 9.6 by issuing the following command:

```
suma-test-srv:~ # rpm -q smdba
smdba-1.5.8-0.2.3.1.x86_64
```



Important

Postgresql 9.6 requires smdba version 1.5.8 or higher



Warning

Always create a database backup before performing a migration

. The database migration begins by executing the following command:

```
$> /usr/lib/susemanager/bin/pg-migrate.sh
```

The `pg-migrate.sh` script will automatically perform the following operations:

- Stop spacewalk services
- Shut down the running database
- Check if postgresql 9.6 is installed and install it if not already present
- Switch from postgresql 9.4 to postgresql 9.6 as the new default
- Initiates the database migration
- Creates a postgresql configuration file tuned for use by SUSE Manager (The reason for the latest version of smdba)
- Start both the database and spacewalk services



Note

Please note that during the migration the data directory of the database is copied for use by the postgresql 9.6. This results in temporarily doubling the amount of required disk space. In case of a failure, the migration script will attempt a restore to the original state. After a successful migration, you may safely delete the old database directory (renamed to `/var/lib/pgsql/data-pg94`) to reclaim lost disk space.

13.3 Performing a Fast Migration

There are two negative aspects to performing a regular migration:

- You temporarily need double the disk space under `/var/lib/pgsql`
- Depending on the size of the database the migration can take up some time because the whole data directory needs to be copied.

It is possible however to perform a fast migration. In this case you do not need the additional disk space as the database files will not be copied but hard linked. This also has the natural effect of greatly increasing the speed of the migration process. The entire migration could be completed in less than one minute.



Warning

Keep in mind if a fast migration fails, database restoration will only be possible with a database backup. Only perform a fast migration if you have an available database backup.

Perform a fast migration with the following command (Ensure you have a database backup):

```
$> /usr/lib/susemanager/bin/pg-migrate.sh fast
```

13.4 Typical Migration Sample Session

A slow migration should provide you with the following output:

```
d235:~ # /usr/lib/susemanager/bin/pg-migrate.sh
15:58:00 Shut down spacewalk services...
Shutting down spacewalk services...
Done.
15:58:03 Checking postgresql version...
15:58:03 Installing postgresql 9.6...
Dienst 'SUSE_Linux_Enterprise_Server_12_SP2_x86_64' wird aktualisiert.
Dienst 'SUSE_Manager_Server_3.1_x86_64' wird aktualisiert.
Repository-Daten werden geladen...
Installierte Pakete werden gelesen...
Paketabhängigkeiten werden aufgelöst...

Die folgenden 3 NEUEN Pakete werden installiert:
  postgresql96 postgresql96-contrib postgresql96-server

3 neue Pakete zu installieren.
Gesamtgröße des Downloads: 5,7 MiB. Bereits im Cache gespeichert: 0 B. Nach der Operation
werden zusätzlich 25,3 MiB belegt.
Fortfahren? [j/n/...? zeigt alle Optionen] (j): j
Paket postgresql96-9.6.3-2.4.x86_64 abrufen (1/3), 1,3 MiB ( 5,1 MiB entpackt)
Abrufen: postgresql96-9.6.3-2.4.x86_64.rpm [fertig]
Paket postgresql96-server-9.6.3-2.4.x86_64 abrufen (2/3), 3,7 MiB ( 17,9 MiB entpackt)
Abrufen: postgresql96-server-9.6.3-2.4.x86_64.rpm [.fertig]
Paket postgresql96-contrib-9.6.3-2.4.x86_64 abrufen (3/3), 648,9 KiB ( 2,2 MiB entpackt)
Abrufen: postgresql96-contrib-9.6.3-2.4.x86_64.rpm [fertig]
Überprüfung auf Dateikonflikte läuft: [.....fertig]
(1/3) Installieren: postgresql96-9.6.3-2.4.x86_64 [.....fertig]
(2/3) Installieren: postgresql96-server-9.6.3-2.4.x86_64 [.....fertig]
(3/3) Installieren: postgresql96-contrib-9.6.3-2.4.x86_64 [.....fertig]
15:58:08 Ensure postgresql 9.6 is being used as default...
```

```
15:58:09 Successfully switched to new postgresql version 9.6.
15:58:09 Create new database directory...
15:58:09 Initialize new postgresql 9.6 database...
The files belonging to this database system will be owned by user "postgres".
This user must also own the server process.
```

The database cluster will be initialized with locale "en_US.UTF-8".
The default database encoding has accordingly been set to "UTF8".
The default text search configuration will be set to "english".

Data page checksums are disabled.

```
fixing permissions on existing directory /var/lib/pgsql/data ... ok
creating subdirectories ... ok
selecting default max_connections ... 100
selecting default shared_buffers ... 128MB
selecting dynamic shared memory implementation ... posix
creating configuration files ... ok
running bootstrap script ... ok
performing post-bootstrap initialization ... ok
syncing data to disk ... ok
```

WARNING: enabling "trust" authentication for local connections
You can change this by editing pg_hba.conf or using the option -A, or
--auth-local and --auth-host, the next time you run initdb.

Success. You can now start the database server using:

```
pg_ctl -D /var/lib/pgsql/data -l logfile start
```

```
15:58:12 Successfully initialized new postgresql 9.6 database.
15:58:12 Upgrade database to new version postgresql 9.6...
Performing Consistency Checks
```

```
-----
Checking cluster versions                                ok
Checking database user is the install user              ok
Checking database connection settings                   ok
Checking for prepared transactions                      ok
Checking for reg* system OID user data types           ok
Checking for contrib/isn with bigint-passing mismatch  ok
Checking for roles starting with 'pg_'                 ok
Creating dump of global objects                         ok
Creating dump of database schemas
  postgres
  susemanager
  templatel
                                                    ok
```

```

Checking for presence of required libraries          ok
Checking database user is the install user          ok
Checking for prepared transactions                  ok

If pg_upgrade fails after this point, you must re-initdb the
new cluster before continuing.

Performing Upgrade
-----
Analyzing all rows in the new cluster                ok
Freezing all rows on the new cluster                ok
Deleting files from new pg_clog                     ok
Copying old pg_clog to new server                   ok
Setting next transaction ID and epoch for new cluster ok
Deleting files from new pg_multixact/offsets         ok
Copying old pg_multixact/offsets to new server       ok
Deleting files from new pg_multixact/members         ok
Copying old pg_multixact/members to new server       ok
Setting next multixact ID and offset for new cluster ok
Resetting WAL archives                             ok
Setting frozenxid and minmxid counters in new cluster ok
Restoring global objects in the new cluster          ok
Restoring database schemas in the new cluster
    postgres
    susemanager
    template1
                                                    ok

Copying user relation files
    /var/lib/pgsql/data-pg94/base/12753/12710

[...]

    /var/lib/pgsql/data-pg94/base/1/12574
                                                    ok

Setting next OID for new cluster                    ok
Sync data directory to disk                        ok
Creating script to analyze new cluster               ok
Creating script to delete old cluster                ok

Upgrade Complete
-----
Optimizer statistics are not transferred by pg_upgrade so,
once you start the new server, consider running:
    ./analyze_new_cluster.sh

Running this script will delete the old cluster's data files:
    ./delete_old_cluster.sh

```

```
15:58:51 Successfully upgraded database to postgresql 9.6.
15:58:51 Tune new postgresql configuration...
INFO: Database configuration has been changed.
INFO: Wrote new general configuration. Backup as /var/lib/pgsql/data/
postgresql.2017-07-26-15-58-51.conf
INFO: Wrote new client auth configuration. Backup as /var/lib/pgsql/data/
pg_hba.2017-07-26-15-58-51.conf
INFO: Configuration has been changed, but your database is right now offline.
Database is offline
System check finished
15:58:51 Successfully tuned new postgresql configuration.
15:58:51 Starting spacewalk services...
Starting spacewalk services...
Done.
```

14 Backup and Restore

Back up your Uyuni installation regularly, in order to prevent data loss. Because Uyuni relies on a database as well as the installed program and configurations, it is important to back up all components of your installation. This chapter contains information on the files you need to back up, and introduces the `smdba` tool to manage database backups. It also contains information about restoring from your backups in the case of a system failure.



Important: Backup Space Requirements

Regardless of the backup method you use, you must have available at least three times the amount of space your current installation uses. Running out of space can result in backups failing, so check this often.

14.1 Backing up Uyuni

The most comprehensive method for backing up your Uyuni installation is to back up the relevant files and directories. This can save you time in administering your backup, and can be faster to reinstall and re-synchronize in the case of failure. However, this method requires significant disk space and could take a long time to perform the backup.

If you want to only back up the required files and directories, use this list:



Note

To make this process simpler, and more comprehensive, we recommend backing up the entire `/etc` and `/root` directories, not just the ones specified here.

- `/etc/cobbler/`
- `/etc/dhcp.conf`
- `/etc/fstab` and any ISO mountpoints you require.
- `/etc/rhn/`
- `/etc/salt`
- `/etc/sudoers`

- /etc/sysconfig/rhn/
- /root/.gnupg/
- /root/.ssh
 - This file exists if you are using an SSH tunnel or SSH push. You will also need to have saved a copy of the id-susemanager key.
- /root/ssl-build/
- /srv/pillar
- /srv/salt
- /srv/tftpboot/
- /srv/www/cobbler
- /srv/www/htdocs/pub/
- /var/lib/cobbler/
- /var/lib/rhn/kickstarts/
- /var/pacewalk/
- Any directories containing custom data (such as scripts, Kickstart profiles, AutoYaST, and custom RPMs).

Important

You will also need to back up your database, which you can do by copying /var/pacewalk/db-backup or by using the smdba tool, which is explained later in this chapter.

PROCEDURE: RESTORE FROM A MANUAL BACKUP

1. Re-install Uyuni
2. Re-synchronize your installation with the mgr-sync tool.
3. You can choose to re-register your product, or skip the registration and SSL certificate generation sections.

4. Re-install the `/root/ssl-build/rhn-org-httpd-ssl-key-pair-MACHINE_NAME-VER-REL.noarch.rpm` package.

5. Schedule the re-creation of search indexes next time the `rhn-search` service is started:

```
rcrhn-search cleanindex
```

6. If you did not have `/var/spacewalk/packages/` in your backup, but the source repository still exists, you can restore it by performing a complete channel synchronization.

14.2 Administering the Database with `smdba`

The `smdba` tool is used for managing a local PostgreSQL database. It allows you to back up and restore your database, and manage backups. It can also be used to check the status of your database, and perform administration tasks, such as restarting.



Note

The `smdba` tool works with local PostgreSQL databases only, it will not work with remotely accessed databases, or Oracle databases.



Important

The `smdba` tool requires `sudo` access, in order to execute system changes. Ensure you have enabled `sudo` access for the `admin` user before you begin, by checking the `/etc/sudoers` file for this line:

```
admin    ALL=(postgres) /usr/bin/smdba
```

Check the runtime status of your database with the `smdba db-status` command. This command will return either `online` or `offline`:

```
smdba db-status
Checking database core...      online
```

To check the full connection to the database, use the `smdba db-check` command. Depending on your environment, this command will report on the status of listeners, in addition to connectivity status.

```
smdba db-check
```

Starting and stopping the database can be performed with `smdba db-start` and `smdba db-stop`

```
smdba db-start
Starting core...      done
```

```
smdba db-stop
Stopping the SUSE Manager database...
Stopping core:        done
```

14.3 Database Backup with `smdba`

The `smdba` tool performs a continuous archiving backup. This backup method combines a log of every change made to the database during the current session, with a series of more traditional backup files. When a crash occurs, the database state is first restored from the most recent backup file on disk, then the log of the current session is replayed exactly, to bring the database back to a current state. A continuous archiving backup with `smdba` is performed with the database running, so there is no need for downtime.

This method of backing up is stable and generally creates consistent snapshots, however it can take up a lot of storage space. Ensure you have at least three times the current database size of space available for backups. You can check your current database size by navigating to `/var/lib/pgsql/` and running `df -h`.

The `smdba` tool also manages your archives, keeping only the most recent backup, and the current archive of logs. The log files can only be a maximum file size of 16 MB, so a new log file will be created once the files reach this size. Every time you create a new backup, previous backups will be purged to release disk space. We recommend you use `cron` to schedule your `smdba` backups to ensure that your storage is managed effectively, and you always have a backup ready in case of failure.

14.3.1 Performing a Manual Database Backup

The `smdba` tool can be run directly from the command line. We recommend you run a manual database backup immediately after installation, or if you have made any significant changes to your configuration.



Note

When `smdba` is run for the first time, or if you have changed the location of the backup, it will need to restart your database before performing the archive. This will result in a small amount of downtime. Note that regular database backups will not require any downtime.

PROCEDURE: PERFORMING A MANUAL DATABASE BACKUP

1. Allocate permanent storage space for your backup. In this procedure, we will be using an NFS share located at `/var/spacewalk/`. This will become a permanent target for your backup, so ensure it will remain accessible by your server at all times.

2. In your backup location, create a directory for the backup:

```
sudo -u postgres mkdir /var/spacewalk/db-backup
```

Or, as root:

```
install -d -o postgres /var/spacewalk/db-backup
```

3. Ensure you have the correct permissions set on the backup location:

```
chown postgres:postgres /var/spacewalk/db-backup
```

4. To run a backup for the first time, run the `smdba backup-hot` command with the `enable` option set. This will create the backup in the specified directory, and, if necessary, restart the database:

```
smdba backup-hot --enable=on --backup-dir=/var/spacewalk/db-backup
```

5. Check that the backup files exist in the `/mnt/backup/database` directory, to ensure that your backup has been successful.

14.3.2 Scheduling Automatic Backups

You do not need to shut down your system in order to perform a database backup with `smdba`. However, because it is a large operation, database performance can slow down while the backup is running. We recommend you schedule regular database backups for a low-traffic period, to minimize disruption.

Important

Ensure you have at least three times the current database size of space available for backups. You can check your current database size by navigating to /var/lib/pgsql/ and running df -h.

PROCEDURE: SCHEDULING AUTOMATIC BACKUPS

1. Create a directory for the backup, and set the appropriate permissions:

```
# mkdir /var/spacewalk/db-backup
# chown -R postgres:postgres /var/spacewalk/db-backup
# chmod 700 /var/spacewalk/db-backup
```

2. Open /etc/cron.d/db-backup-mgr, or create it if it doesn't exist, and add the following line to create the cron job:

```
0 2 * * * root /usr/bin/smdba backup-hot --enable=on --backup-dir=/var/spacewalk/db-backup
```

3. Check the backup directory regularly to ensure the backups are working as expected.

14.4 Restoring from Backup

The smdba tool can be used to restore from backup in the case of failure.

PROCEDURE: RESTORING FROM BACKUP

1. Shutdown the database:

```
smdba db-stop
```

2. Start the restore process and wait for it to complete:

```
smdba backup-restore start
```

3. Restart the database:

```
smdba db-start
```

4. Check if there are differences between the RPMs and the database.

```
spacewalk-data-fsck
```

14.5 Archive Log Settings

In SUSE Manager with an embedded database, archive logging is enabled by default. This feature allows the database management tool `smdba` to perform hot backups.

With archive log enabled, even more data is stored on the hard disk:

- PostgreSQL maintains a limited number of archive logs. Using the default configuration, approx. 64 files with a size of 16 MiB are stored.

Creating a user and syncing the channels:

- SLES12-SP2-Pool-x86_64
- SLES12-SP2-Updates-x86_64
- SLE-Manager-Tools12-Pool-x86_64-SP2
- SLE-Manager-Tools12-Updates-x86_64-SP2

Postgresql will generate an additional ~1 GB of data. So it is important to think about a backup strategy and create a backups in a regular way.

Archive logs are stored at:

- /var/lib/pgsql/data/pg_xlog/ (postgresql)

14.6 Retrieving an Overview of Occupied Database Space

Database administrators may use the subcommand `space-overview` to get a report about occupied table spaces, for example:

```
smdba space-overview
SUSE Manager Database Control. Version 1.5.2
Copyright (c) 2012 by SUSE Linux Products GmbH
```

Tablespace	Size (Mb)	Avail (Mb)	Use %
postgres	7	49168	0.013
susemanager	776	48399	1.602

The following command is available for PostgreSQL. For a more detailed report, use the `space-tables` subcommand. It lists the table and its size, for example:

```
smdba space-tables
SUSE Manager Database Control. Version 1.5.2
Copyright (c) 2012 by SUSE Linux Products GmbH

Table | Size
-----+-----
public.all_primary_keys | 0 bytes
public.all_tab_columns | 0 bytes
public.allserverkeywordsincereboot | 0 bytes
public.dblink_pkey_results | 0 bytes
public.dual | 8192 bytes
public.evr_t | 0 bytes
public.log | 32 kB
...
```

14.7 Moving the Database

It is possible to move the database to another location. For example if your database storage space is running low. The following procedure will guide you through moving the database to a new location for use by SUSE Manager.

PROCEDURE: MOVING THE DATABASE

1. The default storage location for SUSE Manager is: `/var/lib/pgsql/`. You would like to move it, for example to: `/storage/postgres/`. To begin, stop the running database with:

```
# rcpostgresql stop
```

Shutdown running spacewalk services with:

```
# spacewalk-service stop
```

2. Copy the current working directory structure with the following syntax:

```
cp [OPTION]... SOURCE... DIRECTORY
```

using the `-a, --archive` option. For example:

```
# cp -ar /var/lib/pgsql/ /storage/postgres/
```

This command will copy the contents of /var/lib/pgsql/ to /storage/postgres/pgsql/ .



Important

The contents of the /var/lib/pgsql needs to remain the same or the SUSE Manager database may malfunction. You also should ensure there is enough available disk space.

3. Mount the new database directory with:

```
# mount /storage/postgres/pgsql
```

4. Make sure ownership is postgres:postgres and not root:root by changing to the new directory and running the following command:

```
/var/lib/pgsql/ # cd /storage/postgres/pgsql/  
/storage/postgres/pgsql/ # l  
total 8  
drwxr-x--- 4 postgres postgres 47 Jun 2 14:35 ./
```

5. Add the new database mount location to your servers fstab by editing etc/fstab .
6. Start the database with:

```
# rcpostgresql start
```

Start spacewalk-services with:

```
# spacewalk-service start
```

14.8 Recovering from a Crashed Root Partition

This section provides guidance on restoring your server after its root partition has crashed. This section assumes you have setup your server similar to the procedure explained in Getting Started guide with separate partitions for the database and for channels mounted at /var/lib/pgsql and /var/spacewalk/ .

PROCEDURE: RECOVERING FROM A CRASHED ROOT PARTITION

1. Start by installing SLES12 SP2 and the SUSE Manager Extension. Do not mount the /var/spacewalk and /var/lib/pgsql partitions.
2. Once installation of SUSE Manager has completed shutdown services with spacewalk-service shutdown and the database with rcpostgresql stop.
3. Mount your /var/spacewalk and /var/lib/pgsql partitions and restore the directories listed in section one.
4. Start SUSE Manager services and the database with spacewalk-services start and rcpostgresql start
5. SUSE Manager should now operate normally without loss of your database or synced channels.

14.9 Database Connection Information

The information for connecting to the SUSE Manager database is located in /etc/rhn/rhn.conf :

```
db_backend = postgresql
db_user = susemanager
db_password = susemanager
db_name = susemanager
db_host = localhost
db_port = 5432
db_ssl_enabled =
```


15 Authentication Methods

15.1 Authentication Via PAM

As security measures become increasingly complex, SUSE Manager supports network-based authentication systems via Pluggable Authentication Modules (PAM). PAM is a suite of libraries that allows to integrate SUSE Manager with a centralized authentication mechanism, thus eliminating the need to remember multiple passwords. SUSE Manager supports LDAP, Kerberos, and other network-based authentication systems via PAM. To enable SUSE Manager to use PAM in your organization's authentication infrastructure, follow the steps below.

1. Set up a PAM service file (default location: /etc/pam.d/susemanager) then enforce its use by adding the following line to /etc/rhn/rhn.conf :

```
pam_auth_service = susemanager
```



Note

This assumes the PAM service file is named susemanager.

2. To enable a new or existing user to authenticate with PAM, proceed to the *Create User* page and select the checkbox labeled Pluggable Authentication Modules (PAM) positioned below the password and password confirmation fields.
3. To authenticate a SLES system against Kerberos add the following lines to /etc/pam.d/susemanager :

```
##PAM-1.0
auth      include      common-auth
account   include      common-account
password  include      common-password
session   include      common-session
```



Note

To register a Red Hat Enterprise Linux System against Kerberos add the following lines to /etc/pam.d/susemanager

```
#%PAM-1.0
auth      required      pam_env.so
auth      sufficient     pam_krb5.so no_user_check
auth      required      pam_deny.so
account   required      pam_krb5.so no_user_check
```

+

1. YaST can now be used to configure PAM, when packages such as `yast2-ldap-client` and `yast2-kerberos-client` are installed; for detailed information on configuring PAM, see the SUSE Linux Enterprise Server Security Guide https://www.suse.com/documentation/sles-12/book_security/data/part_auth.html. This example is not limited to Kerberos; it is generic and uses the current server configuration. Note that only network based authentication services are supported.

Important: Changing Passwords

Changing the password on the SUSE Manager Web interface changes only the local password on the SUSE Manager server. But this password may not be used at all if PAM is enabled for that user. In the above example, for instance, the Kerberos password will not be changed.

15.2 Authentication Via eDirectory and PAM

1. First check to ensure eDirectory authentication is working with your current OS for example:

```
#getent passwd
```

2. If users are returned from eDirectory then create the following file:

```
# cat /etc/pam.d/susemanager
```

3. And add the following content:

```
#%PAM-1.0
auth      include        common-auth
account   include        common-account
password  include        common-password
session   include        common-session
```

```
#
```

4. Finally add the following lines to the Uyuni conf file:

```
# grep -i pam /etc/rhn/rhn.conf  
pam_auth_service = susemanager
```

5. You may now create users with the same id that appears on eDirectory and mark the Use PAM check-box from the SUSE Manager WebUI.

15.3 Example Quest VAS Active Directory Authentication Template

If you are using Quest VAS for active directory authentication, you can use the following /etc/pam.d/susemanager file.

```
##PAM-1.0  
auth      required      pam_env.so  
auth      sufficient     pam_vas3.so no_user_check  
auth      requisite      pam_vas3.so echo_return  
auth      required       pam_deny.so  
account    required       pam_vas3.so no_user_check
```

16 Using a Custom SSL Certificate

The following section will guide you through using a custom certificate with Uyuni 3.2 and SUSE Manager Proxy 3.2 .

16.1 Prerequisites

The following list provides requirements for using a custom certificate.

- A Certificate Authority (CA) SSL public certificate file
- A Web server SSL private key file
- A Web server SSL public certificate file
- Key and Certificate files must be in PEM format



Important: Hostname and SSL Keys

The hostname of the web server's SSL keys and relevant certificate files must match the hostname of the machine which they will be deployed on.



Tip: Intermediate Certificates

In case you want to use CAs with intermediate certificates, merge the intermediate and root CA certificates into one file. It is important that the intermediate certificate comes first within the combined file.

16.2 Setup

After completing YaST firstboot procedures, export your current environment variables and point them to the correct SSL files to be imported. Running these commands will make the default certificate obsolete after executing the `yast2 susemanagersetup` command. For more information on YaST firstboot, see https://www.suse.com/documentation/suse-manager-3/singlehtml/suse_manager21/book_susemanager_install/book_susemanager_install.html#sec.manager.inst.setup.

1. Export the environment variables and point to the SSL files to be imported:

```
export CA_CERT=`path_to_CA_certificate_file`export
SERVER_KEY=`path_to_web_server_key`export
SERVER_CERT=`path_to_web_server_certificate`
```

2. Execute Uyuni setup with

```
yast2 susemanagersetup
```

Proceed with the default setup. Upon reaching the Certificate Setup window during YaST installation, fill in random values, as these will be overridden with the values specified in *???TITLE???*.



Note: Shell Requirements

Make sure that you execute `yast2 susemanagersetup` from within the same shell the environment variables were exported from.

16.3 Using a Custom Certificate with SUSE Manager Proxy

After completing the installation with `yast` found in *Book "Advanced Topics", Chapter 2 "Uyuni 3.2 Proxy"* continue with a modified procedure:

1. Execute `configure-proxy.sh`.
2. When prompted with:

```
Do you want to import existing certificates?
```

Answer with ☐ y .

3. Continue by following the script prompts.

17 Troubleshooting

This chapter provides guidance on registering cloned systems with SUSE Manager. This includes both Salt and Traditional clients. For more information, see <https://www.novell.com/support/kb/doc.php?id=7012170>.

17.1 Registering Cloned Salt Minions

PROCEDURE: REGISTERING A CLONED SALT MINION WITH SUSE MANAGER

1. Clone your system (for example using the existing cloning mechanism of your favorite Hypervisor)



Note: Quick Tips

Each step in this section is performed on the cloned system, this procedure does not manipulate the original system, which will still be registered to SUSE Manager. The cloned virtual machine should have a different UUID from the original (this UUID is generated by your hypervisor) or SUSE Manager will overwrite the original system data with the new one.

2. Make sure your machines have different hostnames and IP addresses, also check that `/etc/hosts` contains the changes you made and the correct host entries.

The next step you take will depend on the Operating System of the clone.

The following scenario can occur after on-boarding cloned Salt minions. If after accepting all cloned minion keys from the onboarding page and you see only one minion on the System Overview page, this is likely due to these machines being clones of the original and using a duplicate machine-id. Perform the following steps to resolve this conflict based on OS.

PROCEDURE: SLES 12 REGISTERING SALT CLONES

1. SLES 12: If your machines have the same machine ids then delete the file on each minion and recreate it:

```
# rm /etc/machine-id
# rm /var/lib/dbus/machine-id
# dbus-uuidgen --ensure
```

```
# systemd-machine-id-setup
```

PROCEDURE: SLES 11 REGISTERING SALT CLONES

1. SLES 11: As there is no systemd machine id, generate one from dbus:

```
# rm /var/lib/dbus/machine-id  
# dbus-uuidgen --ensure
```

If your machines still have the same minion id then delete the minion_id file on each minion (FQDN will be used when it is regenerated on minion restart):

```
# rm /etc/salt/minion_id
```

Finally delete accepted keys from Onboarding page and system profile from SUSE Manager, and restart the minion with:

```
# systemctl restart salt-minion
```

You should be able to re-register them again, but each minion will use a different '/etc/machine-id' and should now be correctly displayed on the System Overview page.

17.2 Registering Cloned Traditional Systems

This section provides guidance on troubleshooting cloned traditional systems registered via bootstrap.

PROCEDURE: REGISTERING A CLONED SYSTEM WITH SUSE MANAGER (TRADITIONAL SYSTEMS)

1. Clone your system (using your favorite hypervisor.)



Note: Quick Tips

Each step in this section is performed on the cloned system, this procedure does not manipulate the original system, which will still be registered to SUSE Manager. The cloned virtual machine should have a different UUID from the original (this UUID is generated by your hypervisor) or SUSE Manager will overwrite the original system data with the new one.

2. Change the Hostname and IP addresses, also make sure /etc/hosts contains the changes you made and the correct host entries.

3. Stop rhnsd daemon with:

```
# /etc/init.d/rhnsd stop
```

or alternatively:

```
# rcrhnsd stop
```

4. Stop osad with:

```
# /etc/init.d/osad stop
```

or alternatively:

```
# rcosad stop
```

5. Remove the osad authentication configuration file and the systemid with:

```
# rm -f /etc/sysconfig/rhn/{osad-auth.conf,systemid}
```

The next step you take will depend on the Operating System of the clone.

PROCEDURE: SLES 12 REGISTERING A CLONED TRADITIONAL SYSTEM

1. If your machines have the same machine ids then delete the file on each client and recreate it:

```
# rm /etc/machine-id  
# rm /var/lib/dbus/machine-id  
# dbus-uuidgen --ensure  
# systemd-machine-id-setup
```

2. Remove the following credential files:

```
# rm -f /etc/zypp/credentials.d/{SCCcredentials,NCCcredentials}
```

3. Re-run the bootstrap script. You should now see the cloned system in SUSE Manager without overwriting the system it was cloned from.

PROCEDURE: SLES 11 REGISTERING A CLONED TRADITIONAL SYSTEM

1. Continued from section 1 step 5:

```
# suse_register -E
```


(--erase-local-regdata, Erase all local files created from a previous executed registration. This option make the system look like never registered)

2. Re-run the bootstrap script. You should now see the cloned system in SUSE Manager without overwriting the system it was cloned from.

PROCEDURE: SLES 10 REGISTERING A CLONED TRADITIONAL SYSTEM

1. Continued from section 1 step 5:

```
# rm -rf /etc/{zmd,zypp}
```

- 2.

```
# iiii everthing in /var/lib/zypp/ except /var/lib/zypp/db/products/ !!!!!  
# check whether this command works for you  
# rm -rf /var/lib/zypp/!(db)
```

- 3.

```
# rm -rf /var/lib/zmd/
```

4. Re-run the bootstrap script. You should now see the cloned system in SUSE Manager without overwriting the system it was cloned from.

PROCEDURE: RHEL 5,6 AND 7

1. Continued from section 1 step 5:

```
# rm -f /etc/NCCcredentials
```

2. Re-run the bootstrap script. You should now see the cloned system in SUSE Manager without overwriting the system it was cloned from.

17.3 Typical OSAD/jabberd Challenges

This section provides answers for typical issues regarding OSAD and jabberd.

17.3.1 Open File Count Exceeded

SYMPTOMS: OSAD clients cannot contact the SUSE Manager Server, and jabberd requires long periods of time to respond on port 5222.

CAUSE: The number of maximum files that a jabber user can open is lower than the number of connected clients. Each client requires one permanently open TCP connection and each connection requires one file handler. The result is jabberd begins to queue and refuse connections.

CURE: Edit the `/etc/security/limits.conf` to something similar to the following: `jabber-softnofile<#clients + 100> jabberhardnofile<#clients + 1000>`

This will vary according to your setup. For example in the case of 5000 clients: `jabbersoft-nofile5100 jabberhardnofile6000`

Ensure you update the `/etc/jabberd/c2s.xml` `max_fds` parameter as well. For example: `<max_fds>6000</max_fds>`

EXPLANATION: The soft file limit is the limit of the maximum number of open files for a single process. In SUSE Manager the highest consuming process is c2s, which opens a connection per client. 100 additional files are added, here, to accommodate for any non-connection file that c2s requires to work correctly. The hard limit applies to all processes belonging to the jabber user, and accounts for open files from the router, s2s and sm processes additionally.

17.3.2 jabberd Database Corruption

SYMPTOMS: After a disk is full error or a disk crash event, the jabberd database may have become corrupted. jabberd may then fail to start during spacewalk-service start:

```
Starting spacewalk services...
  Initializing jabberd processes...
    Starting router
done
  Starting sm startproc: exit status of parent of /usr/bin/sm: 2
failed
  Terminating jabberd processes...
```

`/var/log/messages` shows more details:

```
jabberd/sm[31445]: starting up
jabberd/sm[31445]: process id is 31445, written to /var/lib/jabberd/pid/sm.pid
jabberd/sm[31445]: loading 'db' storage module
jabberd/sm[31445]: db: corruption detected! close all jabberd processes and run
db_recover
jabberd/router[31437]: shutting down
```

CURE: Remove the jabberd database and restart. Jabberd will automatically re-create the database:

```
spacewalk-service stop
rm -Rf /var/lib/jabberd/db/*
spacewalk-service start
```

An alternative approach would be to test another database, but SUSE Manager does not deliver drivers for this:

```
rcosa-dispatcher stop
rcjabberd stop
cd /var/lib/jabberd/db
rm *
cp /usr/share/doc/packages/jabberd/db-setup.sqlite .
sqlite3 sqlite.db < db-setup.sqlite
chown jabber:jabber *
rcjabberd start
rcosa-dispatcher start
```

17.3.3 Capturing XMPP Network Data for Debugging Purposes

If you are experiencing bugs regarding OSAD, it can be useful to dump network messages in order to help with debugging. The following procedures provide information on capturing data from both the client and server side.

PROCEDURE: SERVER SIDE CAPTURE

1. Install the tcpdump package on the SUSE Manager Server as root: zypper in tcpdump
2. Stop the OSA dispatcher and Jabber processes with rcosa-dispatcher stop and rcjabberd stop.
3. Start data capture on port 5222: tcpdump -s 0 port 5222 -w server_dump.pcap
4. Start the OSA dispatcher and Jabber processes: rcosa-dispatcher start and rcjabberd start.
5. Open a second terminal and execute the following commands: rcosa-dispatcher start and rcjabberd start.

6. Operate the SUSE Manager server and clients so the bug you formerly experienced is reproduced.
7. Once you have finished your capture re-open terminal 1 and stop the capture of data with:



PROCEDURE: CLIENT SIDE CAPTURE

1. Install the tcpdump package on your client as root: zypper in tcpdump
2. Stop the OSA process: rcosad stop.
3. Begin data capture on port 5222: tcpdump -s 0 port 5222 -w client_client_dump.pcap
4. Open a second terminal and start the OSA process: rcosad start
5. Operate the SUSE Manager server and clients so the bug you formerly experienced is reproduced.
6. Once you have finished your capture re-open terminal 1 and stop the capture of data with:



17.3.4 Engineering Notes: Analyzing Captured Data

This section provides information on analyzing the previously captured data from client and server.

1. Obtain the certificate file from your SUSE Manager server: /etc/pki/spacewalk/jabberd/server.pem
2. Edit the certificate file removing all lines before -----BEGIN RSA PRIVATE KEY-----, save it as key.pem
3. Install Wireshark as root with: zypper in wireshark
4. Open the captured file in wireshark.
5. From *Edit > [menu:Preferences]* select SSL from the left pane.
6. Select RSA keys list: *Edit > [menu:New]*
 - IP Address any
 - Port: 5222

- Protocol: xmpp
- Key File: open the key.pem file previously edited.
- Password: leave blank

For more information see also:

- <https://wiki.wireshark.org/SSL> ↗
- https://bugs.wireshark.org/bugzilla/show_bug.cgi?id=3444 ↗

17.4 RPC Connection Timeout Settings

RPC connection timeouts are configurable on the Uyuni server, SUSE Manager Proxy server, and the clients. For example, if package downloads take longer then expected, you can increase timeout values. `spacewalk-proxy restart` should be run after the setting is added or modified. Set the following variables to a value in seconds specifying how long an RPC connection may take at maximum:

Server - `/etc/rhn/rhn.conf`

```
server.timeout =`number`
```

Proxy Server - `/etc/rhn/rhn.conf`

```
proxy.timeout =`number`
```

SUSE Linux Enterprise Server Clients (using zypp-plugin-spacewalk) - `/etc/zypp/zypp.conf`

```
## Valid values:  [0,3600]
## Default value: 180
download.transfer_timeout = 180
```

This is the maximum time in seconds that a transfer operation is allowed to take. This is useful for preventing batch jobs from hanging for hours due to slow networks or links going down. If limiting operations to less than a few minutes, you risk aborting perfectly normal operations.

```
timeout =`number`
```

17.5 Client/Server Package Inconsistency

In some cases, updates are available in the web interface, but not appearing on the client. If you schedule an update on the client, it will fail with an error stating that no updates are available. This can be caused by a metadata regeneration problem, or because update packages have been locked.

The notice that updates are available will appear immediately, but new metadata is only generated on the server after synchronizing. In this case, an inconsistency can occur if taskomatic crashes, or because taskomatic is still running and creating new metadata.

To address this issue, check the /var/log/rhn/rhn_taskomatic_daemon.log file to determine if any processes are still running, or an exception which could indicate a crash. In the case of a crash, restart taskomatic.

Check package locks and exclude lists to determine if packages are locked or excluded on the client:

On Expanded Support Platform, check /etc/yum.conf and search for exclude=.

On SUSE Linux Enterprise Server, use the zypper locks command.

17.6 Corrupted Repository Data

If the information in /var/cache/rhn/repodata/sles12-sp3-updates-x86_64 becomes out of date, it will cause problems with updating the server. The repository data file can be regenerated using the spacecmd command:

PROCEDURE: REBUILD REPODATA FILE

1. Remove all files from /var/cache/rhn/repodata/sles12-sp3-updates-x86_64
2. Regenerate the file with spacecmd softwarechannel_regenerateyumcache sles12-sp3-updates-x86_64

17.7 Unable to Get Local Issuer Certificate

Some older bootstrap scripts will will create a link to the local certificate in the wrong place, which can cause problems with zypper returning an Unrecognized error about the local issuer certificate. In this case, ensure that the link to the local issuer certificate has been created in /etc/ssl/certs/, and consider updating your bootstrap scripts.

18 Additional Resources

This chapter contains links to helpful resources.

18.1 Learning YAML Syntax for Salt States

Learn how to write states and pillars with YAML.

- <https://docs.saltstack.com/en/latest/topics/yaml/index.html> ↗
- <http://yaml-online-parser.appspot.com/> ↗

18.2 Getting Started with Jinja Templates

Learn how to begin writing templates in Jinja

- <https://docs.saltstack.com/en/latest/topics/jinja/index.html> ↗
- <http://jinja.pocoo.org/docs/dev/> ↗

18.3 Salt Best Practices

Best practices from the Salt team.

- https://docs.saltstack.com/en/latest/topics/best_practices.html ↗

19 A Uyuni 2.1 and 3.2 Product Comparison

You may already have experience managing your systems using Uyuni 2.1. The good news is all the features you are used to working with, (regarding the traditional stack) have not changed. The only real exception is that the original built-in monitoring feature has been removed. Icinga, a third party monitoring solution is included within the SUSE Manager Tools Channel for SLES 12. Uyuni 3.2 supports managing systems via the popular IT orchestration engine Salt, in addition to the previously existing traditional management stack.

Important: Selecting a Management Method

You cannot and should not mix a single system with both methods although you can have Salt managed systems and traditionally managed systems coexisting and managed by a Uyuni server. You may for example have a development department and assign all machines in this department as Salt minions, likewise you could also have a production department and assign machines as traditional bootstrap clients. Remember a single machine is either traditionally managed or Salt but never both.

Keep in mind that minions are not traditional clients and their feature set is currently limited. Future maintenance updates will provide feature parity over time and your feedback for prioritization of these features is welcome! The following tables provide a comparison between each feature set. This includes features in development and features available only to their parent management stack.

TABLE 19.1: COMPARING TRADITIONAL MANAGEMENT AND SALT MANAGEMENT

Feature/Function	Traditional Management	Salt Management
Registration	bootstrap/rhnreg_ks	bootstrap scripts and Web UI
Install Packages	Supported	Supported
Install Patches	Supported	Supported
Remote Commands	Supported	Supported
System Package States	Unsupported	Supported
System Custom States	Unsupported	Supported

Feature/Function	Traditional Management	Salt Management
Group Custom States	Unsupported	Supported
Organization Custom States	Unsupported	Supported
System Set Manager	Supported	Supported
Service Pack Migration	Supported	Supported
Virtualization Host Management: <i>Auto-installation/bare metal installation support</i>	Supported	Supported (read-only)
System Redeployment: <i>With Auto-installation</i>	Supported	Coming Soon
Contact Methods: <i>How the server communicates with a client</i>	osad, rhnsd, ssh push	zeromq: Salt default salt-ssh
Red Hat Network Clients <i>RHEL 6, 7</i>	Supported	Supported
SUSE Manager Proxy	Supported	Supported
Action Chains	Supported	Supported
Software Crash Reporting	Supported	Unsupported
Staging	Supported	Supported
Duplicate Package Reporting <i>Example: Multiple Versions of the Linux Kernel</i>	Supported	Supported
SCAP Auditing	Supported	Supported
Support for Multiple Organizations	Supported **	Supported **

Feature/Function	Traditional Management	Salt Management
Package Verification	Supported	Under Review
System Locking	Supported	Under Review
Configuration File Management	Supported	Supported
Snapshots and Profiles	Supported	Under Review (Profiles are supported)
Power Management	Supported	Coming Soon



Warning: Isolation Enforcement **

Uyuni 2.1 is multi-tenant and organizations are completely isolated from one another. This isolation includes both privacy and security.

For example: User A in Org_1 cannot see user B in Org_2. (This relates to any data specific to an organization including: servers, channels, activation keys, configuration channels, files and so on.)

In Uyuni 3.2 Salt currently does not support any level of multi-tenancy and therefore information specific to an organization is accessible across organizations. For example:

```
salt '*' cmd.run "some_dangerous_command"
```

The above command will target all organization, groups and single systems including their files, channels, activation keys etc. This should be kept in mind when working with Salt.

The following table provides an overview of differences in functionality between Uyuni 2.1 and 3.2.

TABLE 19.2: COMPARING UYUNI 2.1 AND 3.2 FUNCTIONALITY

Functionality	SUSE Manager 2.1	SUSE Manager 3.2
Configuration Management	Based on Static Configuration	Redesigned with Salt Integration
Configuration Management	No Concept of States	States are Supported

Functionality	SUSE Manager 2.1	SUSE Manager 3.2
Subscription Management	Limited Functionality	New Design, Full Featured
Monitoring	Traditional Monitoring Supported until End of Life	Nagios Compatible, Icinga Monitoring Server is Included
Installation Approach	Appliance Based	Installed as an Add-on
Compatibility	Compatibility Carried Forward to Uyuni 3	Maintains full Uyuni 2.1 Functionality <i>Traditional Monitoring Removed</i>

A GNU Licenses

This appendix contains the GNU Free Documentation License version 1.2.

20 GNU Free Documentation License

Copyright © 2000, 2001, 2002 Free Software Foundation, Inc. 51 Franklin St, Fifth Floor, Boston, MA 02110-1301 USA. Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

0. PREAMBLE

The purpose of this License is to make a manual, textbook, or other functional and useful document "free" in the sense of freedom: to assure everyone the effective freedom to copy and redistribute it, with or without modifying it, either commercially or noncommercially. Secondly, this License preserves for the author and publisher a way to get credit for their work, while not being considered responsible for modifications made by others.

This License is a kind of "copyleft", which means that derivative works of the document must themselves be free in the same sense. It complements the GNU General Public License, which is a copyleft license designed for free software.

We have designed this License in order to use it for manuals for free software, because free software needs free documentation: a free program should come with manuals providing the same freedoms that the software does. But this License is not limited to software manuals; it can be used for any textual work, regardless of subject matter or whether it is published as a printed book. We recommend this License principally for works whose purpose is instruction or reference.

1. APPLICABILITY AND DEFINITIONS

This License applies to any manual or other work, in any medium, that contains a notice placed by the copyright holder saying it can be distributed under the terms of this License. Such a notice grants a world-wide, royalty-free license, unlimited in duration, to use that work under the conditions stated herein. The "Document", below, refers to any such manual or work. Any member of the public is a licensee, and is addressed as "you". You accept the license if you copy, modify or distribute the work in a way requiring permission under copyright law.

A "Modified Version" of the Document means any work containing the Document or a portion of it, either copied verbatim, or with modifications and/or translated into another language.

A "Secondary Section" is a named appendix or a front-matter section of the Document that deals exclusively with the relationship of the publishers or authors of the Document to the Document's overall subject (or to related matters) and contains nothing that could fall directly within that overall subject. (Thus, if the Document is in part a textbook of mathematics, a Secondary Section may not explain any mathematics.) The relationship could be a matter of historical connection with the subject or with related matters, or of legal, commercial, philosophical, ethical or political position regarding them.

The "Invariant Sections" are certain Secondary Sections whose titles are designated, as being those of Invariant Sections, in the notice that says that the Document is released under this License. If a section does not fit the above definition of Secondary then it is not allowed to be designated as Invariant. The Document may contain zero Invariant Sections. If the Document does not identify any Invariant Sections then there are none.

The "Cover Texts" are certain short passages of text that are listed, as Front-Cover Texts or Back-Cover Texts, in the notice that says that the Document is released under this License. A Front-Cover Text may be at most 5 words, and a Back-Cover Text may be at most 25 words.

A "Transparent" copy of the Document means a machine-readable copy, represented in a format whose specification is available to the general public, that is suitable for revising the document straightforwardly with generic text editors or (for images composed of pixels) generic paint programs or (for drawings) some widely available drawing editor, and that is suitable for input to text formatters or for automatic translation to a variety of formats suitable for input to text formatters. A copy made in an otherwise Transparent file format whose markup, or absence of markup, has been arranged to thwart or discourage subsequent modification by readers is not Transparent. An image format is not Transparent if used for any substantial amount of text. A copy that is not "Transparent" is called "Opaque".

Examples of suitable formats for Transparent copies include plain ASCII without markup, Texinfo input format, LaTeX input format, SGML or XML using a publicly available DTD, and standard-conforming simple HTML, PostScript or PDF designed for human modification. Examples of transparent image formats include PNG, XCF and JPG. Opaque formats include proprietary formats that can be read and edited only by proprietary word processors, SGML or XML for which the DTD and/or processing tools are not generally available, and the machine-generated HTML, PostScript or PDF produced by some word processors for output purposes only.

The "Title Page" means, for a printed book, the title page itself, plus such following pages as are needed to hold, legibly, the material this License requires to appear in the title page. For works in formats which do not have any title page as such, "Title Page" means the text near the most prominent appearance of the work's title, preceding the beginning of the body of the text.

A section "Entitled XYZ" means a named subunit of the Document whose title either is precisely XYZ or contains XYZ in parentheses following text that translates XYZ in another language. (Here XYZ stands for a specific section name mentioned below, such as "Acknowledgements", "Dedications", "Endorsements", or "History".) To "Preserve the Title" of such a section when you modify the Document means that it remains a section "Entitled XYZ" according to this definition. The Document may include Warranty Disclaimers next to the notice which states that this License applies to the Document. These Warranty Disclaimers are considered to be included by reference in this License, but only as regards disclaiming warranties: any other implication that these Warranty Disclaimers may have is void and has no effect on the meaning of this License.

2. VERBATIM COPYING

You may copy and distribute the Document in any medium, either commercially or noncommercially, provided that this License, the copyright notices, and the license notice saying this License applies to the Document are reproduced in all copies, and that you add no other conditions whatsoever to those of this License. You may not use technical measures to obstruct or control the reading or further copying of the copies you make or distribute. However, you may accept compensation in exchange for copies. If you distribute a large enough number of copies you must also follow the conditions in section 3.

You may also lend copies, under the same conditions stated above, and you may publicly display copies.

3. COPYING IN QUANTITY

If you publish printed copies (or copies in media that commonly have printed covers) of the Document, numbering more than 100, and the Document's license notice requires Cover Texts, you must enclose the copies in covers that carry, clearly and legibly, all these Cover Texts: Front-Cover Texts on the front cover, and Back-Cover Texts on the back cover. Both covers must also clearly and legibly identify you as the publisher of these copies. The front cover must present the full title with all words of the title equally prominent and visible. You may add other material on the covers in addition. Copying with changes limited to the covers, as long as they preserve the title of the Document and satisfy these conditions, can be treated as verbatim copying in other respects.

If the required texts for either cover are too voluminous to fit legibly, you should put the first ones listed (as many as fit reasonably) on the actual cover, and continue the rest onto adjacent pages.

If you publish or distribute Opaque copies of the Document numbering more than 100, you must either include a machine-readable Transparent copy along with each Opaque copy, or state in or with each Opaque copy a computer-network location from which the general network-using public has access to download using public-standard network protocols a complete Transparent copy of the Document, free of added material. If you use the latter option, you must take reasonably prudent steps, when you begin distribution of Opaque copies in quantity, to ensure that this Transparent copy will remain thus accessible at the stated location until at least one year after the last time you distribute an Opaque copy (directly or through your agents or retailers) of that edition to the public.

It is requested, but not required, that you contact the authors of the Document well before redistributing any large number of copies, to give them a chance to provide you with an updated version of the Document.

4. MODIFICATIONS

You may copy and distribute a Modified Version of the Document under the conditions of sections 2 and 3 above, provided that you release the Modified Version under precisely this License, with the Modified Version filling the role of the Document, thus licensing distribution and modification of the Modified Version to whoever possesses a copy of it. In addition, you must do these things in the Modified Version:

- A. Use in the Title Page (and on the covers, if any) a title distinct from that of the Document, and from those of previous versions (which should, if there were any, be listed in the History section of the Document). You may use the same title as a previous version if the original publisher of that version gives permission.
- B. List on the Title Page, as authors, one or more persons or entities responsible for authorship of the modifications in the Modified Version, together with at least five of the principal authors of the Document (all of its principal authors, if it has fewer than five), unless they release you from this requirement.
- C. State on the Title page the name of the publisher of the Modified Version, as the publisher.
- D. Preserve all the copyright notices of the Document.

- E. Add an appropriate copyright notice for your modifications adjacent to the other copyright notices.
- F. Include, immediately after the copyright notices, a license notice giving the public permission to use the Modified Version under the terms of this License, in the form shown in the Addendum below.
- G. Preserve in that license notice the full lists of Invariant Sections and required Cover Texts given in the Document's license notice.
- H. Include an unaltered copy of this License.
- I. Preserve the section Entitled "History", Preserve its Title, and add to it an item stating at least the title, year, new authors, and publisher of the Modified Version as given on the Title Page. If there is no section Entitled "History" in the Document, create one stating the title, year, authors, and publisher of the Document as given on its Title Page, then add an item describing the Modified Version as stated in the previous sentence.
- J. Preserve the network location, if any, given in the Document for public access to a Transparent copy of the Document, and likewise the network locations given in the Document for previous versions it was based on. These may be placed in the "History" section. You may omit a network location for a work that was published at least four years before the Document itself, or if the original publisher of the version it refers to gives permission.
- K. For any section Entitled "Acknowledgements" or "Dedications", Preserve the Title of the section, and preserve in the section all the substance and tone of each of the contributor acknowledgements and/or dedications given therein.
- L. Preserve all the Invariant Sections of the Document, unaltered in their text and in their titles. Section numbers or the equivalent are not considered part of the section titles.
- M. Delete any section Entitled "Endorsements". Such a section may not be included in the Modified Version.
- N. Do not retitle any existing section to be Entitled "Endorsements" or to conflict in title with any Invariant Section.
- O. Preserve any Warranty Disclaimers.

If the Modified Version includes new front-matter sections or appendices that qualify as Secondary Sections and contain no material copied from the Document, you may at your option designate some or all of these sections as invariant. To do this, add their titles to the list of Invariant Sections in the Modified Version's license notice. These titles must be distinct from any other section titles.

You may add a section Entitled "Endorsements", provided it contains nothing but endorsements of your Modified Version by various parties—for example, statements of peer review or that the text has been approved by an organization as the authoritative definition of a standard.

You may add a passage of up to five words as a Front-Cover Text, and a passage of up to 25 words as a Back-Cover Text, to the end of the list of Cover Texts in the Modified Version. Only one passage of Front-Cover Text and one of Back-Cover Text may be added by (or through arrangements made by) any one entity. If the Document already includes a cover text for the same cover, previously added by you or by arrangement made by the same entity you are acting on behalf of, you may not add another; but you may replace the old one, on explicit permission from the previous publisher that added the old one.

The author(s) and publisher(s) of the Document do not by this License give permission to use their names for publicity for or to assert or imply endorsement of any Modified Version.

5. COMBINING DOCUMENTS

You may combine the Document with other documents released under this License, under the terms defined in section 4 above for modified versions, provided that you include in the combination all of the Invariant Sections of all of the original documents, unmodified, and list them all as Invariant Sections of your combined work in its license notice, and that you preserve all their Warranty Disclaimers.

The combined work need only contain one copy of this License, and multiple identical Invariant Sections may be replaced with a single copy. If there are multiple Invariant Sections with the same name but different contents, make the title of each such section unique by adding at the end of it, in parentheses, the name of the original author or publisher of that section if known, or else a unique number. Make the same adjustment to the section titles in the list of Invariant Sections in the license notice of the combined work.

In the combination, you must combine any sections Entitled "History" in the various original documents, forming one section Entitled "History"; likewise combine any sections Entitled "Acknowledgements", and any sections Entitled "Dedications". You must delete all sections Entitled "Endorsements".

6. COLLECTIONS OF DOCUMENTS

You may make a collection consisting of the Document and other documents released under this License, and replace the individual copies of this License in the various documents with a single copy that is included in the collection, provided that you follow the rules of this License for verbatim copying of each of the documents in all other respects.

You may extract a single document from such a collection, and distribute it individually under this License, provided you insert a copy of this License into the extracted document, and follow this License in all other respects regarding verbatim copying of that document.

7. AGGREGATION WITH INDEPENDENT WORKS

A compilation of the Document or its derivatives with other separate and independent documents or works, in or on a volume of a storage or distribution medium, is called an "aggregate" if the copyright resulting from the compilation is not used to limit the legal rights of the compilation's users beyond what the individual works permit. When the Document is included in an aggregate, this License does not apply to the other works in the aggregate which are not themselves derivative works of the Document.

If the Cover Text requirement of section 3 is applicable to these copies of the Document, then if the Document is less than one half of the entire aggregate, the Document's Cover Texts may be placed on covers that bracket the Document within the aggregate, or the electronic equivalent of covers if the Document is in electronic form. Otherwise they must appear on printed covers that bracket the whole aggregate.

8. TRANSLATION

Translation is considered a kind of modification, so you may distribute translations of the Document under the terms of section 4. Replacing Invariant Sections with translations requires special permission from their copyright holders, but you may include translations of some or all

Invariant Sections in addition to the original versions of these Invariant Sections. You may include a translation of this License, and all the license notices in the Document, and any Warranty Disclaimers, provided that you also include the original English version of this License and the original versions of those notices and disclaimers. In case of a disagreement between the translation and the original version of this License or a notice or disclaimer, the original version will prevail.

If a section in the Document is Entitled "Acknowledgements", "Dedications", or "History", the requirement (section 4) to Preserve its Title (section 1) will typically require changing the actual title.

9. TERMINATION

You may not copy, modify, sublicense, or distribute the Document except as expressly provided for under this License. Any other attempt to copy, modify, sublicense or distribute the Document is void, and will automatically terminate your rights under this License. However, parties who have received copies, or rights, from you under this License will not have their licenses terminated so long as such parties remain in full compliance.

10. FUTURE REVISIONS OF THIS LICENSE

The Free Software Foundation may publish new, revised versions of the GNU Free Documentation License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns. See <http://www.gnu.org/copyleft/>.

Each version of the License is given a distinguishing version number. If the Document specifies that a particular numbered version of this License "or any later version" applies to it, you have the option of following the terms and conditions either of that specified version or of any later version that has been published (not as a draft) by the Free Software Foundation. If the Document does not specify a version number of this License, you may choose any version ever published (not as a draft) by the Free Software Foundation.

ADDENDUM: How to use this License for your documents

Copyright (c) YEAR YOUR NAME.

Permission is granted to copy, distribute and/or modify this document
under the terms of the GNU Free Documentation License, Version 1.2

```
or any later version published by the Free Software Foundation;  
with no Invariant Sections, no Front-Cover Texts, and no Back-Cover Texts.  
A copy of the license is included in the section entitled“GNU  
Free Documentation License”.
```

If you have Invariant Sections, Front-Cover Texts and Back-Cover Texts, replace the “ with... Texts.” line with this:

```
with the Invariant Sections being LIST THEIR TITLES, with the  
Front-Cover Texts being LIST, and with the Back-Cover Texts being LIST.
```

If you have Invariant Sections without Cover Texts, or some other combination of the three, merge those two alternatives to suit the situation.

If your document contains nontrivial examples of program code, we recommend releasing these examples in parallel under your choice of free software license, such as the GNU General Public License, to permit their use in free software.